
Derivative-Free Policy Optimization for Linear Risk-Sensitive and Robust Control Design: Implicit Regularization and Sample Complexity

Kaiqing Zhang^{*1} Xiangyuan Zhang^{*1} Bin Hu¹ Tamer Başar¹

Abstract

Policy-based model-free reinforcement learning (RL) methods have shown great promise for continuous control applications. However, their performances on risk-sensitive/robust control tasks have not been fully understood, which has been generally considered to be one important open problem in the seminal work (Fazel et al., 2018). We make a step toward addressing this open problem, by providing the first sample complexity results for policy gradient (PG) methods in two fundamental risk-sensitive/robust control settings: the linear exponential quadratic Gaussian, and the linear-quadratic (LQ) disturbance attenuation problems. The optimization landscapes for these problems are by nature more challenging than that of the LQ regulator problem, due to lack of coercivity of their objective functions. To overcome this challenge, we obtain the first *implicit regularization* results for model-free PG methods, certifying that the controller *remains robust* during the learning process, which further lead to the sample complexity guarantees. As a by-product, our results also provide the first sample complexity of PG methods in two-player zero-sum LQ dynamic games, a baseline in multi-agent RL.

1. Introduction

Recent years have witnessed the rapid development of reinforcement learning (RL) methods in handling continuous control tasks (Schulman et al., 2015b; Lillicrap et al., 2015; Recht, 2019). Central to the success of RL are policy optimization (PO) methods, including policy gradient (PG) (Sutton et al., 2000; Kakade, 2002), actor-critic (Konda &

Tsitsiklis, 2000; Bhatnagar et al., 2009), and other variants (Schulman et al., 2017; 2015a). Progress reported in the literature has clearly shown an increasing interest in understanding theoretical properties of PO methods for relatively simple baselines such as various linear control problems (Fazel et al., 2018; Tu & Recht, 2019; Bu et al., 2019a; Zhang et al., 2019; 2021a; Mohammadi et al., 2019; Furieri et al., 2020; Malik et al., 2020; Jansch-Porto et al., 2020a; Fatkhullin & Polyak, 2020). However, the theory of model-free PO methods on *risk-sensitive/robust* control remains underdeveloped in the literature. Since risk-sensitivity and robustness are important issues for designing safety-critical systems, it is natural to bring up the questions of whether and how model-free PO methods would converge for these continuous control tasks.

Our work in this paper is motivated by the above concern, and studies the sample complexity of model-free PG methods on two important baseline problems in risk-sensitive/robust control, namely the linear exponential quadratic Gaussian (LEQG), and the linear quadratic (LQ) disturbance attenuation problems. The former covers a fundamental setting in risk-sensitive control, and the latter is an important baseline for robust control. Based on the well-known equivalence between these problems and LQ dynamic games (Jacobson, 1973; Glover & Doyle, 1988; Başar, 1991; Başar & Bernhard, 1995; Başar & Olsder, 1998), we develop a *unified* PO perspective for both. A common feature for the above two problems is that their optimization landscapes are by nature more challenging than that of the linear quadratic regulator (LQR) problem, and existing proof techniques for model-free PG methods (Fazel et al., 2018; Mohammadi et al., 2019; Furieri et al., 2020; Malik et al., 2020) are no longer effective due to lack of coercivity of the objective functions. Specifically, when applying PG methods to LQR, the feasible set for the resultant constrained optimization problem is the set of all linear state-feedback controllers that *stabilize* the closed-loop dynamics. The objective function of LQR is coercive on this feasible set and serves as a barrier function itself (Bu et al., 2019a), guaranteeing for the PG iterates to stay in the feasible set and converge to the globally optimal controller. For the LEQG and LQ disturbance attenuation problems, the risk-sensitivity/robustness conditions have reshaped the

^{*}Equal contribution ¹Department of Electrical and Computer Engineering and the Coordinated Science Laboratory, University of Illinois at Urbana-Champaign, Urbana, IL, USA. Correspondence to: Kaiqing Zhang <kzhang66@illinois.edu>, Xiangyuan Zhang <xz7@illinois.edu>.

feasible set in a way that the objective function becomes non-coercive, i.e. the objective value can remain finite while approaching the boundary. The objective is no longer a barrier function, and new proof techniques are needed to show that model-free PG iterates will stay in the feasible set. This is significant for safety-critical control systems with model uncertainty, as the iterates' feasibility here is equivalent to *risk-sensitivity/robustness* of the controller (cf. Remark A.9), and the failure to preserve robustness during learning can cause catastrophic effects, e.g., destabilizing the system in face of disturbances.

The most relevant result was developed in (Zhang et al., 2021a), which proposes implicit regularization (IR) arguments to show the convergence of several PG methods on the mixed $\mathcal{H}_2/\mathcal{H}_\infty$ control design problem (which can be viewed as the infinite-horizon variant of the LQ disturbance attenuation problem studied in this paper). The main finding there is that two specific PG search directions (with perfect model information) are automatically biased towards the interior of the robustness-related feasible set. In (Zhang et al., 2021a), it is emphasized that IR is a feature of both the *problem* and the *algorithm*, contrasting to that the stability-preserving nature of PG methods for LQR problems is based on the barrier function property of the objective and hence is *algorithm-agnostic*. Although the idea of IR is relevant to PO problems with non-coercive objective functions, the arguments in (Zhang et al., 2021a) only apply to the setting with a *known* model, since they rely on a specialized perturbation technique which may potentially generate arbitrarily small “margins”. However, in the model-free setting, a uniform margin is required for provable tolerance of statistical errors. See a detailed comparison of the literature in §A.1.

In this paper, for the LEQG and LQ disturbance attenuation problems, we overcome the above margin issue and obtain the first IR result in the model-free setting. This enables the first model-free PG method that provably solves these control problems with polynomial number of samples. We highlight our contributions as follows.

Contributions. We provide the first sample complexity results for model-free PG methods for solving linear control problems with risk-sensitivity/robustness concerns (the LEQG and LQ disturbance attenuation problems), which were viewed as important open problems in the seminal work (Fazel et al., 2018). From the robust control perspective, one feature of our algorithms is that, during the learning process, a certain level of *robustness/risk-sensitivity* of the controller is proved to be preserved. This has generalized the results in (Zhang et al., 2021a) with a known model, and has thus enabled the first finite-sample convergence guarantees of PO methods for risk-sensitive/robust control design. Our algorithms and sample complexity results also address two-player zero-sum LQ dynamic games in the finite-horizon

time-varying setting, which are among the first sample complexity results for the global convergence of policy-based methods for competitive multi-agent RL. Second, in the context of minimax optimization, our results address a class of nonconvex-nonconcave minimax constrained optimization problems, using *zeroth-order* multi-step gradient descent-ascent methods. Finally, part of our results provide the first sample complexity analysis for PG methods that solve the finite-horizon time-varying LQR problem with system noises and a possibly *indefinite* state-weighting matrix.

2. Background

We first introduce LEQG, and the LQ disturbance attenuation problem, and their equivalence to zero-sum LQ dynamic games, deferring the extended formulations to §A.3.

2.1. Linear Exponential Quadratic Gaussian

Consider the LEQG problem (Jacobson, 1973) characterized by $x_{t+1} = A_t x_t + B_t u_t + w_t$, where $x_t \in \mathbb{R}^m$ represents the state; $u_t \in \mathbb{R}^d$ is the control; $w_t \in \mathbb{R}^m$ is an independent (across time) Gaussian random noise with $w_t \sim \mathcal{N}(0, W)$ for some $W > 0$; and the initial state $x_0 \sim \mathcal{N}(0, X_0)$ for some $X_0 > 0$, which is independent of $\{w_t\}$. The objective is to find a control policy that minimizes $\mathcal{J}(\{u_t\}) := \frac{2}{\beta} \log \mathbb{E} \exp \left[\frac{\beta}{2} \left(\sum_{t=0}^{N-1} c_t + x_N^\top Q_N x_N \right) \right]$, where $c_t := x_t^\top Q_t x_t + u_t^\top R_t u_t$; $Q_t, Q_N \geq 0$ and $R_t > 0$; and $\beta > 0$ is the degree of risk-sensitivity.

2.2. LQ Disturbance Attenuation

Second, consider the LQ disturbance attenuation problem described by $x_{t+1} = A_t x_t + B_t u_t + D_t w_t$ and $z_t = C_t x_t + E_t u_t$, where $x_t \in \mathbb{R}^m$ is the state; $u_t \in \mathbb{R}^d$ is the control; $w_t \in \mathbb{R}^n$ is the (unknown) disturbance; and $z_t \in \mathbb{R}^l$ is the output. Subsequently, we denote the ℓ^2 -norms of $\omega := [x_0^\top C_0^\top, w_0^\top, \dots, w_{N-1}^\top]^\top$ and $z := [z_0^\top, \dots, z_{N-1}^\top, x_N^\top Q_N^{1/2}]^\top$ as $\|\omega\|$ and $\|z\| = \{\mathcal{C}(\{u_t\}, \{w_t\})\}^{1/2}$, where $Q_N \geq 0$. Then, our goal is the following: Given a $\gamma > \gamma^*$, where $\gamma^* > 0$ is the *minimax* level of disturbance attenuation (Başar & Bernhard, 1995), find a control policy $\mu_t = -K_t x_t$ that solves $\min_{\{\mu_t\}} \overline{\mathcal{J}}(\{\mu_t(x_t)\}, \{w_t\})$ subject to $\sup_{x_0, \{w_t\}} \frac{\|z\|}{\|\omega\|} < \gamma$, where $\overline{\mathcal{J}}(\{\mu_t(x_t)\})$ is an upper bound of the LQG cost (Mustafa & Bernstein, 1991). A discussion regarding the challenges when attempting to address LEQG and LQ disturbance attenuation directly using derivative-free PG methods is provided in Remark A.3.

2.3. An Equivalent Dynamic Game Formulation

To overcome the challenges in Remark A.3, we introduce an equivalent *dynamic game* formulation, dictated by the dynamics $x_{t+1} = A_t x_t + B_t u_t + D_t w_t + \xi_t$, where

$x_t \in \mathbb{R}^m$ is the system state, and $u_t := -K_t x_t \in \mathbb{R}^d$ (resp., $w_t := -L_t x_t \in \mathbb{R}^n$) is the linear state-feedback controller of the minimizing (resp., maximizing) player, which suffices to attain the saddle point whenever exists (Başar & Olsder, 1998). The initial state x_0 and the noises ξ_t are drawn independently from \mathcal{D} , where \mathcal{D} is zero-mean, has a positive-definite covariance, and satisfies almost surely that $\|x_0\|, \|\xi_t\| \leq \vartheta$ for some ϑ . The goal is to solve the game:

$$\inf_{\{K_t\}} \sup_{\{L_t\}} \mathbb{E}_{x_0, \xi_0, \dots, \xi_{N-1}} \left[\sum_{t=0}^{N-1} g_t + x_N^\top Q_N x_N \right], \quad (2.1)$$

where $g_t := x_t^\top (Q_t + K_t^\top R_t^u K_t - L_t^\top R_t^w L_t) x_t$; $Q_t, Q_N \geq 0$, $R_t^u, R_t^w > 0$ are symmetric weighting matrices.

To characterize the solution to (2.1), define the Riccati difference equation (RDE) $P_t^* = Q_t + A_t^\top P_{t+1}^* \Lambda_t^{-1} A_t$, where $\Lambda_t := I + (B_t(R_t^u)^{-1} B_t^\top - D_t(R_t^w)^{-1} D_t^\top) P_{t+1}^*$ and $P_N^* = Q_N$. By (Başar & Olsder, 1998), whenever exists, the saddle-point control policies are linear state-feedback, with unique $K_t^* = (R_t^u)^{-1} B_t^\top P_{t+1}^* \Lambda_t^{-1} A_t$ and $L_t^* = -(R_t^w)^{-1} D_t^\top P_{t+1}^* \Lambda_t^{-1} A_t$. We introduce an assumption that ensures the existence of the value of the game.

Assumption 2.1 $R_t^w - D_t^\top P_{t+1}^* D_t > 0$ for all t .

Under Assumption 2.1, the value in (2.1) is attained by the sequence $(\{K_t^*\}, \{L_t^*\})$. We provide some notes on Assumption 2.1 in Remark A.5 and state the following equivalence conditions, with a short proof provided in §B.3.

Lemma 2.2 For any fixed $\gamma > \gamma^*$ in §2.2, there exist an equivalent LEQG and an equivalent zero-sum LQ dynamic game. Specifically, if we set $\beta^{-1} I, R_t, C_t^\top C_t, W$ in §2.1, $\gamma^2 I, R_t, C_t^\top C_t, D_t D_t^\top$ in §2.2, and $R_t^u, R_t^w, Q_t, D_t D_t^\top$ in §2.3 to be the same, for all t , then the optimal gain matrices in LEQG, the gain matrix in the LQ disturbance attenuation problem, and the Nash equilibrium gain matrix for the minimizing player in the game are the same.

By Lemma 2.2 and Remark A.3, we will hereafter focus on solving the game formulation in §2.3 using PG methods.

3. Policy Gradient Methods

For ease of analysis, we adopt the compact notations in §A.4 and denote $\mathcal{S}(d, m, N)$ and $\mathcal{S}(n, m, N)$ as the subspaces that we confine our searches of K and L to, respectively. In particular, we only search over $K \in \mathbb{R}^{Nd \times (N+1)m}$ and $L \in \mathbb{R}^{Nn \times (N+1)m}$, and K, L has the form of (A.17), without losing any optimality. Then, we have for any (K, L) , $\mathcal{G}(K, L) = \text{Tr}((Q + K^\top R^u K - L^\top R^w L) \Sigma_{K, L}) = \text{Tr}(P_{K, L} \Sigma_0)$, where $P_{K, L} := \text{diag}(P_{K_0, L_0}, \dots, P_{K_N, L_N})$ and P_{K_t, L_t} solves the recursive Lyapunov equation:

$$P_{K_t, L_t} = (A_t - B_t K_t - D_t L_t)^\top P_{K_{t+1}, L_{t+1}} (A_t - B_t K_t - D_t L_t) + Q_t + K_t^\top R_t^u K_t - L_t^\top R_t^w L_t, \quad P_{K_N, L_N} := Q_N. \quad (3.1)$$

Moreover, $P_{K, L}$ and $\Sigma_{K, L}$ solve the Lyapunov equations

$$P_{K, L} = (A - BK - DL)^\top P_{K, L} (A - BK - DL) + Q + K^\top R^u K - L^\top R^w L, \quad (3.2)$$

$$\Sigma_{K, L} = (A - BK - DL) \Sigma_{K, L} (A - BK - DL)^\top + \Sigma_0, \quad (3.3)$$

where $\Sigma_0 > 0$ is full-rank. Furthermore, we define the following notations that will be useful in our analysis.

$$c_0 := (N+1)\vartheta^2, \quad d_\Sigma := m^2(N+1), \quad \phi := \lambda_{\min}(\Sigma_0) > 0,$$

$$H_{K, L} := R^w - D^\top P_{K, L} D, \quad G_{K, L(K)} := R^u + B^\top \tilde{P}_{K, L(K)} B.$$

$$E_{K, L} := -H_{K, L} L - D^\top P_{K, L} (A - BK),$$

$$F_{K, L} := (R^u + B^\top P_{K, L} B) K - B^\top P_{K, L} (A - DL), \quad (3.4)$$

$$\tilde{P}_{K, L(K)} := P_{K, L(K)} + P_{K, L(K)} D H_{K, L(K)}^{-1} D^\top P_{K, L(K)}, \quad (3.5)$$

Our goal is to solve the minimax optimization problem $\min_K \max_L \mathcal{G}(K, L)$ such that $\mathcal{G}(K^*, L) \leq \mathcal{G}(K^*, L^*) \leq \mathcal{G}(K, L^*)$ for any K and L . Some properties of $\mathcal{G}(K, L)$ are presented here, with proofs in §B.4-§B.5.

Lemma 3.1 There exist zero-sum LQ dynamic games such that the objective function $\mathcal{G}(K, L)$ is nonconcave in L for a fixed K , and nonconvex in K for a fixed L .

Lemma 3.2 The PGs of $\mathcal{G}(K, L)$ can be computed as $\nabla_K \mathcal{G}(K, L) = 2F_{K, L} \Sigma_{K, L}$ and $\nabla_L \mathcal{G}(K, L) = 2E_{K, L} \Sigma_{K, L}$. Also, if at some stationary point (K, L) of $\mathcal{G}(K, L)$, it holds that $P_{K, L} \geq 0$ and $R^w - D^\top P_{K, L} D > 0$, then this (K, L) is the unique Nash equilibrium.

3.1. Double-Loop Scheme

We introduce a double-loop update scheme. First, we fix K and solve for the optimal L , denoted as $L(K)$. Then, we update K to minimize $\mathcal{G}(K, L(K))$ over $K \in \mathcal{S}(d, m, N)$. For each fixed K , the inner-loop is an indefinite LQR. We state the condition for the solution to the inner-loop problem to be well-defined as follows, deferring its proof to §B.6.

Lemma 3.3 For the Riccati equation:

$$P_{K, L(K)} = Q + K^\top R^u K + (A - BK)^\top (P_{K, L(K)} + P_{K, L(K)} D (R^w - D^\top P_{K, L(K)} D)^{-1} D^\top P_{K, L(K)}) (A - BK), \quad (3.6)$$

define the following set of outer-loop gain matrices K :

$$\mathcal{K} := \{K \in \mathcal{S}(d, m, N) \mid (3.6) \text{ admits a solution}$$

$$P_{K, L(K)} \geq 0, \text{ and } R^w - D^\top P_{K, L(K)} D > 0\}. \quad (3.7)$$

Then, $K \in \mathcal{K}$ is sufficient for the unique inner-loop solution $L(K) = (-R^w + D^\top P_{K, L(K)} D)^{-1} D^\top P_{K, L(K)} (A - BK)$ to be well defined. Also, $K \in \mathcal{K}$ is almost necessary, in that if $K \notin \mathcal{K}$, where $\mathcal{K} := \{K \in \mathcal{S}(d, m, N) \mid (3.6) \text{ admits a solution } P_{K, L(K)} \geq 0, \text{ and } R^w - D^\top P_{K, L(K)} D \geq 0\}$, then $L(K)$ is not well defined, i.e., $\mathcal{G}(K, \cdot)$ could be driven to arbitrarily large values. Lastly, the solution to (3.6) satisfies $P_{K, L(K)} \geq P_{K, L}$ in the p.s.d. sense for all L .




Figure 1. Left: Optimization landscape of LQR, where the dashed line represents the boundary of the stabilizing controller set. Right: Optimization landscape of the outer loop, with the dashed line representing the boundary of \mathcal{K} . The solid lines represent the contour lines of the objective function, K denotes the control gain of one iterate, and \star is the global minimizer.

By Lemmas 3.1, 3.3, and a variant of the Danskin's theorem (Bernhard & Rapaport, 1995), we can conclude that $\max_L \mathcal{G}(K, L)$ is differentiable with respect to K and $\nabla_K \{\max_L \mathcal{G}(K, L)\} = \nabla_K \mathcal{G}(K, L(K))$. The outer loop is a nonconvex (cf. Lemma 3.5) constrained optimization problem in K . The PG of the outer loop is computed as $\nabla_K \mathcal{G}(K, L(K)) = 2F_{K,L(K)} \Sigma_{K,L(K)}$. Lastly, we comment on the double-loop scheme in Remark A.7.

3.2. Optimization Landscape

For a fixed K , we present the optimization landscape of the inner-loop and defer the proof to §B.7.

Lemma 3.4 (Inner-Loop Landscape) For a fixed $K \in \mathcal{K}$, the maximization problem over $L \in \mathcal{S}(n, m, N)$ is nonconvex, $\mathcal{G}(K, L)$ is coercive in L , and the superlevel set

$$\mathcal{L}_K(a) := \{L \in \mathcal{S}(n, m, N) \mid \mathcal{G}(K, L) \geq a\} \quad (3.8)$$

is compact for any a such that $\mathcal{L}_K(a) \neq \emptyset$. Moreover, $\mathcal{G}(K, L)$ is $(l_{K,L}, \rho_{K,L})$ locally Lipschitz and $(\psi_{K,L}, \rho_{K,L})$ locally smooth at (K, L) . Also, $\mathcal{G}(K, L)$ is $l_{K,a}$ -globally Lipschitz and $\psi_{K,a}$ -globally smooth over $\mathcal{L}_K(a)$. Lastly, $\mathcal{G}(K, L)$ is μ_K -PL for some $\mu_K > 0$.

A discussion of the inner-loop landscape is provided in Remark A.8. Subsequently, we analyze the optimization landscape of the outer loop (proved in §B.8), subject to \mathcal{K} .

Lemma 3.5 There exist zero-sum LQ dynamic games such that $\mathcal{G}(K, L(K))$ is nonconvex and noncoercive in K . Moreover, the stationary point of $\mathcal{G}(K, L(K))$ in \mathcal{K} , denoted as $(K^*, L(K^*))$, is unique and constitutes the unique Nash equilibrium of the game.

Lack of the coercivity brings up challenges for convergence analysis, as a decrease in the value of the objective function cannot ensure feasibility of the updated gain matrix, in sharp contrast to the LQR problem (Fazel et al., 2018; Bu et al., 2019a). We illustrate the difficult landscape of the outer loop in Figure 1. To address this challenge, we will




Figure 2. Proof idea for Theorem 3.7. For any $K \in \mathcal{K}$ that induces a $P_{K,L(K)}$, denote the next iterate following (3.11) or (3.13) as K' . We construct an iterative argument backward in time to find a constant stepsize such that $P_{K'_t, L(K'_t)} \geq 0$ exists and satisfies $P_{K'_t, L(K'_t)} \leq P_{K_t, L(K_t)}$ for all t . Specifically, for any $t \in \{0, \dots, N-1\}$, a K'_t satisfying \blacksquare also satisfies \blacktriangle . Moreover, \blacklozenge is enforced by Assumption 2.1. Combined, $K' \in \mathcal{K}$ is guaranteed.

show next that the natural PG (NPG) and Gauss-Newton (GN) updates, with some stepsize choices, can automatically preserve the feasibility of the iterates on-the-fly, termed the *implicit regularization* (IR) property.

3.3. Update Rules and Global Convergence

We use $l, k \geq 0$ to represent the iteration indices of the inner- and outer-loop updates, respectively, and introduce:

$$\text{PG: } L_{l+1} = L_l + \eta \nabla_L \mathcal{G}(K_k, L_l), \quad (3.9)$$

$$\text{NPG: } L_{l+1} = L_l + \eta \nabla_L \mathcal{G}(K_k, L_l) \Sigma_{K_k, L_l}^{-1}, \quad (3.10)$$

$$K_{k+1} = K_k - \alpha \nabla_K \mathcal{G}(K_k, L(K_k)) \Sigma_{K_k, L(K_k)}^{-1}, \quad (3.11)$$

$$\text{GN: } L_{l+1} = L_l + \eta H_{K_k, L_l}^{-1} \nabla_L \mathcal{G}(K_k, L_l) \Sigma_{K_k, L_l}^{-1}, \quad (3.12)$$

$$K_{k+1} = K_k - \alpha G_{K_k, L(K_k)}^{-1} \nabla_K \mathcal{G}(K_k, L(K_k)) \Sigma_{K_k, L(K_k)}^{-1}, \quad (3.13)$$

where η, α are constant stepsizes. We first present the convergence results for three inner-loop PG updates.

Theorem 3.6 For a fixed $K \in \mathcal{K}$ and an arbitrary L_0 that induces a finite $\mathcal{G}(K, L_0)$, define $\mathcal{L}_K(a)$ as in (3.8), where $a < \mathcal{G}(K, L_0)$ is arbitrary. Then, iterates L_l of the updates (3.9), (3.10), and (3.12) with **PG**: $\eta \leq 1/\psi_{K,a}$; **NPG**: $\eta \leq 1/(2\|H_{K,L_0}\|)$; and **GN**: $\eta \leq 1/2$ converge to $L(K)$ with globally linear rates. Moreover, (3.12) with $\eta = 1/2$ converges to with a locally Q -quadratic rate.

For the outer loop, we require $K \in \mathcal{K}$. To that end, we introduce the IR property for the NPG (3.11) and GN (3.13) updates in Theorem 3.7, deferring its proof to §B.10.

Theorem 3.7 Let $K_0 \in \mathcal{K}$, **NPG**: $\alpha \leq 1/\|G_{K_0, L(K_0)}\|$, and **GN**: $\alpha \leq 1$. Then, $K_k \in \mathcal{K}$ for all $k \geq 0$. That is, the sequence of solutions to (3.6), $\{P_{K_k, L(K_k)}\}$, exists, and for all $k \geq 0$, $P_{K_k, L(K_k)}$ satisfies (3.7). Furthermore, the sequence $\{P_{K_k, L(K_k)}\}$ is monotonically non-increasing and bounded below by $P_{K^*, L(K^*)}$, in the p.s.d. sense.

We highlight the importance of IR in Remark A.9 and provide an illustration of the proof idea in Figure 2. Based on Theorem 3.7, we establish the convergence result for the outer loop, deferring its proof to §B.11.

Theorem 3.8 *Let $K_0 \in \mathcal{K}$, NPG: $\alpha \leq 1/(2\|G_{K_0,L(K_0)}\|)$, and GN: $\alpha \leq 1/2$. Then, the sequence of average natural gradient norm squares $\{k^{-1} \sum_{\kappa=0}^{k-1} \|F_{K_\kappa,L(K_\kappa)}\|_F^2\}$, $k \geq 1$, converges to 0 with $\mathcal{O}(1/k)$ rate. Moreover, this convergence is towards the unique Nash equilibrium. Lastly, updates (3.11) and (3.13) enjoy locally linear and Q -quadratic rates, respectively, around the Nash equilibrium.*

4. Derivative-Free Policy Gradient Methods

When the exact PG is not accessible, we propose a zeroth-order double-loop algorithm (see Algorithms 1-2 in §D).

4.1. Inner-Loop Maximization Oracle

Our inner-loop maximization oracle extends the sample complexity result in (Fazel et al., 2018) to finite-horizon time-varying indefinite LQR with system noises. The two specific zeroth-order PG updates are as follows:

$$\text{ZO-PG: } L_{l+1} = L_l + \eta \bar{\nabla}_L \mathcal{G}(K, L_l), \quad (4.1)$$

$$\text{ZO-NPG: } L_{l+1} = L_l + \eta \bar{\nabla}_L \mathcal{G}(K, L_l) \bar{\Sigma}_{K,L_l}^{-1}, \quad (4.2)$$

where $\bar{\nabla}_L \mathcal{G}(K, L)$, $\bar{\Sigma}_{K,L}$ are obtained from Algorithm 1. We now present the sample complexity of the inner loop.

Theorem 4.1 *Let $\epsilon_1, \delta_1 \in (0, 1)$, and parameters of Algorithm 1 in §D satisfy **Minibatch Size**: $M_1 \sim \tilde{\Omega}(\epsilon_1^{-2})$; **Stepsize**: $\eta \sim \mathcal{O}(1)$; and **Radius**: $r_1 \sim \Theta(\sqrt{\epsilon_1})$. Then, with probability at least $1 - \delta_1$ and a total of $L \sim \mathcal{O}(\log(\epsilon_1^{-1}))$ iterations, updates (4.1) and (4.2) output some $\bar{L}(K) := L_L$ satisfying $\mathcal{G}(K, \bar{L}(K)) \geq \mathcal{G}(K, L(K)) - \epsilon_1$ and $\|L(K) - \bar{L}(K)\|_F \leq \sqrt{\lambda_{\min}^{-1}(H_{K,L(K)}) \cdot \epsilon_1}$.*

Complete versions of Theorem 4.1 for updates (4.1) and (4.2) are given respectively in §B.12 and §B.13, along with their proofs. The total sample complexity of both (4.1) and (4.2) scales as $M_1 \cdot L \sim \tilde{\mathcal{O}}(\epsilon_1^{-2} \cdot \log(\delta_1^{-1}))$.

4.2. Outer-Loop ZO-NPG

With $\bar{L}(K)$ obtained from Algorithm 1, the outer-loop ZO-NPG algorithm approximately solves the constrained minimization problem $\min_{K \in \mathcal{K}} \mathcal{G}(K, \bar{L}(K))$, following:

$$\text{ZO-NPG: } K_{k+1} = K_k - \alpha \bar{\nabla}_K \mathcal{G}(K_k, \bar{L}(K_k)) \bar{\Sigma}_{K_k, \bar{L}(K_k)}^{-1}, \quad (4.3)$$

where $\bar{\nabla}_K \mathcal{G}(K, \bar{L}(K))$, $\bar{\Sigma}_{K, \bar{L}(K)}$ are obtained from Algorithm 2 in §D. In Theorem 4.2, with its complete version and proof being deferred to §B.14, we have that IR also holds in the derivative-free setting, with high probability. See Figure 3 for an illustration of the proof idea.




Figure 3. Illustrating the proof idea for Theorem 4.2. For $K_0 \in \mathcal{K}$, we define two compact sets $\hat{\mathcal{K}}$ and $\mathcal{K}_0 := \{K \in \mathcal{K}, \text{ and } P_{K,L(K)} \leq P_{K_0,L(K_0)}\}$ as shown in blue and red, respectively. We prove that the iterates following (4.3) stay within $\hat{\mathcal{K}}$ (thus, uniformly separated from $\partial \mathcal{K}$) with high probability.

Theorem 4.2 *For any $K_0 \in \mathcal{K}$, define $\hat{\mathcal{K}} := \{K \mid (3.6) \text{ admits a solution } P_{K,L(K)} \geq 0, \text{ and } P_{K,L(K)} \leq P_{K_0,L(K_0)} + \frac{\lambda_{\min}(H_{K_0,L(K_0)})}{2\|D\|^2} \cdot I\}$. Let $\delta_2 \in (0, 1)$, and parameters of Algorithm 2 in §D satisfy **Minibatch Size**: $M_2 \sim \tilde{\Omega}(K^4)$; **Stepsize**: $\alpha \sim \mathcal{O}(1)$; **Inner-Loop Specs**: $\epsilon_1 \sim \mathcal{O}(K^{-2})$, $\delta_1 \sim \mathcal{O}(\delta_2 M_2^{-1} K^{-1})$; and **Radius**: $r_2 \sim \mathcal{O}(K^{-1})$. Then, it holds with probability at least $1 - \delta_2$ that $K_k \in \hat{\mathcal{K}} \subset \mathcal{K}$ for all $k \in \{1, \dots, K\}$.*

Theorem 4.2 appears to be the first IR result of PO in robust control in the derivative-free setting. With Theorem 4.2 at hand, we now present the sample complexity of the outer-loop ZO-NPG algorithm (4.3), deferring its proof to §B.15.

Theorem 4.3 *For any $K_0 \in \mathcal{K}$, let $\epsilon_2 \leq \phi/2$, $\delta_2, \epsilon_1, \delta_1, M_2, r_2, \alpha$ in Algorithm 2 have the same order as the requirements introduced in Theorem 4.2. Then, it holds with probability at least $1 - \delta_2$ that the sequence $\{K_k\}$, $k \in \{0, \dots, K\}$, converges with $\mathcal{O}(1/K)$ rate in the sense that $K^{-1} \sum_{k=0}^{K-1} \|F_{K_k,L(K_k)}\|_F^2 \leq \epsilon_2$ with $K = \text{Tr}(P_{K_0,L(K_0)} - P_{K^*,L(K^*)})/[\alpha \phi \epsilon_2]$.*

Hence, ZO-NPG converges to the ϵ_2 -neighborhood of the Nash equilibrium with probability at least $1 - \delta_2$, after using $M_2 \cdot K \sim \tilde{\mathcal{O}}(\epsilon_2^{-5} \cdot \log(\delta_2^{-1}))$ samples. Moreover, the convergent gain matrix K_K also solves the LEQG and the LQ disturbance attenuation problems, due to Lemma 2.2.

Acknowledgements

K. Zhang, X. Zhang, and T. Başar were supported in part by the US Army Research Laboratory (ARL) Cooperative Agreement W911NF-17-2-0196, and in part by the Office of Naval Research (ONR) MURI Grant N00014-16-1-2710. B. Hu was supported by the National Science Foundation (NSF) award CAREER-2048168. The authors would like to thank Dhruv Malik and Na Li for helpful discussions.

References

- Başar, T. Minimax disturbance attenuation in LTV plants in discrete time. In *American Control Conference*, pp. 3112–3113, 1990.
- Başar, T. A dynamic games approach to controller design: Disturbance rejection in discrete-time. *IEEE Transactions on Automatic Control*, 36(8):936–952, 1991.
- Başar, T. and Bernhard, P. *H-Infinity Optimal Control and Related Minimax Design Problems: A Dynamic Game Approach*. Birkhäuser, Boston, 1995.
- Başar, T. and Olsder, G. J. *Dynamic Noncooperative Game Theory*. SIAM, 1998.
- Berg, C. *Kompleks Funktionsteori*. Universitetsbogladan, 2001.
- Bernhard, P. and Rapaport, A. On a theorem of Danskin with an application to a theorem of Von Neumann-Sion. *Nonlinear Analysis: Theory, Methods & Applications*, 24(8):1163–1181, 1995.
- Bhatnagar, S., Sutton, R. S., Ghavamzadeh, M., and Lee, M. Natural actor-critic algorithms. *Automatica*, 45(11):2471–2482, 2009.
- Boyd, S. and Vandenberghe, L. *Convex Optimization*. Cambridge University Press, 2004.
- Bu, J. and Mesbahi, M. Global convergence of policy gradient algorithms for indefinite least squares stationary optimal control. *IEEE Control Systems Letters*, 4(3):638–643, 2020.
- Bu, J., Mesbahi, A., Fazel, M., and Mesbahi, M. LQR through the lens of first order methods: Discrete-time case. *arXiv preprint arXiv:1907.08921*, 2019a.
- Bu, J., Ratliff, L. J., and Mesbahi, M. Global convergence of policy gradient for sequential zero-sum linear quadratic dynamic games. *arXiv preprint arXiv:1911.04672*, 2019b.
- Daskalakis, C. and Panageas, I. The limit points of (optimistic) gradient descent in min-max optimization. In *Advances in Neural Information Processing Systems*, pp. 9236–9246, 2018.
- Daskalakis, C., Foster, D. J., and Golowich, N. Independent policy gradient methods for competitive reinforcement learning. In *Advances in Neural Information Processing Systems*, 2020a.
- Daskalakis, C., Skoulakis, S., and Zampetakis, M. The complexity of constrained min-max optimization. *arXiv preprint arXiv:2009.09623*, 2020b.
- Diakonikolas, J., Daskalakis, C., and Jordan, M. I. Efficient methods for structured nonconvex-nonconcave min-max optimization. *arXiv preprint arXiv:2011.00364*, 2020.
- Doyle, J., Glover, K., Khargonekar, P., and Francis, B. State-space solutions to standard \mathcal{H}_2 and \mathcal{H}_∞ control problems. *IEEE Transactions on Automatic Control*, 34:831–847, 1989.
- Fatkhullin, I. and Polyak, B. Optimizing static linear feedback: Gradient method. *arXiv preprint arXiv:2004.09875*, 2020.
- Fazel, M., Ge, R., Kakade, S. M., and Mesbahi, M. Global convergence of policy gradient methods for the linear quadratic regulator. In *International Conference on Machine Learning*, pp. 1467–1476, 2018.
- Fei, Y., Yang, Z., Chen, Y., Wang, Z., and Xie, Q. Risk-sensitive reinforcement learning: Near-optimal risk-sample tradeoff in regret. *arXiv preprint arXiv:2006.13827*, 2020.
- Fiez, T. and Ratliff, L. Gradient descent-ascent provably converges to strict local minmax equilibria with a finite timescale separation. *arXiv preprint arXiv:2009.14820*, 2020.
- Flaxman, A. D., Kalai, A. T., and McMahan, H. B. Online convex optimization in the bandit setting: gradient descent without a gradient. In *ACM-SIAM Symposium on Discrete Algorithms*, pp. 385–394, 2005.
- Furieri, L., Zheng, Y., and Kamgarpour, M. Learning the globally optimal distributed LQ regulator. In *Learning for Dynamics and Control*, pp. 287–297, 2020.
- Glover, K. and Doyle, J. C. State-space formulae for all stabilizing controllers that satisfy an \mathcal{H}_∞ -norm bound and relations to risk sensitivity. *Systems & Control Letters*, 11(3):167–172, 1988.
- Gravell, B., Esfahani, P. M., and Summers, T. Learning robust controllers for linear quadratic systems with multiplicative noise via policy gradient. *arXiv preprint arXiv:1907.03680*, 2019.
- Hambly, B. M., Xu, R., and Yang, H. Policy gradient methods for the noisy linear quadratic regulator over a finite horizon. *arXiv preprint arXiv:2011.10300*, 2020.
- Hsieh, Y.-P., Liu, C., and Cevher, V. Finding mixed Nash equilibria of generative adversarial networks. In *International Conference on Machine Learning*, pp. 2810–2819, 2019.
- Iyengar, G. N. Robust dynamic programming. *Mathematics of Operations Research*, 30(2):257–280, 2005.
- Jacobson, D. Optimal stochastic linear systems with exponential performance criteria and their relation to deterministic differential games. *IEEE Transactions on Automatic Control*, 18(2):124–131, 1973.
- Jansch-Porto, J. P., Hu, B., and Dullerud, G. Convergence guarantees of policy optimization methods for Markovian jump linear systems. In *American Control Conference*, pp. 2882–2887, 2020a.
- Jansch-Porto, J. P., Hu, B., and Dullerud, G. Policy learning of MDPs with mixed continuous/discrete variables: A case study on model-free control of Markovian jump systems. In *Learning for Dynamics and Control*, pp. 947–957, 2020b.
- Jin, C., Netrapalli, P., and Jordan, M. I. What is local optimality in nonconvex-nonconcave minimax optimization? In *International Conference on Machine Learning*, 2020.
- Kakade, S. M. A natural policy gradient. In *Advances in Neural Information Processing Systems*, pp. 1531–1538, 2002.
- Kaminer, I., Khargonekar, P. P., and Rotea, M. A. Mixed $\mathcal{H}_2/\mathcal{H}_\infty$ control for discrete-time systems via convex optimization. *Automatica*, 29(1):57–70, 1993.
- Konda, V. R. and Tsitsiklis, J. N. Actor-critic algorithms. In *Advances in Neural Information Processing Systems*, pp. 1008–1014, 2000.

- Li, Y., Tang, Y., Zhang, R., and Li, N. Distributed reinforcement learning for decentralized linear quadratic control: A derivative-free policy optimization approach. *arXiv preprint arXiv:1912.09135*, 2019.
- Lillicrap, T. P., Hunt, J. J., Pritzel, A., Heess, N., Erez, T., Tassa, Y., Silver, D., and Wierstra, D. Continuous control with deep reinforcement learning. *arXiv preprint arXiv:1509.02971*, 2015.
- Lim, S. H. and Autef, A. Kernel-based reinforcement learning in robust markov decision processes. In *International Conference on Machine Learning*, pp. 3973–3981, 2019.
- Lim, S. H., Xu, H., and Mannor, S. Reinforcement learning in robust Markov decision processes. In *Advances in Neural Information Processing Systems*, pp. 701–709, 2013.
- Lin, T., Jin, C., and Jordan, M. I. On gradient descent ascent for nonconvex-concave minimax problems. In *International Conference on Machine Learning*, 2020.
- Littman, M. L. Markov games as a framework for multi-agent reinforcement learning. In *Machine Learning Proceedings*, pp. 157–163. Elsevier, 1994.
- Malik, D., Pananjady, A., Bhatia, K., Khamaru, K., Bartlett, P. L., and Wainwright, M. J. Derivative-free methods for policy optimization: Guarantees for linear quadratic systems. *Journal of Machine Learning Research*, 21(21):1–51, 2020.
- Mankowitz, D. J., Levine, N., Jeong, R., Abdolmaleki, A., Springenberg, J. T., Mann, T., Hester, T., and Riedmiller, M. Robust reinforcement learning for continuous control with model misspecification. *arXiv preprint arXiv:1906.07516*, 2019.
- Mårtensson, K. and Rantzer, A. Gradient methods for iterative distributed control synthesis. In *IEEE Conference on Decision and Control held jointly with Chinese Control Conference*, pp. 549–554, 2009.
- Mohammadi, H., Zare, A., Soltanolkotabi, M., and Jovanović, M. R. Convergence and sample complexity of gradient methods for the model-free linear quadratic regulator problem. *arXiv preprint arXiv:1912.11899*, 2019.
- Mohammadi, H., Soltanolkotabi, M., and Jovanović, M. R. On the linear convergence of random search for discrete-time LQR. *IEEE Control Systems Letters*, 5(3):989–994, 2020.
- Morimoto, J. and Doya, K. Robust reinforcement learning. *Neural Computation*, 17(2):335–359, 2005.
- Mustafa, D. and Bernstein, D. S. LQG cost bounds in discrete-time $\mathcal{H}_2/\mathcal{H}_\infty$ control. *Transactions of the Institute of Measurement and Control*, 13(5):269–275, 1991.
- Nilim, A. and El Ghaoui, L. Robust control of Markov decision processes with uncertain transition matrices. *Operations Research*, 53(5):780–798, 2005.
- Nishimura, H., Mehr, N., Gaidon, A., and Schwager, M. RAT ILQR: A risk auto-tuning controller to optimally account for stochastic model mismatch. *IEEE Robotics and Automation Letters*, 6(2):763–770, 2021.
- Nouiehed, M., Sanjabi, M., Huang, T., Lee, J. D., and Razaviyayn, M. Solving a class of non-convex min-max games using iterative first order methods. In *Advances in Neural Information Processing Systems*, pp. 14934–14942, 2019.
- Pinto, L., Davidson, J., Sukthankar, R., and Gupta, A. Robust adversarial reinforcement learning. In *International Conference on Machine Learning*, pp. 2817–2826, 2017.
- Rautert, T. and Sachs, E. W. Computational design of optimal output feedback controllers. *SIAM Journal on Optimization*, 7(3):837–852, 1997.
- Recht, B. A tour of reinforcement learning: The view from continuous control. *Annual Review of Control, Robotics, and Autonomous Systems*, 2:253–279, 2019.
- Roulet, V., Fazel, M., Srinivasa, S., and Harchaoui, Z. On the convergence of the iterative linear exponential quadratic Gaussian algorithm to stationary points. In *American Control Conference*, pp. 132–137. IEEE, 2020.
- Schulman, J., Levine, S., Abbeel, P., Jordan, M., and Moritz, P. Trust region policy optimization. In *International Conference on Machine Learning*, pp. 1889–1897, 2015a.
- Schulman, J., Moritz, P., Levine, S., Jordan, M., and Abbeel, P. High-dimensional continuous control using generalized advantage estimation. *arXiv preprint arXiv:1506.02438*, 2015b.
- Schulman, J., Wolski, F., Dhariwal, P., Radford, A., and Klimov, O. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- Shapley, L. S. Stochastic games. *Proceedings of the National Academy of Sciences*, 39(10):1095–1100, 1953.
- Sutton, R. S., McAllester, D. A., Singh, S. P., and Mansour, Y. Policy gradient methods for reinforcement learning with function approximation. In *Advances in Neural Information Processing Systems*, pp. 1057–1063, 2000.
- Tu, S. and Recht, B. The gap between model-based and model-free methods on the linear quadratic regulator: An asymptotic viewpoint. In *Conference on Learning Theory*, pp. 3036–3083, 2019.
- Vlatakis-Gkaragkounis, E.-V., Flokas, L., and Piliouras, G. Poincaré recurrence, cycles and spurious equilibria in gradient-descent-ascent for non-convex non-concave zero-sum games. In *Advances in Neural Information Processing Systems*, pp. 10450–10461, 2019.
- Whittle, P. Risk-sensitive linear/quadratic/gaussian control. *Advances in Applied Probability*, pp. 764–777, 1981.
- Whittle, P. *Risk-Sensitive Optimal Control*, volume 2. Wiley, 1990.
- Yang, J., Kiyavash, N., and He, N. Global convergence and variance-reduced optimization for a class of nonconvex-nonconcave minimax problems. In *Advances in Neural Information Processing Systems*, 2020.
- Zames, G. Feedback and optimal sensitivity: Model reference transformations, multiplicative seminorms, and approximate inverses. *IEEE Transactions on Automatic Control*, 26(2):301–320, 1981.
- Zhang, K., Yang, Z., and Başar, T. Policy optimization provably converges to Nash equilibria in zero-sum linear quadratic games. In *Advances in Neural Information Processing Systems*, pp. 11602–11614, 2019.

Zhang, K., Hu, B., and Başar, T. On the stability and convergence of robust adversarial reinforcement learning: A case study on linear quadratic systems. In *Advances in Neural Information Processing Systems*, 2020.

Zhang, K., Hu, B., and Başar, T. Policy optimization for \mathcal{H}_2 linear control with \mathcal{H}_∞ robustness guarantee: Implicit regularization and global convergence. *SIAM Journal on Control and Optimization*, 2021a.

Zhang, X., Chen, Y., Zhu, X., and Sun, W. Robust policy gradient against strong data corruption. *arXiv preprint arXiv:2102.05800*, 2021b.

Supplementary Materials for “Derivative-Free Policy Optimization for Linear Risk-Sensitive and Robust Control Design: Implicit Regularization and Sample Complexity”

A. Appendix

The code for our simulations in §D are provided in the supplementary materials.

A.1. Related Works

Risk-Sensitive/Robust Control and Dynamic Games. \mathcal{H}_∞ -robust control has been one of the most fundamental research fields in control theory, addressing worst-case controller design for linear plants in the presence of unknown disturbances and uncertainties. Frequency-domain and time-domain/state-space formulations of the \mathcal{H}_∞ -robust control problem were first introduced in (Zames, 1981) and (Doyle et al., 1989), respectively. Based on the time-domain approach, the precise equivalence relationships between controllers in the LQ disturbance attenuation (which belongs to a class of \mathcal{H}_∞ -robust control) problem, risk-sensitive linear control, and zero-sum LQ dynamic games have been studied extensively (Jacobson, 1973; Glover & Doyle, 1988; Başar, 1990; 1991; Başar & Bernhard, 1995; Başar & Olsder, 1998). Specifically, (Jacobson, 1973) first demonstrated the equivalence of the controllers between LEQG and zero-sum LQ (differential) games; (Glover & Doyle, 1988) studied the relationship between the mixed $\mathcal{H}_2/\mathcal{H}_\infty$ design problem, a sub-problem of the \mathcal{H}_∞ -robust control problem, and LEQG; (Başar, 1991; 1990) investigated the LQ disturbance attenuation problem through a dynamic games approach, for both finite- and infinite-horizons, time-invariant and -varying, deterministic and stochastic settings. We refer the readers to (Başar & Bernhard, 1995; Başar & Olsder, 1998) for comprehensive studies of \mathcal{H}_∞ -robust control, dynamic game theory, risk-sensitive control, and also their precise interconnections.

Policy Optimization for LQ Control. Applying PG methods to LQ control problems has been investigated in both control and learning communities extensively. In contrast to the early work on this topic (Rautert & Sachs, 1997; Mårtensson & Rantzer, 2009), the recent study focuses more on theoretical aspects such as global convergence and sample complexity of these PG methods (Fazel et al., 2018; Tu & Recht, 2019; Bu et al., 2019a; Zhang et al., 2019; Bu et al., 2019b; Mohammadi et al., 2019; Zhang et al., 2021a; Gravell et al., 2019; Malik et al., 2020; Bu & Mesbahi, 2020; Mohammadi et al., 2020). Specifically, (Fazel et al., 2018) was the first work to show the global convergence of PG methods for the LQR problem. Initial sample complexity results for the derivative-free PG methods based on zeroth-order optimization techniques were also reported in (Fazel et al., 2018). Subsequently, (Bu et al., 2019a) characterized the optimization landscape of the LQR in detail, and provided an initial extension toward the distributive LQR. Building upon these works, (Malik et al., 2020) enhanced the sample complexity result by adopting a two-point zeroth-order optimization method, in contrast to the single-point method adopted in (Fazel et al., 2018). The sample complexity result of the two-point method was further improved recently in both continuous- (Mohammadi et al., 2019) and discrete-time LQR (Mohammadi et al., 2020). Moreover, (Bu & Mesbahi, 2020) studied the global convergence of PG methods for the LQR where both the state and control weighting matrices are indefinite, but without addressing the derivative-free setting; (Zhang et al., 2019; Bu et al., 2019b) derived global convergence of PG methods for zero-sum LQ dynamic games, also pointing to the need to investigate the indefinite LQR. The most relevant works that addressed LQ control with *robustness/risk-sensitivity concerns* using policy optimization methods were (Gravell et al., 2019; Roulet et al., 2020; Zhang et al., 2021a). The work in (Gravell et al., 2019) considered LQR with multiplicative noises, which enjoys a similar landscape as standard LQR. The work in (Roulet et al., 2020) studied finite-horizon risk-sensitive nonlinear control, and established stationary-point convergence when using iterative LEQG and performing optimization over control actions directly. (Zhang et al., 2021a) examined the mixed $\mathcal{H}_2/\mathcal{H}_\infty$ control design problem, a class of \mathcal{H}_∞ -robust control problems, with a different landscape from LQR. However, no sample complexity results were provided in either (Roulet et al., 2020) or (Zhang et al., 2021a). Very recently, (Hambly et al., 2020) has also studied the convergence of policy gradient method for finite-horizon LQR problems. Interestingly, the landscape analysis independently developed in (Hambly et al., 2020) matches some of our findings for the inner-loop subproblem of zero-sum LQ dynamic games, while we consider a more general setting with time-varying system dynamics and a possibly indefinite state-weighting matrix. Other recent results on PO for LQ control include (Li et al., 2019; Furieri et al., 2020; Tu & Recht, 2019; Jansch-Porto et al., 2020a;b; Fatkhullin & Polyak, 2020).

Nonconvex-Nonconcave Minimax Optimization & MARL. Solving minimax problems using PG methods, especially

gradient descent-ascent (GDA), has been one of the most popular sub-fields in the optimization community recently. Convergence of the first-order GDA algorithm (and its variants) has been studied in (Daskalakis & Panageas, 2018; Hsieh et al., 2019; Jin et al., 2020; Vlatakis-Gkaragkounis et al., 2019) for general nonconvex-nonconcave objectives, in (Yang et al., 2020) for the setting with an additional two-sided Polyak-Łojasiewicz (PL) condition, and in (Diakonikolas et al., 2020) for the setting under the weak Minty variational inequality condition. In fact, (Daskalakis et al., 2020b) has shown very recently the computational hardness of solving nonconvex-nonconcave constrained minimax problems with first-order methods. To the best of our knowledge, *zeroth-order/derivative-free* methods, as studied in our paper, have not yet been investigated for nonconvex-nonconcave minimax optimization problems. On the other hand, in the multi-agent RL (MARL) regime, policy optimization for solving zero-sum Markov/dynamic games naturally leads to a nonconvex-nonconcave minimax problem (Zhang et al., 2019; Daskalakis et al., 2020a), whose convergence guarantee has remained open until very recently, except for the aforementioned LQ setting (Zhang et al., 2019; Bu et al., 2019b). In particular, (Daskalakis et al., 2020a) has established the first non-asymptotic global convergence of independent policy gradient methods for tabular (finite-state-action) zero-sum Markov games, where GDA with two-timescale stepsizes is used. In stark contrast to these works, the crucial global *smoothness* assumption (or property that automatically holds in the tabular MARL setting) on the objective therein does not hold in our control setting with *unbounded* and *continuous* spaces. A careful characterization of the iterate trajectory is thus required in order to establish global convergence and sample complexity results.

Risk-Sensitive/Robust RL. Risk-sensitivity/robustness to model uncertainty/misspecification in RL has attracted significant research efforts, following the frameworks of robust adversarial RL (RARL) and robust Markov decision process (RMDP). Tracing back to (Morimoto & Doya, 2005), early works on RARL exploited the same game-theoretic perspective as in \mathcal{H}_∞ -robust control, by modeling the uncertainty as a fictitious adversary against the nominal agent. This worst-case (minimax) design concept then enabled actor-critic type algorithms with exceptional empirical performance, but rather sparse theoretical footprints (Pinto et al., 2017). Very recently, theoretical investigations on the convergence and stability of RARL have been carried out in (Zhang et al., 2020) within the LQ setting. In parallel, the RMDP framework has been proposed in (Nilim & El Ghaoui, 2005; Iyengar, 2005), and the robustness of RL algorithms under such framework has been further studied in (Lim et al., 2013; Lim & Autef, 2019; Mankowitz et al., 2019) for the *finite* MDP settings, in contrast to our continuous control tasks. Other recent advances on risk-sensitive/robust RL/MDP include (Fei et al., 2020; Zhang et al., 2021b).

A.2. Notations

For a square matrix X of proper dimension, we use $\text{Tr}(X)$ to denote its trace. We also use $\|X\|$ and $\|X\|_F$ to denote, respectively, the operator norm and the Frobenius norm of X . If X is further symmetric, we use $X > 0$ to denote that X is positive definite. Similarly, $X \geq 0$, $X \leq 0$, and $X < 0$ are used to denote X being positive semi-definite, negative semi-definite, and negative definite, respectively. Further, again for a symmetric matrix X , $\lambda_{\min}(X)$ and $\lambda_{\max}(X)$ are used to denote, respectively, the smallest and the largest eigenvalues of X . Moreover, we use $\langle X, Y \rangle := \text{Tr}(X^\top Y)$ to denote the standard matrix inner product and use $\text{diag}(X_0, \dots, X_N)$ to denote the block-diagonal matrix with X_0, \dots, X_N on the diagonal block entries. In the case where $X_0 = X_1 = \dots = X_N = X$, we further denote it as $\text{diag}(X^N)$. We use $x \sim \mathcal{N}(\mu, \Sigma)$ to denote a Gaussian random variable with mean μ and covariance Σ and use $\|v\|$ to denote the Euclidean norm of a vector v . Lastly, we use \mathbf{I} and $\mathbf{0}$ to denote the identity and zero matrices with appropriate dimensions.

A.3. Extended Backgrounds

In this section, we first introduce two classic settings in risk-sensitive and robust control, namely LEQG, and LQ disturbance attenuation. We then discuss, in §2.2, the challenges one confronts when attempting to address the above two problems directly using derivative-free PG methods by sampling system trajectories. Fortunately, solving zero-sum LQ (stochastic) dynamic games, a benchmark setting in MARL, via derivative-free PG methods by sampling system trajectories provides a workaround to address these problems all in a unified way, due to the well-known equivalence relationships between zero-sum LQ dynamic games and the two aforementioned classes of problems (Başar & Bernhard, 1995), which we will also discuss in §A.3.3.

A.3.1. LINEAR EXPONENTIAL QUADRATIC GAUSSIAN

We first consider a fundamental setting of risk-sensitive optimal control, known as the LEQG problem (Jacobson, 1973; Whittle, 1981; 1990), in the finite-horizon setting. The time-varying (linear) systems dynamics are described by:

$$x_{t+1} = A_t x_t + B_t u_t + w_t, \quad t \in \{0, \dots, N-1\},$$

where $x_t \in \mathbb{R}^m$ represents the system state; $u_t \in \mathbb{R}^d$ is the control input; $w_t \in \mathbb{R}^m$ is an independent (across time) Gaussian random noise drawn from $w_t \sim \mathcal{N}(\mathbf{0}, W)$ for some $W > 0$; the initial state $x_0 \sim \mathcal{N}(\mathbf{0}, X_0)$ is a Gaussian random vector for some $X_0 > 0$, independent of the sequence $\{w_t\}$; and A_t, B_t are time-varying system matrices with appropriate dimensions. The objective function is given by

$$\mathcal{J}(\{u_t\}) := \frac{2}{\beta} \log \mathbb{E} \exp \left[\frac{\beta}{2} \left(\sum_{t=0}^{N-1} (x_t^\top Q_t x_t + u_t^\top R_t u_t) + x_N^\top Q_N x_N \right) \right],$$

where $Q_t \geq 0$ and $R_t > 0$, for all $t \in \{0, \dots, N-1\}$, and $Q_N \geq 0$ are symmetric weighting matrices; and $\beta > 0$ is a parameter capturing the degree of risk-sensitivity, which is upper-bounded by some $\beta^* > 0$ (Glover & Doyle, 1988; Whittle, 1990; Nishimura et al., 2021).

The goal in the LEQG problem is to find the \mathcal{J} -minimizing optimal control policy $\mu_t^* : (\mathbb{R}^m \times \mathbb{R}^d)^t \times \mathbb{R}^m \rightarrow \mathbb{R}^d$ that maps, at time t , the history of state-control pairs up to time t and the current state x_t to the control u_t , this being so for all $t \in \{0, \dots, N-1\}$. It has been shown in (Jacobson, 1973) that μ_t^* has a linear state-feedback form $\mu_t^*(x_t) = -K_t^* x_t$, where $K_t^* \in \mathbb{R}^{d \times m}$, for all $t \in \{0, \dots, N-1\}$. Therefore, it suffices to search K_t^* in the matrix space $\mathbb{R}^{d \times m}$ for all $t \in \{0, \dots, N-1\}$, without losing any optimality. The resulting policy optimization problem is then represented as (by a slight abuse of notation with regard to \mathcal{J}):

$$\min_{\{K_t\}} \mathcal{J}(\{K_t\}) := \frac{2}{\beta} \log \mathbb{E} \exp \left[\frac{\beta}{2} \left(\sum_{t=0}^{N-1} (x_t^\top (Q_t + K_t^\top R_t K_t) x_t) + x_N^\top Q_N x_N \right) \right]. \quad (\text{A.1})$$

To characterize the solution to (A.1), we first introduce the following time-varying Riccati difference equation (RDE):

$$P_{K_t} = Q_t + K_t^\top R_t K_t + (A_t - B_t K_t)^\top \tilde{P}_{K_{t+1}} (A_t - B_t K_t), \quad t \in \{0, \dots, N-1\}, \quad P_{K_N} = Q_N, \quad (\text{A.2})$$

where $\tilde{P}_{K_{t+1}} := P_{K_{t+1}} + \beta P_{K_{t+1}} (W^{-1} - \beta P_{K_{t+1}})^{-1} P_{K_{t+1}}$. Then, (A.1) can be expressed by the solution to (A.2) and the exact expression of its PG can be analyzed, as follows, with its proof being deferred to §B.1.

Lemma A.1 (Closed-Form Objective Function and PG) *For any sequence of $\{K_t\}$ such that (A.2) generates a sequence of positive semi-definite (p.s.d.) sequence $\{P_{K_t}\}$ satisfying $X_0^{-1} - \beta P_{K_0} > 0$ and $W^{-1} - \beta P_{K_t} > 0$, for all $t \in \{1, \dots, N\}$, the objective function $\mathcal{J}(\{K_t\})$ can be expressed as*

$$\mathcal{J}(\{K_t\}) = -\frac{1}{\beta} \log \det(\mathbf{I} - \beta P_{K_0} X_0) - \frac{1}{\beta} \sum_{t=1}^N \log \det(\mathbf{I} - \beta P_{K_t} W). \quad (\text{A.3})$$

As $\beta \rightarrow 0$, (A.3) reduces to $\text{Tr}(P_{K_0} X_0) + \sum_{t=1}^N \text{Tr}(P_{K_t} W)$, which is the objective function of the finite-horizon linear-quadratic-Gaussian (LQG) problem. Moreover, the PG of (A.3) at time t has the form:

$$\nabla_{K_t} \mathcal{J}(\{K_t\}) := \frac{\partial \mathcal{J}(\{K_t\})}{\partial K_t} = 2[(R_t + B_t^\top \tilde{P}_{K_{t+1}} B_t) K_t - B_t^\top \tilde{P}_{K_{t+1}} A_t] \Sigma_{K_t}, \quad t \in \{0, \dots, N-1\}, \quad (\text{A.4})$$

where

$$\begin{aligned} \Sigma_{K_t} := & \prod_{i=0}^{t-1} [(\mathbf{I} - \beta P_{K_{i+1}} W)^{-\top} (A_i - B_i K_i)] \cdot X_0^{\frac{1}{2}} (\mathbf{I} - \beta X_0^{\frac{1}{2}} P_{K_0} X_0^{\frac{1}{2}})^{-1} X_0^{\frac{1}{2}} \\ & \cdot \prod_{i=0}^{t-1} [(A_i - B_i K_i)^{\top} (\mathbf{I} - \beta P_{K_{i+1}} W)^{-1}] \\ & + \sum_{\tau=1}^t \left\{ \prod_{i=\tau}^{t-1} [(\mathbf{I} - \beta P_{K_{i+1}} W)^{-\top} (A_i - B_i K_i)] \cdot W^{\frac{1}{2}} (\mathbf{I} - \beta W^{\frac{1}{2}} P_{K_{\tau}} W^{\frac{1}{2}})^{-1} W^{\frac{1}{2}} \right. \\ & \left. \cdot \prod_{i=\tau}^{t-1} [(A_i - B_i K_i)^{\top} (\mathbf{I} - \beta P_{K_{i+1}} W)^{-1}] \right\}. \end{aligned}$$

A.3.2. LQ DISTURBANCE ATTENUATION

Second, we introduce the *optimal* LQ disturbance attenuation problem, again for finite-horizon settings, with time-varying (deterministic) dynamical systems described by

$$x_{t+1} = A_t x_t + B_t u_t + D_t w_t, \quad z_t = C_t x_t + E_t u_t, \quad t \in \{0, \dots, N-1\},$$

where $x_t \in \mathbb{R}^m$ is the system state; $u_t \in \mathbb{R}^d$ is the control input; $w_t \in \mathbb{R}^n$ is the (unknown) disturbance input; $z_t \in \mathbb{R}^l$ is the controlled output; A_t, B_t, C_t, D_t, E_t are system matrices with appropriate dimensions; and $x_0 \in \mathbb{R}^m$ is unknown. In addition, we assume that $E_t^{\top} [C_t \ E_t] = [0 \ R_t]$ for some $R_t > 0$ to eliminate the cross-weightings between control input u and system state x , which provides no loss of generality, and is standard in the literature (Glover & Doyle, 1988; Başar & Bernhard, 1995). As shown in §3.5.1 of (Başar & Bernhard, 1995), a simple procedure can transform the general problem into a form where there are no coupling terms between u and x , satisfying the above “normalization” assumption that we have made. Subsequently, we introduce the ℓ^2 -norms of the vectors $\omega := [x_0^{\top} C_0^{\top}, w_0^{\top}, \dots, w_{N-1}^{\top}]^{\top}$ and $z := [z_0^{\top}, \dots, z_{N-1}^{\top}, x_N^{\top} Q_N^{1/2}]^{\top}$ as $\|\omega\| := \{x_0^{\top} C_0^{\top} C_0 x_0 + \sum_{t=0}^{N-1} \|w_t\|^2\}^{1/2}$ and $\|z\| = \{\mathcal{C}(\{u_t\}, \{w_t\})\}^{1/2} := \{x_N^{\top} Q_N x_N + \sum_{t=0}^{N-1} \|z_t\|^2\}^{1/2}$, where $Q_N \geq 0$ is symmetric, and $\|\cdot\|$ in the two expressions denote appropriate Euclidean norms. Then, the robustness of a designed controller can be guaranteed by a constraint on the ratio between $\|z\|$ and $\|\omega\|$. Specifically, the goal of the optimal LQ disturbance attenuation problem is to find the optimal control policy $\mu_t^* : (\mathbb{R}^m \times \mathbb{R}^d)^t \times \mathbb{R}^m \rightarrow \mathbb{R}^d$ that maps, at time t , the history of state-control pairs until time t and the current state x_t , denoted as $h_t := \{x_0, \dots, x_t, u_0, \dots, u_{t-1}\}$, to the control u_t , such that

$$\gamma^* = \inf_{\{\mu_t\}} \sup_{x_0, \{w_t\}} \frac{\{\mathcal{C}(\{\mu_t(h_t)\}, \{w_t\})\}^{1/2}}{\|\omega\|} = \sup_{x_0, \{w_t\}} \frac{\{\mathcal{C}(\{\mu_t^*(h_t)\}, \{w_t\})\}^{1/2}}{\|\omega\|}, \quad (\text{A.5})$$

where $\gamma^* = \sqrt{(\beta^*)^{-1}} > 0$ is the *optimal (minimax)* level of disturbance attenuation at the output (Glover & Doyle, 1988; Başar & Bernhard, 1995), and recall that β^* is the upper bound for the risk-sensitivity parameter in LEQG. It is known that under closed-loop perfect-state information pattern, the existence of a sequence $\{\mu_t^*\}$, $t \in \{0, \dots, N-1\}$, is always guaranteed and μ_t^* is linear state-feedback, i.e., $\mu_t^*(h_t) = -K_t^* x_t$ ((Başar & Bernhard, 1995); Theorem 3.5). Thus, it suffices to search for $K_t^* \in \mathbb{R}^{d \times m}$, for all $t \in \{0, \dots, N-1\}$, to achieve the optimal attenuation.

The problem (A.5) can be challenging to solve to the optimum (Başar & Bernhard, 1995). Moreover, in practice, due to the inevitable uncertainty in modeling the system, the robust controller that achieves the optimal attenuation level can be too sensitive to the model uncertainty to use. Hence, a reasonable surrogate of the optimal LQ disturbance attenuation problem is the following: Given a $\gamma > \gamma^*$, find a control policy $\mu_t = -K_t x_t$, for all $t \in \{0, \dots, N-1\}$, that solves

$$\min_{\{\mu_t\}} \mathcal{J}(\{\mu_t(x_t)\}) := \mathbb{E}_{\{\xi_t\}} \left[\sum_{t=0}^{N-1} (x_t^{\top} Q_t x_t + u_t^{\top} R_t u_t) + x_N^{\top} Q_N x_N \right] \quad \text{s.t.} \quad \sup_{x_0, \{w_t\}} \frac{\{\mathcal{C}(\{\mu_t(x_t)\}, \{w_t\})\}^{1/2}}{\|\omega\|} < \gamma, \quad (\text{A.6})$$

where $Q_t := C_t^{\top} C_t$ for $t \in \{0, \dots, N-1\}$, $\mathcal{J}(\{\mu_t(x_t)\})$ is the LQG cost of the system $x_{t+1} = A_t x_t + B_t u_t + D_t \xi_t$, and $\xi_t \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ is independent Gaussian noise with unit intensity (Mustafa & Bernstein, 1991). Note that the LQ disturbance attenuation problem in (A.6), and particularly with $x_0 = \mathbf{0}$, is also called the mixed $\mathcal{H}_2/\mathcal{H}_{\infty}$ problem (Doyle et al., 1989;

Mustafa & Bernstein, 1991; Kaminer et al., 1993; Zhang et al., 2021a), when the system is time-invariant, and the horizon of the problem N is ∞ . To solve (A.6) using policy optimization methods, we introduce the time-varying RDE:

$$P_{K_t} = Q_t + K_t^\top R_t K_t + (A_t - B_t K_t)^\top \tilde{P}_{K_{t+1}} (A_t - B_t K_t), \quad t \in \{0, \dots, N-1\}, \quad P_{K_N} = Q_N, \quad (\text{A.7})$$

where $\tilde{P}_{K_{t+1}} := P_{K_{t+1}} + P_{K_{t+1}}(\gamma^2 \mathbf{I} - D_t^\top P_{K_{t+1}} D_t)^{-1} P_{K_{t+1}}$. Then, two upper bounds of \mathcal{J} and their corresponding PGs can be expressed in terms of solutions to (A.7) (if exist) (Mustafa & Bernstein, 1991). In particular, (A.8) below is closely related to the objective function of LEQG in (A.1), and (A.9) is connected to the objective function of zero-sum LQ dynamic games to be introduced shortly.

Lemma A.2 (Closed-Forms for the Objective Function and PGs) *For any sequence of control gains $\{K_t\}$ such that (A.7) generates a sequence of p.s.d. solutions $\{P_{K_{t+1}}\}$ satisfying $\gamma^2 \mathbf{I} - D_t^\top P_{K_{t+1}} D_t > 0$, for all $t \in \{0, \dots, N-1\}$, and $\gamma^2 \mathbf{I} - P_{K_0} > 0$, then two common upper bounds of the objective function \mathcal{J} in (A.6) are:*

$$\overline{\mathcal{J}}(\{K_t\}) = -\gamma^2 \log \det(\mathbf{I} - \gamma^{-2} P_{K_0}) - \gamma^2 \sum_{t=1}^N \log \det(\mathbf{I} - \gamma^{-2} P_{K_t} D_{t-1} D_{t-1}^\top), \quad (\text{A.8})$$

$$\overline{\mathcal{J}}(\{K_t\}) = \text{Tr}(P_{K_0}) + \sum_{t=1}^N \text{Tr}(P_{K_t} D_{t-1} D_{t-1}^\top). \quad (\text{A.9})$$

The PGs of (A.8) and (A.9) at time t can be expressed as

$$\nabla_{K_t} \overline{\mathcal{J}}(\{K_t\}) := \frac{\partial \overline{\mathcal{J}}(\{K_t\})}{\partial K_t} = 2[(R_t + B_t^\top \tilde{P}_{K_{t+1}} B_t) K_t - B_t^\top \tilde{P}_{K_{t+1}} A_t] \Sigma_{K_t}, \quad t \in \{0, \dots, N-1\}, \quad (\text{A.10})$$

where Σ_t , $t \in \{0, \dots, N-1\}$, for the PGs of (A.8) and (A.9) are expressed, respectively, as

$$\begin{aligned} \Sigma_{K_t} &:= \prod_{i=0}^{t-1} [(\mathbf{I} - \gamma^{-2} P_{K_{i+1}} D_i D_i^\top)^{-\top} (A_i - B_i K_i)] \cdot (\mathbf{I} - \gamma^{-2} P_{K_0})^{-1} \cdot \prod_{i=0}^{t-1} [(A_i - B_i K_i)^\top (\mathbf{I} - \gamma^{-2} P_{K_{i+1}} D_i D_i^\top)^{-1}] \\ &\quad + \sum_{\tau=1}^t \left\{ \prod_{i=\tau}^{t-1} [(\mathbf{I} - \gamma^{-2} P_{K_{i+1}} D_i D_i^\top)^{-\top} (A_i - B_i K_i)] \cdot D_{\tau-1} (\mathbf{I} - \gamma^{-2} D_{\tau-1}^\top P_{K_\tau} D_{\tau-1})^{-1} D_{\tau-1}^\top \right. \\ &\quad \left. \cdot \prod_{i=\tau}^{t-1} [(A_i - B_i K_i)^\top (\mathbf{I} - \gamma^{-2} P_{K_{i+1}} D_i D_i^\top)^{-1}] \right\}, \\ \Sigma_{K_t} &:= \prod_{i=0}^{t-1} [(\mathbf{I} - \gamma^{-2} P_{K_{i+1}} D_i D_i^\top)^{-\top} (A_i - B_i K_i) (A_i - B_i K_i)^\top (\mathbf{I} - \gamma^{-2} P_{K_{i+1}} D_i D_i^\top)^{-1}] \\ &\quad + \sum_{\tau=1}^t \left\{ \prod_{i=\tau}^{t-1} [(\mathbf{I} - \gamma^{-2} P_{K_{i+1}} D_i D_i^\top)^{-\top} (A_i - B_i K_i)] \cdot D_{\tau-1} D_{\tau-1}^\top \cdot \prod_{i=\tau}^{t-1} [(A_i - B_i K_i)^\top (\mathbf{I} - \gamma^{-2} P_{K_{i+1}} D_i D_i^\top)^{-1}] \right\}. \end{aligned}$$

The proof of Lemma A.2 is deferred to §B.2. Before moving on to the next section, we provide a remark regarding the challenges of addressing LEQG or the LQ disturbance attenuation problem directly using derivative-free PG methods.

Remark A.3 *It seems tempting to tackle LEQG or the LQ disturbance attenuation problem directly using derivative-free PG methods, e.g., vanilla PG or natural PG methods by sampling the system trajectories as in (Fazel et al., 2018). However, there are several challenges. First, when attempting to sample the LEQG objective (A.1) using system trajectories (and then to estimate PGs), the bias of gradient estimates can be difficult to control uniformly with a fixed sample size per iterate, as the log function is not globally Lipschitz¹. Moreover, for the vanilla PG method, even with exact PG accesses, the iterates may not preserve a certain level of risk-sensitivity/disturbance attenuation, i.e., remain feasible, along the iterations, which can make the disturbance input drive the cost to arbitrarily large values (see the numerical examples in (Zhang et al., 2021a))*

¹One workaround for mitigating this bias might be to remove the log operator in the objective (A.1), and only minimize the terms after the $\mathbb{E} \exp(\cdot)$ function. However, it is not clear yet if derivative-free methods provably converge for this objective, as exp function is neither globally smooth nor Lipschitz. We have left this direction in our future work, as our goal in the present work is to provide a unified way to solve all three classes of problems.

for the infinite-horizon LTI setting). This failure will only be exacerbated in the derivative-free setting with accesses to only noisy PG estimates. Lastly, developing a derivative-free natural PG method (which we will show in §3.3 that it can preserve a prescribed attenuation level along the iterations) is also challenging, because the expressions of Σ_{K_t} in Lemmas A.1 and A.2 cannot be sampled from system trajectories directly. Therefore, the game-theoretic approach to be introduced next is rather one workaround.

A.3.3. AN EQUIVALENT DYNAMIC GAME FORMULATION

Lastly, and more importantly, we describe an equivalent dynamic game formulation to the LEQG and the LQ disturbance attenuation problems introduced above. Under certain conditions to be introduced in Lemma A.6, the saddle-point gain matrix of the minimizing player in the dynamic game (if exists) also addresses the LEQG and the LQ disturbance attenuation problem, providing an alternative route to overcome the challenges reported in Remark A.3. Specifically, we consider a zero-sum LQ stochastic dynamic game (henceforth, game) model with closed-loop perfect-state information pattern, which can also be viewed as a benchmark setting of MARL for two competing agents (Shapley, 1953; Littman, 1994; Zhang et al., 2019; Bu et al., 2019b), as the role played by LQR for single-agent RL (Recht, 2019). The linear time-varying system dynamics follow

$$x_{t+1} = A_t x_t + B_t u_t + D_t w_t + \xi_t, \quad t \in \{0, \dots, N-1\}, \quad (\text{A.11})$$

where $x_t \in \mathbb{R}^m$ is the system state, $u_t := \mu_{u,t}(h_{u,t}) \in \mathbb{R}^d$ (resp., $w_t := \mu_{w,t}(h_{w,t}) \in \mathbb{R}^n$) is the control input of the minimizing (resp., maximizing) player² and $\mu_{u,t} : (\mathbb{R}^m \times \mathbb{R}^d)^t \times \mathbb{R}^m \rightarrow \mathbb{R}^d$ (resp., $\mu_{w,t} : (\mathbb{R}^m \times \mathbb{R}^n)^t \times \mathbb{R}^m \rightarrow \mathbb{R}^n$) is the control policy that maps, at time t , the history of state-control pairs up to time t , and state at time t , denoted as $h_{u,t} := \{x_0, \dots, x_t, u_0, \dots, u_{t-1}\}$ (resp., $h_{w,t} := \{x_0, \dots, x_t, w_0, \dots, w_{t-1}\}$), to the control u_t (resp., w_t). The independent process noises are denoted by $\xi_t \sim \mathcal{D}$ and A_t, B_t, D_t are system matrices with proper dimensions. Further, we assume $x_0 \sim \mathcal{D}$ to be independent of the process noises, and the distribution \mathcal{D} has zero mean and a positive-definite covariance. We also assume that $\|x_0\|, \|\xi_t\| \leq \vartheta$, for all $t \in \{0, \dots, N-1\}$, almost surely, for some constant ϑ ³. The goal of the minimizing (resp. maximizing) player is to minimize (resp. maximize) a quadratic objective function, namely to solve the zero-sum game

$$\inf_{\{u_t\}} \sup_{\{w_t\}} \mathbb{E}_{x_0, \xi_0, \dots, \xi_{N-1}} \left[\sum_{t=0}^{N-1} (x_t^\top Q_t x_t + u_t^\top R_t^u u_t - w_t^\top R_t^w w_t) + x_N^\top Q_N x_N \right], \quad (\text{A.12})$$

where $Q_t \geq 0, R_t^u, R_t^w > 0$ are symmetric weighting matrices and the system transitions follow (A.11). Whenever the solution to (A.12) exists such that the inf and sup operators in (A.12) are interchangeable, then the value (A.12) is the *value* of the game, and the corresponding policies of the players are known as saddle-point policies. To characterize the solution to (A.12), we first introduce the following time-varying RDE:

$$P_t^* = Q_t + A_t^\top P_{t+1}^* \Lambda_t^{-1} A_t, \quad t \in \{0, \dots, N-1\}, \quad (\text{A.13})$$

where $\Lambda_t := \mathbf{I} + (B_t(R_t^u)^{-1}B_t^\top - D_t(R_t^w)^{-1}D_t^\top)P_{t+1}^*$ and $P_N^* = Q_N$. For the stochastic game with closed-loop perfect-state information pattern that we considered, we already know from (Başar & Olsder, 1998) that whenever a saddle point exists, the saddle-point control policies are linear state-feedback (i.e., $\mu_{u,t}^*(h_{u,t}) = -K_t^* x_t$ and $\mu_{w,t}^*(h_{w,t}) = -L_t^* x_t$), and the gain matrices K_t^* and L_t^* , for $t \in \{0, \dots, N-1\}$, are unique and can be expressed by

$$K_t^* = (R_t^u)^{-1} B_t^\top P_{t+1}^* \Lambda_t^{-1} A_t, \quad (\text{A.14})$$

$$L_t^* = -(R_t^w)^{-1} D_t^\top P_{t+1}^* \Lambda_t^{-1} A_t, \quad (\text{A.15})$$

where $P_{t+1}^*, t \in \{0, \dots, N-1\}$, is the sequence of p.s.d. matrices generated by (A.13). We next introduce a standard assumption that suffices to ensure the existence of the value of the game, following from Theorem 3.2 of (Başar & Bernhard, 1995) and Theorem 6.7 of (Başar & Olsder, 1998).

Assumption A.4 $R_t^w - D_t^\top P_{t+1}^* D_t > 0$, for all $t \in \{0, \dots, N-1\}$, where $P_{t+1}^* \geq 0$ is generated by (A.13).

²Hereafter we will use player and agent interchangeably.

³The assumption on the boundedness of distribution is only for ease of analysis (Malik et al., 2020; Furieri et al., 2020). Extensions to sub-Gaussian distributions are standard and do not affect the sample complexity result, as noted by (Malik et al., 2020).

Under Assumption A.4, the value in (A.12) is attained by the controller sequence $(\{-K_t^* x_t^*\}, \{-L_t^* x_t^*\})$, where K_t^* and L_t^* are given in (A.14) and (A.15), respectively, and x_{t+1}^* is the corresponding state trajectory generated by $x_{t+1}^* = \Lambda_t^{-1} A_t x_t^* + \xi_t$, for $t \in \{0, \dots, N-1\}$ and with $x_0^* = x_0$ (Başar & Bernhard, 1995). Thus, the solution to (A.12) can be found by searching K_t^* and L_t^* in $\mathbb{R}^{d \times m} \times \mathbb{R}^{n \times m}$, for all $t \in \{0, \dots, N-1\}$. The resulting minimax policy optimization problem is

$$\min_{\{K_t\}} \max_{\{L_t\}} \mathcal{G}(\{K_t\}, \{L_t\}) := \mathbb{E}_{x_0, \xi_0, \dots, \xi_{N-1}} \left[\sum_{t=0}^{N-1} x_t^\top (Q_t + K_t^\top R_t^u K_t - L_t^\top R_t^w L_t) x_t + x_N^\top Q_N x_N \right], \quad (\text{A.16})$$

subject to the system transition $x_{t+1} = (A_t - B_t K_t - D_t L_t) x_t + \xi_t$, for $t \in \{0, \dots, N-1\}$. Some further notes regarding Assumption A.4 are provided in the remark below.

Remark A.5 (Unique Feedback Nash (Saddle-point) Equilibrium) Assumption A.4 is sufficient to guarantee the existence of feedback Nash (equivalently, saddle-point) equilibrium in zero-sum LQ dynamic games, under which it is also unique. The condition is also almost necessary as a Nash equilibrium will not exist and the upper value of the game would be unbounded if $R_t^w - D_t^\top P_{t+1}^* D_t$ for $t \in \{0, \dots, N-1\}$ has at least one negative eigenvalue (Başar & Bernhard, 1995).

Lastly, we formally state the well-known equivalence conditions between LEQG, LQ disturbance attenuation, and zero-sum LQ stochastic dynamic games in the following lemma, and provide a short proof in §B.3.

Lemma A.6 (Connections) For any fixed $\gamma > \gamma^*$ in the LQ disturbance attenuation problem, we can introduce an equivalent LEQG problem and an equivalent zero-sum LQ dynamic game. Specifically, if we set $\beta^{-1} \mathbf{I}, R_t, C_t^\top C_t, W$ in LEQG, $\gamma^2 \mathbf{I}, R_t, C_t^\top C_t, D_t D_t^\top$ in the LQ disturbance attenuation problem, and $R_t^w, R_t^u, Q_t, D_t D_t^\top$ in the game to be the same, for all $t \in \{0, \dots, N-1\}$, then the optimal gain matrices in LEQG, the gain matrix in the LQ disturbance attenuation problem, and the saddle-point gain matrix for the minimizing player in the game are the same.

By the above connections and Remark A.3, we will hereafter focus on the stochastic zero-sum LQ dynamic game between a minimizing controller and a maximizing disturbance, using PG methods with exact gradient accesses in §3, and derivative-free PG methods by sampling system trajectories in §4. As a result, the minimizing controller we obtain solves all three classes of problems introduced above altogether.

A.4. Compact Formulations

In this section, we re-derive the game formulations in §A.3.3, with linear state-feedback policies for the two players, using the following compact notations:

$$\begin{aligned} \mathbf{x} &= [x_0^\top, \dots, x_N^\top]^\top, \quad \mathbf{u} = [u_0^\top, \dots, u_{N-1}^\top]^\top, \quad \mathbf{w} = [w_0^\top, \dots, w_{N-1}^\top]^\top, \quad \boldsymbol{\xi} = [\xi_0^\top, \xi_0^\top, \dots, \xi_{N-1}^\top]^\top, \\ \mathbf{A} &= \begin{bmatrix} \mathbf{0}_{m \times mN} & \mathbf{0}_{m \times m} \\ \text{diag}(A_0, \dots, A_{N-1}) & \mathbf{0}_{mN \times m} \end{bmatrix}, \quad \mathbf{B} = \begin{bmatrix} \mathbf{0}_{m \times dN} \\ \text{diag}(B_0, \dots, B_{N-1}) \end{bmatrix}, \quad \mathbf{D} = \begin{bmatrix} \mathbf{0}_{m \times nN} \\ \text{diag}(D_0, \dots, D_{N-1}) \end{bmatrix}, \\ \mathbf{Q} &= [\text{diag}(Q_0, \dots, Q_N)], \quad \mathbf{R}^u = [\text{diag}(R_0^u, \dots, R_{N-1}^u)], \quad \mathbf{R}^w = [\text{diag}(R_0^w, \dots, R_{N-1}^w)], \\ \mathbf{K} &= [\text{diag}(K_0, \dots, K_{N-1}) \quad \mathbf{0}_{dN \times m}], \quad \mathbf{L} = [\text{diag}(L_0, \dots, L_{N-1}) \quad \mathbf{0}_{nN \times m}]. \end{aligned} \quad (\text{A.17})$$

With the above definitions, the game in §A.3.3 is characterized by the transition dynamics

$$\mathbf{x} = \mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{u} + \mathbf{D}\mathbf{w} + \boldsymbol{\xi} = (\mathbf{A} - \mathbf{B}\mathbf{K} - \mathbf{D}\mathbf{L})\mathbf{x} + \boldsymbol{\xi}, \quad (\text{A.18})$$

where \mathbf{x} is the concatenated system state, $\mathbf{u} = -\mathbf{K}\mathbf{x}$ (respectively, $\mathbf{w} = -\mathbf{L}\mathbf{x}$) is the linear state-feedback controller of the minimizing (resp., maximizing) player, and $\mathbf{A}, \mathbf{B}, \mathbf{D}$ are the system matrices in their corresponding compact forms. Note that \mathbf{A} is essentially a nilpotent matrix with degree of $N+1$ (i.e. a lower triangular matrix with zeros along the main diagonal) such that $\mathbf{A}^{N+1} = \mathbf{0}$. Also, $\boldsymbol{\xi}$ is a vector that concatenates the independently sampled initial state $x_0 \in \mathcal{D}$ and process noises $\xi_t \in \mathcal{D}$, $t \in \{0, \dots, N-1\}$. The objective of the minimizing (resp. maximizing) player is to minimize (resp. maximize) the objective function re-written in terms of the compact notation, that is to solve

$$\inf_{\mathbf{K}} \sup_{\mathbf{L}} \mathcal{G}(\mathbf{K}, \mathbf{L}) := \mathbb{E}_{\boldsymbol{\xi}} \left[\mathbf{x}^\top (\mathbf{Q} + \mathbf{K}^\top \mathbf{R}^u \mathbf{K} - \mathbf{L}^\top \mathbf{R}^w \mathbf{L}) \mathbf{x} \right] \quad (\text{A.19})$$

subject to the transition dynamics (A.18).

An Illustrative Example: Let us now consider a scalar example with $N, n, m, d = 1$ with the following concatenated notations:

$$\begin{aligned} \mathbf{x} &= \begin{bmatrix} x_0 \\ x_1 \end{bmatrix}, \quad \boldsymbol{\xi} = \begin{bmatrix} x_0 \\ \xi_0 \end{bmatrix}, \quad (\mathbf{u}, \mathbf{w}, \mathbf{R}^u, \mathbf{R}^w) = (u_0, w_0, R_0^u, R_0^w), \quad \mathbf{A} = \begin{bmatrix} 0 & 0 \\ A_0 & 0 \end{bmatrix}, \quad \mathbf{B} = \begin{bmatrix} 0 \\ B_0 \end{bmatrix}, \quad \mathbf{D} = \begin{bmatrix} 0 \\ D_0 \end{bmatrix}, \\ \mathbf{Q} &= \begin{bmatrix} Q_0 & 0 \\ 0 & Q_1 \end{bmatrix}, \quad \mathbf{K} = [K_0 \quad 0], \quad \mathbf{L} = [L_0 \quad 0]. \end{aligned}$$

Then, the system dynamics of the game and the linear state-feedback controllers are represented as

$$\begin{aligned} \mathbf{x} &= \mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{u} + \mathbf{D}\mathbf{w} + \boldsymbol{\xi} \iff \begin{bmatrix} x_0 \\ x_1 \end{bmatrix} = \begin{bmatrix} 0 & 0 \\ A_0 & 0 \end{bmatrix} \cdot \begin{bmatrix} x_0 \\ x_1 \end{bmatrix} + \begin{bmatrix} 0 \\ B_0 \end{bmatrix} \cdot u_0 + \begin{bmatrix} 0 \\ D_0 \end{bmatrix} \cdot w_0 + \begin{bmatrix} x_0 \\ \xi_0 \end{bmatrix} = \begin{bmatrix} A_0 x_0 + B_0 u_0 + D_0 w_0 + \xi_0 \\ x_0 \end{bmatrix} \\ \mathbf{u} &= -\mathbf{K}\mathbf{x} \iff u_0 = -[K_0 \quad 0] \cdot \begin{bmatrix} x_0 \\ x_1 \end{bmatrix}, \quad \mathbf{w} = -\mathbf{L}\mathbf{x} \iff w_0 = -[L_0 \quad 0] \cdot \begin{bmatrix} x_0 \\ x_1 \end{bmatrix} \\ \mathbf{x} &= (\mathbf{A} - \mathbf{B}\mathbf{K} - \mathbf{D}\mathbf{L})\mathbf{x} + \boldsymbol{\xi} \iff \begin{bmatrix} x_0 \\ x_1 \end{bmatrix} = \left(\begin{bmatrix} 0 & 0 \\ A_0 & 0 \end{bmatrix} - \begin{bmatrix} 0 \\ B_0 \end{bmatrix} \cdot [K_0 \quad 0] - \begin{bmatrix} 0 \\ D_0 \end{bmatrix} \cdot [L_0 \quad 0] \right) \cdot \begin{bmatrix} x_0 \\ x_1 \end{bmatrix} + \begin{bmatrix} x_0 \\ \xi_0 \end{bmatrix} \\ &\iff \begin{bmatrix} x_0 \\ x_1 \end{bmatrix} = \begin{bmatrix} 0 & 0 \\ A_0 - B_0 K_0 - D_0 L_0 & 0 \end{bmatrix} \cdot \begin{bmatrix} x_0 \\ x_1 \end{bmatrix} + \begin{bmatrix} x_0 \\ \xi_0 \end{bmatrix} = \begin{bmatrix} (A_0 - B_0 K_0 - D_0 L_0)x_0 + \xi_0 \\ x_0 \end{bmatrix}. \end{aligned}$$

Moreover, we define $A_{K_0, L_0} := A_0 - B_0 K_0 - D_0 L_0$ and establish the equivalence between the recursive Lyapunov equation (3.1) and its compact form (3.2) such that

$$\begin{aligned} \mathbf{P}_{K, L} &= (\mathbf{A} - \mathbf{B}\mathbf{K} - \mathbf{D}\mathbf{L})^\top \mathbf{P}_{K, L} (\mathbf{A} - \mathbf{B}\mathbf{K} - \mathbf{D}\mathbf{L}) + \mathbf{Q} + \mathbf{K}^\top \mathbf{R}^u \mathbf{K} - \mathbf{L}^\top \mathbf{R}^w \mathbf{L} \\ \iff \begin{bmatrix} P_{K_0, L_0} & 0 \\ 0 & P_{K_1, L_1} \end{bmatrix} &= \begin{bmatrix} 0 & 0 \\ A_{K_0, L_0} & 0 \end{bmatrix}^\top \begin{bmatrix} P_{K_0, L_0} & 0 \\ 0 & P_{K_1, L_1} \end{bmatrix} \begin{bmatrix} 0 & 0 \\ A_{K_0, L_0} & 0 \end{bmatrix} + \begin{bmatrix} Q_0 + K_0^\top R_0^u K_0 - L_0 R_0^w L_0 & 0 \\ 0 & Q_1 \end{bmatrix} \\ \iff \begin{bmatrix} P_{K_0, L_0} & 0 \\ 0 & P_{K_1, L_1} \end{bmatrix} &= \begin{bmatrix} A_{K_0, L_0}^\top P_{K_1, L_1} A_{K_0, L_0} + Q_0 + K_0^\top R_0^u K_0 - L_0 R_0^w L_0 & 0 \\ 0 & Q_1 \end{bmatrix}. \end{aligned}$$

Next, we note that the solution to (3.3), $\Sigma_{K, L}$, is also a recursive formula in the compact form, such that

$$\begin{aligned} \Sigma_{K, L} &= (\mathbf{A} - \mathbf{B}\mathbf{K} - \mathbf{D}\mathbf{L})\Sigma_{K, L}(\mathbf{A} - \mathbf{B}\mathbf{K} - \mathbf{D}\mathbf{L})^\top + \Sigma_0 \\ \iff \mathbb{E} \begin{bmatrix} x_0 x_0^\top & 0 \\ 0 & x_1 x_1^\top \end{bmatrix} &= \begin{bmatrix} 0 & 0 \\ A_{K_0, L_0} & 0 \end{bmatrix} \mathbb{E} \begin{bmatrix} x_0 x_0^\top & 0 \\ 0 & x_1 x_1^\top \end{bmatrix} \begin{bmatrix} 0 & 0 \\ A_{K_0, L_0} & 0 \end{bmatrix}^\top + \mathbb{E} \begin{bmatrix} x_0 x_0^\top & 0 \\ 0 & \xi_0 \xi_0^\top \end{bmatrix} \\ \iff \mathbb{E} \begin{bmatrix} x_0 x_0^\top & 0 \\ 0 & x_1 x_1^\top \end{bmatrix} &= \mathbb{E} \begin{bmatrix} x_0 x_0^\top & 0 \\ 0 & A_{K_0, L_0} x_0 x_0^\top A_{K_0, L_0}^\top + \xi_0 \xi_0^\top \end{bmatrix}. \end{aligned}$$

This completes the illustrative example.

A.5. Discussions

Remark A.7 (Double-Loop Scheme) Our double-loop scheme, which has also been used before in (Zhang et al., 2019; Bu et al., 2019b), is pertinent to the descent-multi-step-ascent scheme (Nouiehed et al., 2019) and the gradient descent-ascent (GDA) scheme with two-timescale stepsizes (Jin et al., 2020; Lin et al., 2020; Fiez & Ratliff, 2020) for solving nonconvex-(non)concave minimax optimization problems, with the number of multi-steps of ascent going to infinity, or the ratio between the fast and slow stepsizes going to infinity (the τ -GDA scheme with $\tau \rightarrow \infty$ (Jin et al., 2020; Fiez & Ratliff, 2020)). Another variant is the alternating GDA (AGDA) scheme, which has been investigated in (Yang et al., 2020) to address nonconvex-nonconcave problems with two-sided PL condition. However, unlike this literature, one key condition for establishing convergence, the global smoothness of the objective function, does not hold in our control setting with unbounded decision spaces and objective functions. In fact, as per Lemma 3.3, a non-judicious choice of \mathbf{K} may lead to an undefined inner-loop maximization problem with unbounded objective. Hence, one has to carefully control the update-rule here, to ensure that the iterates do not yield unbounded/undefined values along iterations. As we will show in §D.4, it is not hard to construct cases where descent-multi-step-ascent/AGDA/ τ -GDA diverges even with infinitesimal stepsizes, similar to the negative results reported in the infinite-horizon setting in (Zhang et al., 2020). These results suggest that it seems challenging to provably show that other candidate update schemes converge to the global NE in zero-sum LQ games, except the double-loop one, as we will show next.

Remark A.8 (*Inner-Loop Landscape*) The optimization landscape of the inner loop is similar to that of the standard LQR (Fazel et al., 2018). In our setting, independent process noises (together with the random initial state) with a positive-definite covariance matrix guarantee the full-rankness of $\Sigma_{\mathbf{K}, \mathbf{L}}$, which leads to the stationary point of the inner-loop objective function $\mathcal{G}(\mathbf{K}, \cdot)$ also being the unique optimal solution (cf. Lemma 3.2). Further, the full-rankness of $\Sigma_{\mathbf{K}, \mathbf{L}}$ also ensures that $\mathcal{G}(\mathbf{K}, \cdot)$ is PL (cf. Lemma 3.4). These two properties together are essential for establishing the global convergence of PG methods in our nonconvex-nonconcave setting.

Remark A.9 (*IR*) Suppose that $\mathbf{K}_0 \in \mathcal{K}$. Then Lemma A.6 shows that \mathbf{K}_0 is the control gain matrix that attains a γ -level of disturbance attenuation. By the IR property in Theorem 3.7, every iterate $\mathbf{K}_k \in \mathcal{K}$ for all $k \geq 0$ following the NPG (3.11) or the GN (3.13) update rules will thus preserve this γ -level of disturbance attenuation throughout the policy optimization (learning) process. Theorem 3.7 thus provides some provable robustness guarantees for two specific policy search directions, (3.11) and (3.13), which is important for safety-critical control systems in the presence of adversarial disturbances, since otherwise, the system performance index can be driven to arbitrarily large values.

B. Proofs of Main Results

B.1. Proof of Lemma A.1

Proof Since $x_0 \sim \mathcal{N}(\mathbf{0}, X_0)$, $w_0, \dots, w_{N-1} \sim \mathcal{N}(\mathbf{0}, W)$, $P_{K_t} \geq 0$ for all t , $X_0^{-1} - \beta P_{K_0} > 0$, and $W^{-1} - \beta P_{K_t} > 0$, for all $t \in \{1, \dots, N\}$, we have by Lemma C.1 of (Zhang et al., 2021a) that the objective function of LEQG can be represented by (A.3). Also, the conditions that $X_0^{-1} - \beta P_{K_0} > 0$ and $W^{-1} - \beta P_{K_t} > 0$, for all $t \in \{0, \dots, N\}$, ensure the existence of the expression Σ_{K_t} , which can be verified through applying matrix inversion lemma as in (B.1). Thus, the existence of the expression $\nabla_{K_t} \mathcal{J}(\{K_t\})$ is also ensured, for all $t \in \{0, \dots, N-1\}$. Let us now use ∇_{K_t} to denote the derivative w.r.t. K_t . Then, for $t \in \{1, \dots, N-1\}$, we have

$$\begin{aligned} \nabla_{K_t} \mathcal{J}(\{K_t\}) &= -\frac{1}{\beta} \text{Tr} \left\{ (\mathbf{I} - \beta P_{K_0} X_0)^{-\top} \cdot [\nabla_{K_t} (\mathbf{I} - \beta P_{K_t} X_0)]^\top \right\} - \frac{1}{\beta} \sum_{\tau=1}^t \text{Tr} \left\{ (\mathbf{I} - \beta P_{K_\tau} W)^{-\top} \cdot [\nabla_{K_t} (\mathbf{I} - \beta P_{K_\tau} W)]^\top \right\} \\ &= \text{Tr} [(\mathbf{I} - \beta P_{K_0} X_0)^{-1} \cdot \nabla_{K_t} (P_{K_0} X_0)] + \sum_{\tau=1}^t \text{Tr} [(\mathbf{I} - \beta P_{K_\tau} W)^{-1} \cdot \nabla_{K_t} (P_{K_\tau} W)] \\ &= \text{Tr} [\nabla_{K_t} P_{K_0} \cdot X_0 (\mathbf{I} - \beta P_{K_0} X_0)^{-1}] + \sum_{\tau=1}^t \text{Tr} [\nabla_{K_t} P_{K_\tau} \cdot W (\mathbf{I} - \beta P_{K_\tau} W)^{-1}] \\ &= \text{Tr} [\nabla_{K_t} P_{K_0} \cdot X_0^{\frac{1}{2}} (\mathbf{I} - \beta X_0^{\frac{1}{2}} P_{K_0} X_0^{\frac{1}{2}})^{-1} X_0^{\frac{1}{2}}] + \sum_{\tau=1}^t \text{Tr} [\nabla_{K_t} P_{K_\tau} \cdot W^{\frac{1}{2}} (\mathbf{I} - \beta W^{\frac{1}{2}} P_{K_\tau} W^{\frac{1}{2}})^{-1} W^{\frac{1}{2}}], \end{aligned}$$

where the first equality is due to $\nabla_X \log \det X = X^{-\top}$, the chain rule, and the fact that P_{K_τ} , for all $\tau \in \{t+1, \dots, N\}$, is independent of K_t . The second and third equalities are due to the cyclic property of matrix trace, $\text{Tr}(A^\top B^\top) = \text{Tr}(AB)$, and that X_0, W are independent of K_t . The last equality uses matrix inversion lemma, such that for $V \in \{X_0, W\}$:

$$V \cdot (\mathbf{I} - \beta P_K V)^{-1} = V^{\frac{1}{2}} (\mathbf{I} - \beta V^{\frac{1}{2}} P_K V^{\frac{1}{2}})^{-1} V^{\frac{1}{2}} \quad (\text{B.1})$$

Then, by the definition of the RDE in (A.2) and defining $M_0 := X_0^{\frac{1}{2}} (\mathbf{I} - \beta X_0^{\frac{1}{2}} P_{K_0} X_0^{\frac{1}{2}})^{-1} X_0^{\frac{1}{2}}$, $M_t := W^{\frac{1}{2}} (\mathbf{I} - \beta W^{\frac{1}{2}} P_{K_t} W^{\frac{1}{2}})^{-1} W^{\frac{1}{2}}$ for all $t \in \{1, \dots, N-1\}$, we have:

$$\begin{aligned} \forall t \in \{0, \dots, N-1\}, \quad \nabla_{K_t} \mathcal{J}(\{K_t\}) &= \sum_{\tau=0}^t \text{Tr} [\nabla_{K_t} P_{K_\tau} \cdot M_\tau] \\ &= \sum_{\tau=0}^{t-1} \text{Tr} [\nabla_{K_t} P_{K_\tau} \cdot M_\tau] + 2[(R_t + B_t^\top \tilde{P}_{K_{t+1}} B_t) K_t - B_t^\top \tilde{P}_{K_{t+1}} A_t M_t] \\ &\quad + \text{Tr} [(A_t - B_t K_t)^\top \cdot \nabla_{K_t} \tilde{P}_{K_{t+1}} \cdot (A_t - B_t K_t) M_t] \\ &= \sum_{\tau=0}^{t-1} \text{Tr} [\nabla_{K_t} P_{K_\tau} \cdot M_\tau] + 2[(R_t + B_t^\top \tilde{P}_{K_{t+1}} B_t) K_t - B_t^\top \tilde{P}_{K_{t+1}} A_t M_t], \end{aligned}$$

where the last equality is due to $\tilde{P}_{K_{t+1}}$ not depending on K_t and thus $\nabla_{K_t} \tilde{P}_{K_{t+1}} = \mathbf{0}$. Then, for all $\tau \in \{1, \dots, t-1\}$, we first recall the definition that $\tilde{P}_{K_\tau} = P_{K_\tau} + \beta P_{K_\tau} (W^{-1} - \beta P_{K_\tau})^{-1} P_{K_\tau}$, and then take the derivatives on both sides with respect to K_t to get

$$\begin{aligned} \nabla_{K_t} \tilde{P}_{K_\tau} &= \nabla_{K_t} [P_{K_\tau} + \beta P_{K_\tau} (W^{-1} - \beta P_{K_\tau})^{-1} P_{K_\tau}] = \nabla_{K_t} [(\mathbf{I} - \beta P_{K_\tau} W)^{-1} P_{K_\tau}] \\ &= (\mathbf{I} - \beta P_{K_\tau} W)^{-1} \cdot \nabla_{K_t} P_{K_\tau} \cdot \beta W (\mathbf{I} - \beta P_{K_\tau} W)^{-1} P_{K_\tau} + (\mathbf{I} - \beta P_{K_\tau} W)^{-1} \cdot \nabla_{K_t} P_{K_\tau} \\ &= (\mathbf{I} - \beta P_{K_\tau} W)^{-1} \cdot \nabla_{K_t} P_{K_\tau} \cdot [W^{\frac{1}{2}} (\beta^{-1} \mathbf{I} - W^{\frac{1}{2}} P_{K_\tau} W^{\frac{1}{2}})^{-1} W^{\frac{1}{2}} P_{K_\tau} + \mathbf{I}] \\ &= (\mathbf{I} - \beta P_{K_\tau} W)^{-1} \cdot \nabla_{K_t} P_{K_\tau} \cdot (\mathbf{I} - \beta W P_{K_\tau})^{-1} = (\mathbf{I} - \beta P_{K_\tau} W)^{-1} \cdot \nabla_{K_t} P_{K_\tau} \cdot (\mathbf{I} - \beta P_{K_\tau} W)^{-\top}, \end{aligned}$$

where the second and the fifth equalities use matrix inversion lemma, the third equality uses $\nabla_X (P^{-1}) = -P^{-1} \cdot \nabla_X P \cdot P^{-1}$, and the fourth equality uses (B.1). Now, for any $\tau \in \{0, \dots, t-1\}$, we can iteratively unroll the RDE (A.2) to obtain that

$$\begin{aligned} \text{Tr}[\nabla_{K_t} P_{K_\tau} \cdot M_\tau] &= 2[(R_t + B_t^\top \tilde{P}_{K_{t+1}} B_t) K_t - B_t^\top \tilde{P}_{K_{t+1}} A_t] \\ &\quad \cdot \prod_{i=\tau}^{t-1} [(\mathbf{I} - \beta P_{K_{i+1}} W)^{-\top} (A_i - B_i K_i)] \cdot M_\tau \cdot \prod_{i=\tau}^{t-1} [(A_i - B_i K_i)^\top (\mathbf{I} - \beta P_{K_{i+1}} W)^{-1}]. \end{aligned}$$

Taking a summation proves the PG expression in (A.4). ■

B.2. Proof of Lemma A.2

Proof We first prove the policy gradient of (A.8). By the conditions that $\gamma^2 \mathbf{I} - D_t^\top P_{K_{t+1}} D_t > 0$, for all $t \in \{0, \dots, N-1\}$, and $\gamma^2 \mathbf{I} - P_{K_0} > 0$, the expression Σ_{K_t} exists and thus the expression $\nabla_{K_t} \bar{\mathcal{J}}(\{K_t\})$ exists, for all $t \in \{0, \dots, N-1\}$. Then, the PG expression of (A.8) follows from the proof of Lemma A.1 by replacing β^{-1} and W therein with γ^2 and $D_t D_t^\top$, respectively, and also replacing the W in the first term of the summation in the expression of Σ_{K_t} by \mathbf{I} . Next, we prove the PG of (A.9) w.r.t. K_t . Let us now use ∇_{K_t} to denote the derivative w.r.t. K_t . Then, for $t \in \{0, \dots, N-1\}$, we have

$$\nabla_{K_t} \bar{\mathcal{J}}(\{K_t\}) = \text{Tr}[\nabla_{K_t} P_{K_0}] + \sum_{\tau=1}^t \text{Tr}[\nabla_{K_t} (P_{K_\tau} D_{\tau-1} D_{\tau-1}^\top)] = \text{Tr}[\nabla_{K_t} P_{K_0}] + \sum_{\tau=1}^t \text{Tr}[D_{\tau-1}^\top \cdot \nabla_{K_t} P_{K_\tau} \cdot D_{\tau-1}],$$

where the first equality is due to that P_{K_τ} , for all $\tau \in \{t+1, \dots, N\}$, does not depend on K_t . Recalling the RDE in (A.7), we can unroll the recursive formula to obtain

$$\begin{aligned} \text{Tr}[\nabla_{K_t} P_{K_0}] &= 2[(R_t + B_t^\top \tilde{P}_{K_{t+1}} B_t) K_t - B_t^\top \tilde{P}_{K_{t+1}} A_t] \\ &\quad \cdot \prod_{i=0}^{t-1} [(\mathbf{I} - \gamma^{-2} P_{K_{i+1}} D_i D_i^\top)^{-\top} (A_i - B_i K_i) (A_i - B_i K_i)^\top (\mathbf{I} - \gamma^{-2} P_{K_{i+1}} D_i D_i^\top)^{-1}]. \end{aligned}$$

Moreover, for any $\tau \in \{1, \dots, t\}$, we have

$$\begin{aligned} \text{Tr}[D_{\tau-1}^\top \cdot \nabla_{K_t} P_{K_\tau} \cdot D_{\tau-1}] &= 2[(R_t + B_t^\top \tilde{P}_{K_{t+1}} B_t) K_t - B_t^\top \tilde{P}_{K_{t+1}} A_t] \\ &\quad \cdot \prod_{i=\tau}^{t-1} [(\mathbf{I} - \gamma^{-2} P_{K_{i+1}} D_i D_i^\top)^{-\top} (A_i - B_i K_i)] \cdot D_{\tau-1} D_{\tau-1}^\top \cdot \prod_{i=\tau}^{t-1} [(A_i - B_i K_i)^\top (\mathbf{I} - \gamma^{-2} P_{K_{i+1}} D_i D_i^\top)^{-1}]. \end{aligned}$$

Taking a summation proves the expression of the policy gradient in (A.10). ■

B.3. Proof of Lemma A.6

Proof The equivalence relationships are proved in three parts. Firstly, by Corollary 6.4 of (Başar & Olsder, 1998), the discrete-time zero-sum LQ dynamic game with additive stochastic disturbance (zero-mean, independent sequence), and

with closed-loop perfect-state information pattern, admits the same saddle-point gain matrices for the minimizer and the maximizer as the deterministic version of the game (that is, without the stochastic driving term).

Subsequently, by Theorem 4.1 of (Başar, 1991), we can introduce a corresponding (deterministic) zero-sum LQ dynamic game with $R_t^w = \gamma^2 \mathbf{I}$ for any fixed $\gamma > \gamma^*$ in the LQ disturbance attenuation problem. Then, if we replace R_t^w, R_t^u, Q_t in the deterministic version of the game presented in §A.3.3 by $\gamma^2 \mathbf{I}, R_t, C_t^\top C_t$ in §2.2, respectively, for all $t \in \{0, \dots, N-1\}$, and also set other shared parameters in §A.3.3 and §2.2 to be the same, the saddle-point gain matrix for the minimizing player in that game is equivalent to the gain matrix in the disturbance attenuation problem.

Lastly, by §VI.B of (Jacobson, 1973) and if we replace $R_t^w, R_t^u, Q_t, D_t D_t^\top$ in the deterministic version of the game presented in §A.3.3 by $\beta^{-1} \mathbf{I}, R_t, C_t^\top C_t, W$ in §2.1 for all $t \in \{0, \dots, N-1\}$, and also set other shared parameters in §A.3.3 to be the same as the ones in §2.1, then the deterministic zero-sum LQ dynamic game is equivalent to the LEQG problem with perfect state measurements, in the sense that the optimal gain matrix in the latter is the same as the saddle-point gain matrix for the minimizing player in the former. This completes the proof. ■

B.4. Proof of Lemma 3.1

Proof We first prove that for a fixed $\mathbf{K} \in \mathcal{K}$, the optimization problem with respect to \mathbf{L} can be nonconcave, by establishing an example of the game. Here, \mathcal{K} is the feasible set as defined in (3.7), which ensures that for a fixed \mathbf{K} , the solution to the maximization problem with respect to \mathbf{L} is well defined. Specifically, we consider a 3-dimensional linear time-invariant system with $A_t = A, B_t = B, D_t = D, Q_t = Q, R_t^u = R^u$, and $R_t^w = R^w$ for all t , where

$$A = \begin{bmatrix} 1 & 0 & -5 \\ -1 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}, \quad B = \begin{bmatrix} 1 & -10 & 0 \\ 0 & 3 & 1 \\ -1 & 0 & 2 \end{bmatrix}, \quad D = \begin{bmatrix} 0.5 & 0 & 0 \\ 0 & 0.2 & 0 \\ 0 & 0 & 0.2 \end{bmatrix}, \quad Q = \begin{bmatrix} 2 & -1 & 0 \\ -1 & 2 & -1 \\ 0 & -1 & 2 \end{bmatrix}, \quad R^u = \begin{bmatrix} 4 & -1 & 0 \\ -1 & 4 & -2 \\ 0 & -2 & 3 \end{bmatrix},$$

and $R^w = 5 \cdot \mathbf{I}$. Also, we set $\Sigma_0 = \mathbf{I}, N = 5$, and choose the time-invariant control gain matrices $K_t = K, L_t^1 = L^1, L_t^2 = L^2, L_t^3 = L^3$ for all t , where

$$K = \begin{bmatrix} -0.12 & -0.01 & 0.62 \\ -0.21 & 0.14 & 0.15 \\ -0.06 & 0.05 & 0.42 \end{bmatrix}, \quad L^1 = \begin{bmatrix} -0.86 & 0.97 & 0.14 \\ -0.82 & 0.36 & 0.51 \\ 0.98 & 0.08 & -0.20 \end{bmatrix}, \quad L^2 = \begin{bmatrix} -0.70 & -0.37 & 0.09 \\ -0.54 & -0.28 & 0.23 \\ 0.74 & 0.62 & -0.51 \end{bmatrix}, \quad L^3 = \frac{L^1 + L^2}{2}.$$

The concatenated matrices $A, B, D, Q, R^u, R^w, K, L^1, L^2, L^3$ are generated following the definitions in §3. Then, we first note that $\mathbf{K} \in \mathcal{K}$ holds, as $\mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})} \geq 0$ exists and one can calculate that $\lambda_{\min}(\mathbf{R}^w - \mathbf{D}^\top \mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})} \mathbf{D}) = 0.5041 > 0$. Subsequently, we can verify that $(\mathcal{G}(\mathbf{K}, L^1) + \mathcal{G}(\mathbf{K}, L^2))/2 - \mathcal{G}(\mathbf{K}, L^3) = 6.7437 > 0$. Thus, we can conclude that there exists a $\mathbf{K} \in \mathcal{K}$ such that the objective function is nonconcave with respect to \mathbf{L} . Next, we prove the other argument by choosing the following time-invariant control gain matrices $L_t = L, K_t^1 = K^1, K_t^2 = K^2, K_t^3 = K^3$ for all t , where

$$L = \mathbf{0}, \quad K^1 = \begin{bmatrix} 1.44 & 0.31 & -1.18 \\ 0.03 & -0.13 & -0.39 \\ 0.36 & -1.71 & 0.24 \end{bmatrix}, \quad K^2 = \begin{bmatrix} -0.08 & -0.16 & -1.96 \\ -0.13 & -1.12 & 1.28 \\ 1.67 & -0.91 & 1.71 \end{bmatrix}, \quad K^3 = \frac{K^1 + K^2}{2}.$$

Following similar steps, we can verify that $(\mathcal{G}(K^1, L) + \mathcal{G}(K^2, L))/2 - \mathcal{G}(K^3, L) = -1.2277 \times 10^5 < 0$. Therefore, there exists \mathbf{L} such that the objective function is nonconvex with respect to \mathbf{K} . ■

B.5. Proof of Lemma 3.2

Proof For $t \in \{0, \dots, N-1\}$, we can define the cost-to-go function starting from time t to N as \mathcal{G}_t . Recalling the definition of P_{K_t, L_t} in (3.1), we can write \mathcal{G}_t as

$$\begin{aligned} \mathcal{G}_t &= \mathbb{E}_{x_t, \xi_t, \dots, \xi_{N-1}} \left\{ x_t^\top P_{K_t, L_t} x_t + \sum_{\tau=t}^{N-1} \xi_\tau^\top P_{K_{\tau+1}, L_{\tau+1}} \xi_\tau \right\} = \mathbb{E}_{x_t} \{ x_t^\top P_{K_t, L_t} x_t \} + \mathbb{E}_{\xi_t, \dots, \xi_{N-1}} \left\{ \sum_{\tau=t}^{N-1} \xi_\tau^\top P_{K_{\tau+1}, L_{\tau+1}} \xi_\tau \right\} \\ &= \mathbb{E}_{x_t} \{ x_t^\top (Q_t + K_t^\top R_t^u K_t - L_t^\top R_t^w L_t) x_t + x_t^\top (A_t - B_t K_t - D_t L_t)^\top P_{K_{t+1}, L_{t+1}} (A_t - B_t K_t - D_t L_t) x_t \} \\ &\quad + \mathbb{E}_{\xi_t, \dots, \xi_{N-1}} \left\{ \sum_{\tau=t}^{N-1} \xi_\tau^\top P_{K_{\tau+1}, L_{\tau+1}} \xi_\tau \right\}. \end{aligned}$$

We denote $\nabla_{L_t} \mathcal{G}_t := \frac{\partial \mathcal{G}_t}{\partial L_t}$ for notational simplicity, and note that the last term of the above equation does not depend on L_t . Then, we can show that $\nabla_{L_t} \mathcal{G}_t = 2[(-R_t^w + D_t^\top P_{K_{t+1}, L_{t+1}} D_t) L_t - D_t^\top P_{K_{t+1}, L_{t+1}} (A_t - B_t K_t)] \mathbb{E}_{x_t} \{x_t x_t^\top\}$. By writing the gradient of each time step compactly and taking expectations over x_0 and the additive noises $\{\xi_0, \dots, \xi_{N-1}\}$ along the trajectory, we can derive the exact PG of $\mathcal{G}(K, L)$ with respect to L as $\nabla_L \mathcal{G}(K, L) = 2[(-R^w + D^\top P_{K,L} D)L - D^\top P_{K,L} (A - BK)] \Sigma_{K,L}$, where $\Sigma_{K,L}$ is the solution to (3.3). Using similar techniques, we can derive, for a fixed L , the PG of $\mathcal{G}(K, L)$ with respect to K to be $\nabla_K \mathcal{G}(K, L) = 2[(R^u + B^\top P_{K,L} B)K - B^\top P_{K,L} (A - DL)] \Sigma_{K,L}$. Since Σ_0 being full-rank ensures that $\Sigma_{K,L}$ is also full-rank and we have $\nabla_K \mathcal{G}(K, L) = \nabla_L \mathcal{G}(K, L) = 0$ at the stationary points, the PG forms derived above imply

$$K = (R^u + B^\top P_{K,L} B)^{-1} B^\top P_{K,L} (A - DL), \quad L = (-R^w + D^\top P_{K,L} D)^{-1} D^\top P_{K,L} (A - BK), \quad (\text{B.2})$$

where by assumption, we have $R^w - D^\top P_{K,L} D > 0$, and $R^u + B^\top P_{K,L} B > 0$ is further implied by $P_{K,L} \geq 0$ (and thus they are both invertible). Therefore, we can solve (B.2) to obtain

$$\begin{aligned} K &= (R^u)^{-1} B^\top P_{K,L} [I + (B(R^u)^{-1} B^\top - D(R^w)^{-1} D^\top) P_{K,L}]^{-1} A, \\ L &= -(R^w)^{-1} D^\top P_{K,L} [I + (B(R^u)^{-1} B^\top - D(R^w)^{-1} D^\top) P_{K,L}]^{-1} A, \end{aligned}$$

which are (A.14) and (A.15) in their corresponding compact forms. Moreover, at the Nash equilibrium, the solution to (A.13) in its compact form, denoted as P^* , is also the solution to the Lyapunov equation

$$P^* = Q + (K^*)^\top R^u K^* - (L^*)^\top R^w L^* + (A - BK^* - DL^*)^\top P^* (A - BK^* - DL^*).$$

Lastly, under Assumption A.4, the solution to the RDE (A.13) is uniquely computed. As a result, the stationary point of $\mathcal{G}(K, L)$ is also the unique Nash equilibrium of the game. This completes the proof. \blacksquare

B.6. Proof of Lemma 3.3

Proof Let the sequence of outer-loop control gains $\{K_t\}$, $t \in \{0, \dots, N-1\}$, be fixed. To facilitate the analysis, we first construct a sequence of auxiliary zero-sum LQ dynamic games, denoted as $\{\Gamma_{d_t}\}$, $t \in \{0, \dots, N-1\}$, where each Γ_{d_t} has closed-loop state-feedback information structure. Specifically, for any t , Γ_{d_t} starts from some arbitrary \bar{x}_0 and follows a deterministic system dynamics $\bar{x}_{\tau+1} = \bar{A}_\tau \bar{x}_\tau + \bar{B}_\tau \bar{u}_\tau + \bar{D}_\tau \bar{w}_\tau$ for a finite horizon of $N-t$, where

$$\bar{A}_\tau := A_{t+\tau} - B_{t+\tau} K_{t+\tau}, \quad \bar{B}_\tau := 0, \quad \bar{D}_\tau := D_{t+\tau}, \quad \bar{u}_\tau = -\bar{K}_\tau \bar{x}_\tau = -K_{t+\tau} \bar{x}_\tau, \quad \bar{w}_\tau = -\bar{L}_\tau \bar{x}_\tau.$$

The weighting matrices of Γ_{d_t} are chosen to be $\bar{Q}_\tau := Q_{t+\tau}$, $\bar{R}_\tau^u := R_{t+\tau}^u$, and $\bar{R}_\tau^w := R_{t+\tau}^w$ for all $\tau \in \{0, \dots, N-t-1\}$ and $\bar{Q}_{N-t} = Q_N$. Moreover, for any sequence of control gains $\{\bar{L}_\tau\}$, $\tau \in \{0, \dots, N-t-1\}$ in Γ_{d_t} , we define $\{\bar{P}_{\bar{L}_\tau}\}$, $\tau \in \{0, \dots, N-t\}$, as the sequence of solutions generated by the recursive Lyapunov equation

$$\bar{P}_{\bar{L}_\tau} = \bar{Q}_\tau + \bar{K}_\tau^\top \bar{R}_\tau^u \bar{K}_\tau - \bar{L}_\tau^\top \bar{R}_\tau^w \bar{L}_\tau + (\bar{A}_\tau - \bar{D}_\tau \bar{L}_\tau)^\top \bar{P}_{\bar{L}_{\tau+1}} (\bar{A}_\tau - \bar{D}_\tau \bar{L}_\tau), \quad \bar{P}_{\bar{L}_{N-t}} = \bar{Q}_{N-t}. \quad (\text{B.3})$$

Then, for any $t \in \{0, \dots, N-1\}$, Theorem 3.2 of (Başar & Bernhard, 1995) suggests that Γ_{d_t} admits a unique feedback saddle-point if the sequence of p.s.d. solutions $\{\bar{P}_\tau\}$, $\tau \in \{0, \dots, N-t\}$, computed by the recursive Riccati equation

$$\bar{P}_\tau = \bar{Q}_\tau + \bar{K}_\tau^\top \bar{R}_\tau^u \bar{K}_\tau + \bar{A}_\tau^\top \bar{P}_{\tau+1} \bar{A}_\tau + \bar{A}_\tau^\top \bar{P}_{\tau+1} \bar{D}_\tau (\bar{R}_\tau^w - \bar{D}_\tau^\top \bar{P}_{\tau+1} \bar{D}_\tau)^{-1} \bar{D}_\tau^\top \bar{P}_{\tau+1} \bar{A}_\tau, \quad \bar{P}_{N-t} = \bar{Q}_{N-t}. \quad (\text{B.4})$$

exists and satisfies

$$\bar{R}_\tau^w - \bar{D}_\tau^\top \bar{P}_{\tau+1} \bar{D}_\tau > 0, \quad \tau \in \{0, \dots, N-t-1\}. \quad (\text{B.5})$$

Whenever exist, the saddle-point solutions at time τ , for any $\tau \in \{0, \dots, N-t-1\}$, are given by

$$\bar{u}_\tau^* = 0, \quad \bar{w}_\tau^* = -\bar{L}_\tau^* \bar{x}_\tau = (\bar{R}_\tau^w - \bar{D}_\tau^\top \bar{P}_{\tau+1} \bar{D}_\tau)^{-1} \bar{D}_\tau^\top \bar{P}_{\tau+1} \bar{A}_\tau \bar{x}_\tau. \quad (\text{B.6})$$

Moreover, the saddle-point value of Γ_{d_t} is $\bar{x}_0^\top \bar{P}_0 \bar{x}_0$ and the upper value of Γ_{d_t} becomes unbounded when the matrix $\bar{R}_\tau^w - \bar{D}_\tau^\top \bar{P}_{\tau+1} \bar{D}_\tau$ has at least one negative eigenvalue for at least one $\tau \in \{0, \dots, N-t-1\}$. Further, since the choice of

\bar{x}_0 was arbitrary, it holds that $\bar{P}_0 \geq \bar{P}_{\bar{L}_0}$ for any sequence of $\{\bar{L}_\tau\}$, $\tau \in \{0, \dots, N-t-1\}$. Because the above properties hold for all Γ_{d_t} , $t \in \{0, \dots, N-1\}$, we can generate a compact matrix $P_{K,L(K)} \in \mathbb{R}^{m(N+1) \times m(N+1)}$ by putting \bar{P}_0 for each Γ_{d_t} , $t \in \{0, \dots, N-1\}$, sequentially on the diagonal and putting Q_N at the end of the diagonal. The resulting $P_{K,L(K)}$ then solves (3.6), and satisfies $P_{K,L(K)} \geq P_{K,L}$, for all $L \in \mathcal{S}(n, m, N)$, where $P_{K,L}$ is computed by (3.2).

Subsequently, we construct an auxiliary zero-sum LQ dynamic game Γ_s with closed-loop state-feedback information pattern, using compact notations. In particular, Γ_s has a horizon of N and a (stochastic) system dynamics being $\bar{x} = \bar{A}\bar{x} + \bar{B}\bar{u} + \bar{D}\bar{w} + \bar{\xi}$, where the initial state is some arbitrary \bar{x}_0 and

$$\bar{A} := A - BK, \quad \bar{B} := 0, \quad \bar{D} := D, \quad \bar{\xi} := [x_0^\top, \xi_0^\top, \dots, \xi_{N-1}^\top]^\top, \quad \bar{u} = -\bar{K}\bar{x}, \quad \bar{w} = -\bar{L}\bar{x}.$$

We assume that $x_0, \xi_0, \dots, \xi_{N-1} \sim \mathcal{D}$ are zero-mean and independent random variables. The weighting matrices of Γ_s are chosen to be $\bar{Q} := Q$, $\bar{R}^u := R^u$, and $\bar{R}^w := R^w$. By Corollary 6.4 of (Başar & Olsder, 1998), Γ_s is equivalent to Γ_{d_0} above in that the Riccati equation (B.4) and the saddle-point gain matrices (B.6) are the same for Γ_s and Γ_{d_0} .

Note that for a fixed outer-loop control policy $u = -Kx$, the inner-loop subproblem is the same as Γ_s , in that Γ_s essentially absorbs the fixed outer-loop control input into system dynamics. Therefore, the properties of Γ_s , which are equivalent to those of Γ_{d_0} , apply to our inner-loop subproblem. Specifically, (3.6) is a compact representation of (B.4) and (3.7) is the set of K such that (B.4) admits a sequence of p.s.d. solutions satisfying (B.5). Therefore, $K \in \mathcal{K}$ is a sufficient and almost necessary condition for the solution to the inner-loop to be well defined, in that if $K \notin \mathcal{K}$, the inner-loop objective $\mathcal{G}(K, \cdot)$ can be driven to arbitrarily large values. Whenever the solution to the inner-loop exists, it is unique and takes the form in Lemma 3.3. This completes the proof. ■

B.7. Proof of Lemma 3.4

Proof Let $K \in \mathcal{K}$ be fixed. First, the nonconcavity of the inner-loop objective $\mathcal{G}(K, \cdot)$ follows directly from Lemma 3.1. Then, we have by (3.7) and the definitions of the compact matrices that $P_{K_t, L(K_t)} \geq 0$ exists and satisfies $R_t^w - D_t^\top P_{K_{t+1}, L(K_{t+1})} D_t > 0$, for all $t \in \{0, \dots, N-1\}$. Also, for any sequence of $\{L_t\}$, the sequence of solutions to the recursive Lyapunov equation (3.1), $\{P_{K_t, L_t}\}$, always exists and is unique: by Lemma 3.3, it holds that $P_{K_t, L(K_t)} \geq P_{K_t, L_t}$ and thus $R_t^w - D_t^\top P_{K_{t+1}, L_{t+1}} D_t > 0$ for all $t \in \{0, \dots, N-1\}$. Now, for any $t \in \{0, \dots, N-1\}$, we have by (3.1) that

$$\begin{aligned} P_{K_t, L_t} &= (A_t - B_t K_t - D_t L_t)^\top P_{K_{t+1}, L_{t+1}} (A_t - B_t K_t - D_t L_t) + Q_t + K_t^\top R_t^u K_t - L_t^\top R_t^w L_t \\ &= (A_t - B_t K_t)^\top P_{K_{t+1}, L_{t+1}} (A_t - B_t K_t) + Q_t + K_t^\top R_t^u K_t - (A_t - B_t K_t)^\top P_{K_{t+1}, L_{t+1}} D_t L_t \\ &\quad - L_t^\top D_t^\top P_{K_{t+1}, L_{t+1}} (A_t - B_t K_t) - L_t^\top (R_t^w - D_t^\top P_{K_{t+1}, L_{t+1}} D_t) L_t. \end{aligned} \quad (\text{B.7})$$

Note that as $\|L_t\| \rightarrow +\infty$, the quadratic term in (B.7), with leading matrix $-(R_t^w - D_t^\top P_{K_{t+1}, L_{t+1}} D_t) < 0$, dominates other terms. Thus, as $\|L_t\| \rightarrow +\infty$ for some $t \in \{0, \dots, N-1\}$, $\lambda_{\min}(P_{K_t, L_t}) \rightarrow -\infty$, which further makes $\lambda_{\min}(P_{K, L}) \rightarrow -\infty$. In compact forms, this means that $\lambda_{\min}(P_{K, L}) \rightarrow -\infty$ as $\|L\| \rightarrow +\infty$. Then, the inner-loop objective for a fixed $K \in \mathcal{K}$ and any $L \in \mathcal{S}(n, m, N)$ has the form of $\mathcal{G}(K, L) = \text{Tr}(P_{K, L} \Sigma_0)$, where $\Sigma_0 > 0$ is full-rank. This proves the coercivity of the inner-loop objective $\mathcal{G}(K, \cdot)$.

Moreover, for a fixed $K \in \mathcal{K}$, as $P_{K, L}$ is a polynomial of L , $\mathcal{G}(K, \cdot)$ is continuous in L . Combined with the coercivity property and the upper-boundedness of $\mathcal{G}(K, \cdot)$ for $K \in \mathcal{K}$, we can conclude the compactness of the superlevel set (3.8). The local Lipschitz and smoothness properties of $\mathcal{G}(K, \cdot)$ follow from Lemmas 4 and 5 of (Malik et al., 2020), with their A, B, Q, R, K matrices replaced by our $A - BK, D, Q + K^\top R^u K, -R^w, L$, respectively. Note that the property $\|\Sigma_K\| \leq \frac{c(K)}{\sigma_{\min}(Q)}$ in the proof of Lemma 16 in (Malik et al., 2020) does not hold in our setting as we only assume $Q \geq 0$ and $Q + K^\top R^u K$ may not be full-rank. Instead, we utilize the fact that $\mathcal{L}_K(a)$ is compact, and thus there exists a uniform constant $c_{\Sigma, a} := \max_{L \in \mathcal{L}_K(a)} \|\Sigma_{K, L}\|$ such that $\|\Sigma_{K, L}\| \leq c_{\Sigma, a}$ for all $L \in \mathcal{L}_K(a)$. Subsequently, we note that for all $L \in \mathcal{L}_K(a)$, the operator norms of the PG (3.9) and the Hessian, which are continuous functions of L , can be uniformly bounded by some constants $l_{K, a} > 0$ and $\psi_{K, a} > 0$, respectively. That is,

$$l_{K, a} := \max_{L \in \mathcal{L}_K(a)} \|\nabla_L \mathcal{G}(K, L)\|, \quad \psi_{K, a} := \max_{L \in \mathcal{L}_K(a)} \|\nabla_L^2 \mathcal{G}(K, L)\| = \max_{L \in \mathcal{L}_K(a)} \sup_{\|X\|_F=1} \|\nabla_L^2 \mathcal{G}(K, L)[X, X]\|,$$

where $\nabla_L^2 \mathcal{G}(K, L)[X, X]$ denotes the action of the Hessian on a matrix $X \in \mathcal{S}(n, m, N)$. The expression of $\nabla_L^2 \mathcal{G}(K, L)[X, X]$ follows from Proposition 3.10 of (Bu et al., 2019a), with their R, B, X, Y, A_K being replaced by our

$-R^w, D, P_{K,L}, \Sigma_{K,L}, A - BK - DL$, respectively. This proves the global Lipschitzness and smoothness of $\mathcal{G}(K, \cdot)$ over the compact superlevel set $\mathcal{L}_K(a)$.

To prove the PL condition, we first characterize the difference between $P_{K,L}$ and $P_{K,L(K)}$ as

$$\begin{aligned} P_{K,L(K)} - P_{K,L} &= A_{K,L(K)}^\top (P_{K,L(K)} - P_{K,L}) A_{K,L(K)} + (L(K) - L)^\top E_{K,L} \\ &\quad + E_{K,L}^\top (L(K) - L) - (L(K) - L)^\top H_{K,L} (L(K) - L), \end{aligned}$$

where $A_{K,L(K)} := A - BK - DL(K)$. By Lemma 3.3 and $K \in \mathcal{K}$, we can conclude that for all L , the inequality $H_{K,L} > 0$ holds because $P_{K,L(K)} \geq P_{K,L}$ for all $L \in \mathcal{S}(n, m, N)$. Then, by Proposition 2.1(b) of (Bu et al., 2019a), we have that for every $\varphi > 0$,

$$(L(K) - L)^\top E_{K,L} + E_{K,L}^\top (L(K) - L) \leq \frac{1}{\varphi} (L(K) - L)^\top (L(K) - L) + \varphi E_{K,L}^\top E_{K,L}.$$

Choosing $\varphi = 1/\lambda_{\min}(H_{K,L})$ yields

$$(L(K) - L)^\top E_{K,L} + E_{K,L}^\top (L(K) - L) - (L(K) - L)^\top H_{K,L} (L(K) - L) \leq \frac{E_{K,L}^\top E_{K,L}}{\lambda_{\min}(H_{K,L})} \leq \frac{E_{K,L}^\top E_{K,L}}{\lambda_{\min}(H_{K,L(K)})}, \quad (\text{B.8})$$

where the last inequality is due to $H_{K,L(K)} \leq H_{K,L}, \forall L \in \mathcal{S}(n, m, N)$. Let Y be the solution to $Y = A_{K,L(K)}^\top Y A_{K,L(K)} + \frac{E_{K,L}^\top E_{K,L}}{\lambda_{\min}(H_{K,L(K)})}$. From (B.8), we further know that $P_{K,L(K)} - P_{K,L} \leq Y$ and $Y = \sum_{t=0}^N [A_{K,L(K)}^\top]^t \frac{E_{K,L}^\top E_{K,L}}{\lambda_{\min}(H_{K,L(K)})} [A_{K,L(K)}]^t$. Therefore, letting $\Sigma_{K,L(K)} = \sum_{t=0}^N [A_{K,L(K)}]^t \Sigma_0 [A_{K,L(K)}]^\top$, we have

$$\begin{aligned} \mathcal{G}(K, L(K)) - \mathcal{G}(K, L) &= \text{Tr}(\Sigma_0(P_{K,L(K)} - P_{K,L})) \leq \text{Tr}(\Sigma_0 Y) \\ &= \text{Tr}\left(\Sigma_0 \left(\sum_{t=0}^N [A_{K,L(K)}^\top]^t \frac{E_{K,L}^\top E_{K,L}}{\lambda_{\min}(H_{K,L(K)})} [A_{K,L(K)}]^t\right)\right) \leq \frac{\|\Sigma_{K,L(K)}\|}{\lambda_{\min}(H_{K,L(K)})} \text{Tr}(E_{K,L}^\top E_{K,L}) \\ &\leq \frac{\|\Sigma_{K,L(K)}\|}{\lambda_{\min}(H_{K,L(K)}) \lambda_{\min}^2(\Sigma_{K,L})} \text{Tr}(\Sigma_{K,L}^\top E_{K,L}^\top E_{K,L} \Sigma_{K,L}) \leq \frac{\|\Sigma_{K,L(K)}\|}{4\lambda_{\min}(H_{K,L(K)}) \phi^2} \text{Tr}(\nabla_L \mathcal{G}(K, L)^\top \nabla_L \mathcal{G}(K, L)). \end{aligned} \quad (\text{B.9})$$

Inequality (B.9) follows from the cyclic property of matrix trace and the last inequality is due to $\Sigma_0 \leq \Sigma_{K,L}$. Hence,

$$\mathcal{G}(K, L(K)) - \mathcal{G}(K, L) \leq \frac{\|\Sigma_{K,L(K)}\|}{4\lambda_{\min}(H_{K,L(K)}) \phi^2} \text{Tr}(\nabla_L \mathcal{G}(K, L)^\top \nabla_L \mathcal{G}(K, L)).$$

Substituting in $\mu_K := 4\lambda_{\min}(H_{K,L(K)}) \phi^2 / \|\Sigma_{K,L(K)}\|$ completes the proof. \blacksquare

B.8. Proof of Lemma 3.5

Proof The nonconvexity proof is done by constructing a time-invariant example, which chooses $\Sigma_0 = I$ and $N = 5$. We further choose the system matrices to be $A_t = A, B_t = B, D_t = D, Q_t = Q, R_t^u = R^u$, and $R_t^w = R^w$ for all t , where

$$A = \begin{bmatrix} 1 & 0 & -5 \\ -1 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}, \quad B = \begin{bmatrix} 1 & -10 & 0 \\ 0 & 3 & 1 \\ -1 & 0 & 2 \end{bmatrix}, \quad D = \begin{bmatrix} 0.5 & 0 & 0 \\ 0 & 0.5 & 0 \\ 0 & 0 & 0.5 \end{bmatrix}, \quad Q = \begin{bmatrix} 3 & -1 & 0 \\ -1 & 2 & -1 \\ 0 & -1 & 1 \end{bmatrix}, \quad R^u = \begin{bmatrix} 2 & 1 & 1 \\ 1 & 3 & -1 \\ 1 & -1 & 3 \end{bmatrix},$$

and $R^w = 7.22543 \cdot I$. We also choose the time-invariant gain matrices $K_t^1 = K^1, K_t^2 = K^2$, and $K_t^3 = K^3$, where

$$K^1 = \begin{bmatrix} -0.8750 & 1.2500 & -2.5000 \\ -0.1875 & 0.1250 & 0.2500 \\ -0.4375 & 0.6250 & -0.7500 \end{bmatrix}, \quad K^2 = \begin{bmatrix} -0.8786 & 1.2407 & -2.4715 \\ -0.1878 & 0.1237 & 0.2548 \\ -0.4439 & 0.5820 & -0.7212 \end{bmatrix}, \quad K^3 = \frac{K^1 + K^2}{2}.$$

The concatenated matrices $A, B, D, Q, R^u, R^w, K^1, K^2, K^3$ are generated following the definitions in §3. Subsequently, we can prove that $K^1, K^2, K^3 \in \mathcal{K}$ by verifying that the recursive Riccati equation (3.6) yields p.s.d. solutions for

K^1, K^2, K^3 , respectively, and $\lambda_{\min}(H_{K^1, L(K^1)}) = 4.3496 \times 10^{-6}$, $\lambda_{\min}(H_{K^2, L(K^2)}) = 0.1844$, $\lambda_{\min}(H_{K^3, L(K^3)}) = 0.0926$. Then, we can compute that $(\mathcal{G}(K^1, L(K^1)) + \mathcal{G}(K^2, L(K^2)))/2 - \mathcal{G}(K^3, L(K^3)) = -0.0224 < 0$. Thus, we can conclude that $\mathcal{G}(K, L(K))$ is nonconvex in K . Subsequently, we show that the outer loop is noncoercive on \mathcal{K} by a scalar example with $N = 2$ and time-invariant system matrices being $A_t = 2$, $B_t = D_t = Q_t = R_t^u = 1$, and $R_t^w = 5$, for all t . Then, we consider the gain matrix $K^\epsilon = \begin{bmatrix} 2 - \epsilon & 0 & 0 \\ 0 & 2 - \epsilon & 0 \end{bmatrix} \in \mathcal{K}$, $\forall \epsilon \in (0, \frac{16}{9})$. It can be observed that $\lim_{\epsilon \rightarrow 0+} \mathcal{G}(K^\epsilon, L(K^\epsilon)) = 11 < \infty$, while $\lim_{\epsilon \rightarrow 0+} K^\epsilon \in \partial\mathcal{K}$, for $\epsilon \in (0, \frac{16}{9})$. Therefore, $\mathcal{G}(K, L(K))$ is not coercive. Lastly, by Lemma 3.2, the stationary point of the outer loop in \mathcal{K} , denoted as $(K^*, L(K^*))$, is unique and constitutes the unique Nash equilibrium of the game. ■

B.9. Proof of Theorem 3.6

Proof We prove the global convergence of three inner-loop PG updates as follows:

PG: For a fixed $K \in \mathcal{K}$ and an arbitrary L_0 that induces a finite $\mathcal{G}(K, L_0)$, we first define superlevel sets $\mathcal{L}_K(\mathcal{G}(K, L_0))$ and $\mathcal{L}_K(a)$ as in (3.8), where $a < \mathcal{G}(K, L_0)$ is an arbitrary constant. Clearly, it holds that $L_0 \in \mathcal{L}_K(\mathcal{G}(K, L_0)) \subset \mathcal{L}_K(a)$ and thus $\mathcal{L}_K(\mathcal{G}(K, L_0))$ is nonempty as well as compact (as shown before in Lemma 3.4). Next, denote the closure of the complement of $\mathcal{L}_K(a)$ as $(\mathcal{L}_K(a))^c$, which is again nonempty and also disjoint with $\mathcal{L}_K(\mathcal{G}(K, L_0))$, i.e., $\mathcal{L}_K(\mathcal{G}(K, L_0)) \cap (\mathcal{L}_K(a))^c = \emptyset$, due to a being strictly less than $\mathcal{G}(K, L_0)$. Hence, one can deduce (see for example Lemma A.1 of (Berg, 2001)) that, there exists a Hausdorff distance $\delta_a > 0$ between $\mathcal{L}_K(\mathcal{G}(K, L_0))$ and $(\mathcal{L}_K(a))^c$ such that for a given $L \in \mathcal{L}_K(\mathcal{G}(K, L_0))$, all L' satisfying $\|L' - L\|_F \leq \delta_a$ also satisfy $L' \in \mathcal{L}_K(a)$.

Now, since $L_0 \in \mathcal{L}_K(\mathcal{G}(K, L_0))$, we have $\|\nabla_L \mathcal{G}(K, L_0)\|_F \leq l_{K, \mathcal{G}(K, L_0)}$, where $l_{K, \mathcal{G}(K, L_0)}$ is the global Lipschitz constant over $\mathcal{L}_K(\mathcal{G}(K, L_0))$. Therefore, it suffices to choose $\eta_0 \leq \frac{\delta_a}{l_{K, \mathcal{G}(K, L_0)}}$ to ensure that the one-step “fictitious” PG update satisfies $L_0 + \eta_0 \nabla_L \mathcal{G}(K, L_0) \in \mathcal{L}_K(a)$. By Lemma 3.4, we can apply the descent lemma (for minimization problems) (Boyd & Vandenberghe, 2004) to derive:

$$\mathcal{G}(K, L_0) - \mathcal{G}(K, L_0 + \eta_0 \nabla_L \mathcal{G}(K, L_0)) \leq -\eta_0 \langle \nabla_L \mathcal{G}(K, L_0), \nabla_L \mathcal{G}(K, L_0) \rangle + \frac{\eta_0^2 \psi_{K,a}}{2} \|\nabla_L \mathcal{G}(K, L_0)\|_F^2.$$

Thus, we can additionally require $\eta_0 \leq 1/\psi_{K,a}$ to guarantee that the objective is non-decreasing (i.e., $L_0 + \eta_0 \nabla_L \mathcal{G}(K, L_0) \in \mathcal{L}_K(\mathcal{G}(K, L_0))$). This also implies that starting from $L_0 + \eta_0 \nabla_L \mathcal{G}(K, L_0)$, taking another “fictitious” PG update step of $\eta_0 \nabla_L \mathcal{G}(K, L_0)$ with $\eta_0 \leq \min\{\delta_a/l_{K, \mathcal{G}(K, L_0)}, 1/\psi_{K,a}\}$ ensures that $L_0 + 2\eta_0 \nabla_L \mathcal{G}(K, L_0) \in \mathcal{L}_K(\mathcal{G}(K, L_0))$. Applying the same argument iteratively certifies that we can actually take a much larger stepsize $\tilde{\eta}_0$, with the only requirement being $\tilde{\eta}_0 \leq 1/\psi_{K,a}$, which guarantees that after one-step of the “actual” PG update, $L_1 = L_0 + \tilde{\eta}_0 \nabla_L \mathcal{G}(K, L_0)$, it holds that $L_1 \in \mathcal{L}_K(\mathcal{G}(K, L_0))$. This is because $1/\psi_{K,a}$ can be covered by a finite times of $\delta_a/l_{K, \mathcal{G}(K, L_0)} > 0$. Now, since $L_1 \in \mathcal{L}_K(\mathcal{G}(K, L_0))$, we can repeat the above arguments for all future iterations and show that with a fixed stepsize η satisfying $\eta \leq 1/\psi_{K,a}$, the iterates of the PG update (3.9) satisfies for all $l \geq 0$ that

$$\mathcal{G}(K, L_l) - \mathcal{G}(K, L_{l+1}) \leq -\eta \langle \nabla_L \mathcal{G}(K, L_l), \nabla_L \mathcal{G}(K, L_l) \rangle + \frac{\eta^2 \psi_{K,a}}{2} \|\nabla_L \mathcal{G}(K, L_l)\|_F^2 \leq 0.$$

By the choice of η , we characterize the convergence rate of the PG update (3.9) by exploiting the PL condition from Lemma 3.4, such that

$$\mathcal{G}(K, L_l) - \mathcal{G}(K, L_{l+1}) \leq -\frac{1}{2\psi_{K,a}} \|\nabla_L \mathcal{G}(K, L_l)\|_F^2 \leq -\frac{\mu_K}{2\psi_{K,a}} (\mathcal{G}(K, L(K)) - \mathcal{G}(K, L_l)),$$

where μ_K is the global PL constant for a given $K \in \mathcal{K}$. Thus, $\mathcal{G}(K, L(K)) - \mathcal{G}(K, L_{l+1}) \leq (1 - \frac{\mu_K}{2\psi_{K,a}}) (\mathcal{G}(K, L(K)) - \mathcal{G}(K, L_l))$. This completes the proof for the linear convergence of the objective. Next, we show the convergence of the

control gain matrix to $L(K)$. We let $q := 1 - \frac{\mu_K}{2\psi_{K,a}}$ and present the following comparison

$$\begin{aligned} \mathcal{G}(K, L(K)) - \mathcal{G}(K, L_l) &= \text{Tr}[(P_{K,L(K)} - P_{K,L_l})\Sigma_0] \\ &= \text{Tr}\left[\sum_{t=0}^N (A_{K,L_l}^\top)^t ((L(K) - L_l)^\top H_{K,L(K)} (L(K) - L_l)) (A_{K,L_l})^t \Sigma_0\right] \\ &= \text{Tr}\left[(L(K) - L_l)^\top H_{K,L(K)} (L(K) - L_l) \Sigma_{K,L_l}\right] \\ &\geq \phi \cdot \lambda_{\min}(H_{K,L(K)}) \cdot \|L(K) - L_l\|_F^2, \end{aligned}$$

where the second equality is due to that $E_{K,L(K)} = \mathbf{0}$ from the property of stationary points, and the last inequality is due to $\Sigma_{K,L_l} \geq \Sigma_0, \forall l \geq 0$. Then, we can conclude that

$$\begin{aligned} \|L(K) - L_l\|_F &\leq \sqrt{\lambda_{\min}^{-1}(H_{K,L(K)}) \cdot (\mathcal{G}(K, L(K)) - \mathcal{G}(K, L_l))} \\ &\leq q^{\frac{l}{2}} \cdot \sqrt{\lambda_{\min}^{-1}(H_{K,L(K)}) \cdot (\mathcal{G}(K, L(K)) - \mathcal{G}(K, L_0))}. \end{aligned} \quad (\text{B.10})$$

This completes the convergence proof of the gain matrix.

NPG: We argue that for a fixed $K \in \mathcal{K}$, the inner-loop iterates following the NPG update (3.10) with a certain stepsize choice satisfy $P_{K,L_{l+1}} \geq P_{K,L_l}$ in the p.s.d. sense, for all $l \geq 0$. By the definition of $P_{K,L}$ in (3.2), we can derive the following comparison between P_{K,L_l} and $P_{K,L_{l+1}}$:

$$\begin{aligned} P_{K,L_{l+1}} - P_{K,L_l} &= A_{K,L_{l+1}}^\top (P_{K,L_{l+1}} - P_{K,L_l}) A_{K,L_{l+1}} + 4\eta E_{K,L_l}^\top E_{K,L_l} - 4\eta^2 E_{K,L_l}^\top H_{K,L_l} E_{K,L_l} \\ &= A_{K,L_{l+1}}^\top (P_{K,L_{l+1}} - P_{K,L_l}) A_{K,L_{l+1}} + 4\eta E_{K,L_l}^\top (I - \eta H_{K,L_l}) E_{K,L_l}, \end{aligned}$$

where $A_{K,L} := A - BK - DL$. As a result, we can require $\eta \leq \frac{1}{\|H_{K,L_l}\|}$ to guarantee that $P_{K,L_{l+1}} \geq P_{K,L_l}$. Moreover, Lemma 3.3 suggests that for a fixed $K \in \mathcal{K}$, it holds that $P_{K,L} \leq P_{K,L(K)}$ in the p.s.d. sense for all L . Thus, if we require $\eta \leq \min_l \frac{1}{\|H_{K,L_l}\|}$, then $\{P_{K,L_l}\}_{l \geq 0}$ constitutes a monotonically non-decreasing sequence in the p.s.d. sense. Moreover, since $K \in \mathcal{K}$ and $\{P_{K,L_l}\}_{l \geq 0}$ monotonically non-decreasing, the sequence $\{H_{K,L_l}\}_{l \geq 0}$ is monotonically non-increasing (in the p.s.d. sense) and each H_{K,L_l} is both symmetric and positive-definite. Therefore, we can choose a uniform stepsize $\eta = \frac{1}{2\|H_{K,L_0}\|} \leq \min_l \frac{1}{\|H_{K,L_l}\|}$ and show the convergence rate of the NPG update (3.10) as follows:

$$\begin{aligned} \mathcal{G}(K, L_{l+1}) - \mathcal{G}(K, L_l) &= \text{Tr}((P_{K,L_{l+1}} - P_{K,L_l})\Sigma_0) \geq \text{Tr}((4\eta E_{K,L_l}^\top E_{K,L_l} - 4\eta^2 E_{K,L_l}^\top H_{K,L_l} E_{K,L_l})\Sigma_0) \\ &\geq \frac{\phi}{\|H_{K,L_0}\|} \text{Tr}(E_{K,L_l}^\top E_{K,L_l}) \geq \frac{\phi \lambda_{\min}(H_{K,L(K)})}{\|H_{K,L_0}\| \|\Sigma_{K,L(K)}\|} (\mathcal{G}(K, L(K)) - \mathcal{G}(K, L_l)), \end{aligned}$$

where the last inequality follows from (B.9). Thus,

$$\mathcal{G}(K, L(K)) - \mathcal{G}(K, L_{l+1}) \leq \left(1 - \frac{\phi \lambda_{\min}(H_{K,L(K)})}{\|H_{K,L_0}\| \|\Sigma_{K,L(K)}\|}\right) (\mathcal{G}(K, L(K)) - \mathcal{G}(K, L_l)).$$

This implies the globally linear convergence of the objective. The convergence proof of the gain matrix is similar to the one presented for the vanilla PG update. The only difference is to have $q := 1 - \frac{\phi \lambda_{\min}(H_{K,L(K)})}{\|H_{K,L_0}\| \|\Sigma_{K,L(K)}\|}$ instead.

GN: Similar to the proof for NPG, we first find a stepsize such that $\{P_{K,L_l}\}_{l \geq 0}$ constitutes a monotonically non-decreasing sequence bounded above by $P_{K,L(K)}$ in the p.s.d. sense (based on Lemma 3.3). Taking the difference between P_{K,L_l} and $P_{K,L_{l+1}}$ and substituting in (3.12) yield

$$\begin{aligned} P_{K,L_{l+1}} - P_{K,L_l} &= A_{K,L_{l+1}}^\top (P_{K,L_{l+1}} - P_{K,L_l}) A_{K,L_{l+1}} + 4\eta E_{K,L_l}^\top H_{K,L_l}^{-1} E_{K,L_l} - 4\eta^2 E_{K,L_l}^\top H_{K,L_l}^{-1} E_{K,L_l} \\ &= A_{K,L_{l+1}}^\top (P_{K,L_{l+1}} - P_{K,L_l}) A_{K,L_{l+1}} + 4\eta E_{K,L_l}^\top (H_{K,L_l}^{-1} - \eta H_{K,L_l}^{-1}) E_{K,L_l}. \end{aligned}$$

Therefore, $P_{K,L_{l+1}} \geq P_{K,L_l}$ can be ensured by choosing $\eta \leq 1$. Subsequently, we characterize the convergence rate of the GN update with $\eta \leq 1/2$ as follows

$$\begin{aligned} \mathcal{G}(K, L_{l+1}) - \mathcal{G}(K, L_l) &= \text{Tr}((P_{K,L_{l+1}} - P_{K,L_l})\Sigma_0) \geq \text{Tr}((4\eta E_{K,L_l}^\top H_{K,L_l}^{-1} E_{K,L_l} - 4\eta^2 E_{K,L_l}^\top H_{K,L_l}^{-1} E_{K,L_l})\Sigma_0) \\ &\geq \frac{\phi}{\|H_{K,L_l}\|} \text{Tr}(E_{K,L_l}^\top E_{K,L_l}) \geq \frac{\phi \lambda_{\min}(H_{K,L(K)})}{\|H_{K,L_0}\| \|\Sigma_{K,L(K)}\|} (\mathcal{G}(K, L(K)) - \mathcal{G}(K, L_l)), \end{aligned}$$

where the last inequality follows from (B.9) and the fact that $\{H_{K,L_l}\}_{l \geq 0}$ is a monotonically non-increasing sequence lower bounded by $H_{K,L(K)} > 0$, in the p.s.d. sense. Thus,

$$\mathcal{G}(K, L(K)) - \mathcal{G}(K, L_{l+1}) \leq \left(1 - \frac{\phi \lambda_{\min}(H_{K,L(K)})}{\|H_{K,L_0}\| \|\Sigma_{K,L(K)}\|}\right) (\mathcal{G}(K, L(K)) - \mathcal{G}(K, L_l))$$

and the globally linear convergence rate of the objective is proved. Lastly, the locally Q-quadratic convergence rate for the GN update directly follows from (Bu et al., 2019a; Bu & Mesbahi, 2020) and the convergence proof of the gain matrix is the same as the one for the NPG update. This completes the proof. \blacksquare

B.10. Proof of Theorem 3.7

Proof We first introduce the following cost difference lemma for the outer-loop problem, whose proof is deferred to §C.1.

Lemma B.1 (*Cost Difference Lemma for Outer-Loop*) Suppose that for two sequences of control gain matrices $\{K_t\}$ and $\{K'_t\}$ and a given $t \in \{0, \dots, N-1\}$, there exist p.s.d. solutions to (3.1) at time $t+1$, denoted as $P_{K_{t+1}, L(K_{t+1})}$ and $P_{K'_{t+1}, L(K'_{t+1})}$, respectively. Also, suppose that the following inequalities are satisfied:

$$R_t^w - D_t^\top P_{K_{t+1}, L(K_{t+1})} D_t > 0, \quad R_t^w - D_t^\top P_{K'_{t+1}, L(K'_{t+1})} D_t > 0.$$

Then, there exist p.s.d. solutions to (3.1) at time t , denoted as $P_{K_t, L(K_t)}$ and $P_{K'_t, L(K'_t)}$, and their difference can be quantified as

$$\begin{aligned} P_{K'_t, L(K'_t)} - P_{K_t, L(K_t)} &= A_{K'_t, L(K'_t)}^\top (P_{K'_{t+1}, L(K'_{t+1})} - P_{K_{t+1}, L(K_{t+1})}) A_{K'_t, L(K'_t)} \\ &\quad + \mathcal{R}_{K_t, K'_t} - \Xi_{K_t, K'_t}^\top (R_t^w - D_t^\top P_{K_{t+1}, L(K_{t+1})} D_t)^{-1} \Xi_{K_t, K'_t}, \end{aligned} \quad (\text{B.11})$$

where

$$\begin{aligned} \tilde{P}_{K_{t+1}, L(K_{t+1})} &:= P_{K_{t+1}, L(K_{t+1})} + P_{K_{t+1}, L(K_{t+1})} D_t (R_t^w - D_t^\top P_{K_{t+1}, L(K_{t+1})} D_t)^{-1} D_t^\top P_{K_{t+1}, L(K_{t+1})} \quad (\text{B.12}) \\ F_{K_t, L(K_t)} &:= (R_t^u + B_t^\top \tilde{P}_{K_{t+1}, L(K_{t+1})} B_t) K_t - B_t^\top \tilde{P}_{K_{t+1}, L(K_{t+1})} A_t \\ L(K_t) &:= (-R_t^w + D_t^\top P_{K_{t+1}, L(K_{t+1})} D_t)^{-1} D_t^\top P_{K_{t+1}, L(K_{t+1})} (A_t - B_t K_t) \\ \Xi_{K_t, K'_t} &:= -(R_t^w - D_t^\top P_{K_{t+1}, L(K_{t+1})} D_t) L(K'_t) - D_t^\top P_{K_{t+1}, L(K_{t+1})} (A_t - B_t K'_t) \\ \mathcal{R}_{K_t, K'_t} &:= (A_t - B_t K'_t)^\top \tilde{P}_{K_{t+1}, L(K_{t+1})} (A_t - B_t K'_t) - P_{K_t, L(K_t)} + Q_t + (K'_t)^\top R_t^u (K'_t) \\ &= (K'_t - K_t)^\top F_{K_t, L(K_t)} + F_{K_t, L(K_t)}^\top (K'_t - K_t) + (K'_t - K_t)^\top (R_t^u + B_t^\top \tilde{P}_{K_{t+1}, L(K_{t+1})} B_t) (K'_t - K_t). \end{aligned} \quad (\text{B.13})$$

Subsequently, we prove that starting from any $K \in \mathcal{K}$, the next iterate K' is guaranteed to satisfy $K' \in \mathcal{K}$ following the NPG update (3.11) or the GN update (3.13) with proper stepsizes. Note that (3.6) is the compact form of the following recursive Riccati equation

$$P_{K_t, L(K_t)} = Q_t + K_t^\top R_t^u K_t + (A_t - B_t K_t)^\top \tilde{P}_{K_{t+1}, L(K_{t+1})} (A_t - B_t K_t), \quad t \in \{0, \dots, N-1\} \quad (\text{B.14})$$

where $\tilde{P}_{K_{t+1}, L(K_{t+1})}$ is as defined in (B.12) and $P_{K_N, L(K_N)} = Q_N$. Thus, $K \in \mathcal{K}$ is equivalent to that (B.14) admits a solution $P_{K_{t+1}, L(K_{t+1})} \geq 0$ and $R_t^w - D_t^\top P_{K_{t+1}, L(K_{t+1})} D_t > 0$, for all $t \in \{0, \dots, N-1\}$. Since $K \in \mathcal{K}$ and $P_{K_N, L(K_N)} = P_{K'_N, L(K'_N)} = Q_N$ for any K' , we have

$$R_{N-1}^w - D_{N-1}^\top P_{K'_N, L(K'_N)} D_{N-1} = R_{N-1}^w - D_{N-1}^\top P_{K_N, L(K_N)} D_{N-1} = R_{N-1}^w - D_{N-1}^\top Q_N D_{N-1} > 0.$$

That is, $R_{N-1}^w - D_{N-1}^\top P_{K'_{N-1}, L(K'_{N-1})} D_{N-1}$ is invertible and a solution to (B.14) with $t = N - 1$, denoted as $P_{K'_{N-1}, L(K'_{N-1})}$, exists and is both p.s.d. and unique.

Subsequently, we invoke Lemma B.1 to find a stepsize such that a solution to (B.14) with $t = N - 2$ also exists after one-step NPG update (3.11). By (B.14), it suffices to ensure $R_{N-2}^w - D_{N-2}^\top P_{K'_{N-1}, L(K'_{N-1})} D_{N-2} > 0$ (thus invertible) in order to guarantee the existence of $P_{K'_{N-2}, L(K'_{N-2})}$. By Lemma B.1 and $P_{K_N, L(K_N)} = P_{K'_{N-1}, L(K'_{N-1})} = Q_N$, we have

$$\begin{aligned} P_{K'_{N-1}, L(K'_{N-1})} - P_{K_{N-1}, L(K_{N-1})} &= \mathcal{R}_{K_{N-1}, K'_{N-1}} - \Xi_{K_{N-1}, K'_{N-1}}^\top (R_{N-1}^w - D_{N-1}^\top P_{K_N, L(K_N)} D_{N-1})^{-1} \Xi_{K_{N-1}, K'_{N-1}} \\ &\leq \mathcal{R}_{K_{N-1}, K'_{N-1}}, \end{aligned}$$

where the last inequality is due to the fact that $R_{N-1}^w - D_{N-1}^\top P_{K_N, L(K_N)} D_{N-1} > 0$. Then, we substitute the NPG update rule $K'_{N-1} = K_{N-1} - 2\alpha_{K_{N-1}} F_{K_{N-1}, L(K_{N-1})}$ for $t = N - 1$ into (B.13) to get

$$\begin{aligned} P_{K'_{N-1}, L(K'_{N-1})} - P_{K_{N-1}, L(K_{N-1})} &\leq (K'_{N-1} - K_{N-1})^\top F_{K_{N-1}, L(K_{N-1})} + F_{K_{N-1}, L(K_{N-1})}^\top (K'_{N-1} - K_{N-1}) \\ &\quad + (K'_{N-1} - K_{N-1})^\top (R_{N-1}^u + B_{N-1}^\top \tilde{P}_{K_N, L(K_N)} B_{N-1}) (K'_{N-1} - K_{N-1}) \\ &= -4\alpha_{K_{N-1}} F_{K_{N-1}, L(K_{N-1})}^\top F_{K_{N-1}, L(K_{N-1})} + 4\alpha_{K_{N-1}}^2 F_{K_{N-1}, L(K_{N-1})}^\top (R_{N-1}^u + B_{N-1}^\top \tilde{P}_{K_N, L(K_N)} B_{N-1}) F_{K_{N-1}, L(K_{N-1})} \\ &= -4\alpha_{K_{N-1}} F_{K_{N-1}, L(K_{N-1})}^\top (\mathbf{I} - \alpha_{K_{N-1}} (R_{N-1}^u + B_{N-1}^\top \tilde{P}_{K_N, L(K_N)} B_{N-1})) F_{K_{N-1}, L(K_{N-1})}. \end{aligned}$$

Therefore, choosing $\alpha_{K_{N-1}} \in [0, 1/\|R_{N-1}^u + B_{N-1}^\top \tilde{P}_{K_N, L(K_N)} B_{N-1}\|]$ suffices to ensure that $P_{K'_{N-1}, L(K'_{N-1})} \leq P_{K_{N-1}, L(K_{N-1})}$. Hence,

$$R_{N-2}^w - D_{N-2}^\top P_{K'_{N-1}, L(K'_{N-1})} D_{N-2} \geq R_{N-2}^w - D_{N-2}^\top P_{K_{N-1}, L(K_{N-1})} D_{N-2} > 0,$$

where the last inequality comes from $\mathbf{K} \in \mathcal{K}$. Therefore, the existence of $P_{K'_{N-2}, L(K'_{N-2})} \geq 0$ is proved if the above requirement on $\alpha_{K_{N-1}}$ is satisfied. We can apply this argument iteratively backward since for all $t \in \{1, \dots, N\}$, $P_{K'_t, L(K'_t)} \leq P_{K_t, L(K_t)}$ implies $R_{t-1}^w - D_{t-1}^\top P_{K'_t, L(K'_t)} D_{t-1} \geq R_{t-1}^w - D_{t-1}^\top P_{K_t, L(K_t)} D_{t-1} > 0$ and thus Lemma B.1 can be applied for all t . Moreover, $P_{K'_t, L(K'_t)} \leq P_{K_t, L(K_t)}$ is guaranteed to hold for all $t \in \{1, \dots, N\}$ if we require the stepsize of the NPG update to satisfy

$$\alpha_{K_t} \in [0, 1/\|R_t^u + B_t^\top \tilde{P}_{K_{t+1}, L(K_{t+1})} B_t\|], \quad \forall t \in \{0, \dots, N-1\}.$$

Equivalently, the above conditions can also be represented using the compact forms introduced in §3. In particular, $\mathbf{G}_{\mathbf{K}, L(\mathbf{K})}$ is a concatenated matrix with blocks of $R_t^u + B_t^\top \tilde{P}_{K_{t+1}, L(K_{t+1})} B_t$, $t \in \{0, \dots, N-1\}$, on the diagonal. Thus, we have $\|\mathbf{G}_{\mathbf{K}, L(\mathbf{K})}\| \geq \|R_t^u + B_t^\top \tilde{P}_{K_{t+1}, L(K_{t+1})} B_t\|$, for all $t \in \{0, \dots, N-1\}$. Adopting the compact notations, we have for any $\mathbf{K} \in \mathcal{K}$, if the stepsize $\alpha_{\mathbf{K}} \in [0, 1/\|\mathbf{G}_{\mathbf{K}, L(\mathbf{K})}\|]$, then

$$\begin{aligned} P_{\mathbf{K}', L(\mathbf{K}')} - P_{\mathbf{K}, L(\mathbf{K})} &\leq (\mathbf{K}' - \mathbf{K})^\top \mathbf{F}_{\mathbf{K}, L(\mathbf{K})} + \mathbf{F}_{\mathbf{K}, L(\mathbf{K})}^\top (\mathbf{K}' - \mathbf{K}) + (\mathbf{K}' - \mathbf{K})^\top \mathbf{G}_{\mathbf{K}, L(\mathbf{K})} (\mathbf{K}' - \mathbf{K}) \\ &= -4\alpha_{\mathbf{K}} \mathbf{F}_{\mathbf{K}, L(\mathbf{K})}^\top \mathbf{F}_{\mathbf{K}, L(\mathbf{K})} + 4\alpha_{\mathbf{K}}^2 \mathbf{F}_{\mathbf{K}, L(\mathbf{K})}^\top \mathbf{G}_{\mathbf{K}, L(\mathbf{K})} \mathbf{F}_{\mathbf{K}, L(\mathbf{K})} \\ &= -4\alpha_{\mathbf{K}} \mathbf{F}_{\mathbf{K}, L(\mathbf{K})}^\top (\mathbf{I} - \alpha_{\mathbf{K}} \mathbf{G}_{\mathbf{K}, L(\mathbf{K})}) \mathbf{F}_{\mathbf{K}, L(\mathbf{K})} \leq 0. \end{aligned} \tag{B.15}$$

Since we have already shown that $P_{K'_t, L(K'_t)} \geq 0$ exists for all $t \in \{0, \dots, N-1\}$ and $\mathbf{P}_{\mathbf{K}', L(\mathbf{K}')}$ is a concatenation of blocks of $P_{K'_t, L(K'_t)}$, we can conclude that $\mathbf{K}' \in \mathcal{K}$. Applying the above analysis iteratively proves that the sequence $\{\mathbf{P}_{\mathbf{K}_k, L(\mathbf{K}_k)}\}_{k \geq 0}$ following the NPG update (3.11) with a given $\mathbf{K}_0 \in \mathcal{K}$ and a constant stepsize $\alpha \leq \min_{k \geq 0} \frac{1}{\|\mathbf{G}_{\mathbf{K}_k, L(\mathbf{K}_k)}\|}$ is monotonically non-increasing in the p.s.d. sense, and thus $\mathbf{K}_k \in \mathcal{K}$ for all $k \geq 0$ given a $\mathbf{K}_0 \in \mathcal{K}$. Furthermore, we can choose $\alpha \leq \frac{1}{\|\mathbf{G}_{\mathbf{K}_0, L(\mathbf{K}_0)}\|} \leq \min_{k \geq 0} \frac{1}{\|\mathbf{G}_{\mathbf{K}_k, L(\mathbf{K}_k)}\|}$ because $\{\mathbf{G}_{\mathbf{K}_k, L(\mathbf{K}_k)}\}$ is monotonically non-increasing and each $\mathbf{G}_{\mathbf{K}_k, L(\mathbf{K}_k)}$ is symmetric and positive-definite. This completes the proof for the NPG update (3.11).

For the GN update (3.13), all our arguments for the proof above still hold but instead we invoke Lemma B.1 with the recursive update rule being $K'_t = K_t - 2\alpha_{K_t} (R_t^u + B_t^\top \tilde{P}_{K_{t+1}, L(K_{t+1})} B_t)^{-1} F_{K_t, L(K_t)}$, for all $t \in \{0, \dots, N-1\}$. In particular, we have the matrix difference at time $t = N-1$, denoted as $P_{K'_{N-1}, L(K'_{N-1})} - P_{K_{N-1}, L(K_{N-1})}$, being

$$\begin{aligned} P_{K'_{N-1}, L(K'_{N-1})} - P_{K_{N-1}, L(K_{N-1})} &\leq (K'_{N-1} - K_{N-1})^\top F_{K_{N-1}, L(K_{N-1})} + F_{K_{N-1}, L(K_{N-1})}^\top (K'_{N-1} - K_{N-1}) \\ &\quad + (K'_{N-1} - K_{N-1})^\top (R_{N-1}^u + B_{N-1}^\top \tilde{P}_{K_N, L(K_N)} B_{N-1}) (K'_{N-1} - K_{N-1}) \\ &= -4\alpha_{K_{N-1}} F_{K_{N-1}, L(K_{N-1})}^\top ((1 - \alpha_{K_{N-1}}) \cdot (R_{N-1}^u + B_{N-1}^\top \tilde{P}_{K_N, L(K_N)} B_{N-1})^{-1}) F_{K_{N-1}, L(K_{N-1})}. \end{aligned}$$

Therefore, choosing $\alpha_{K_{N-1}} \in [0, 1]$ suffices to ensure that $P_{K'_{N-1}, L(K'_{N-1})} \leq P_{K_{N-1}, L(K_{N-1})}$. Applying the iterative arguments for the proof of the NPG update and using the compact notations yields that for any fixed $\mathbf{K} \in \mathcal{K}$, if the stepsize satisfies $\alpha_{\mathbf{K}} \in [0, 1]$, then $P_{\mathbf{K}', L(\mathbf{K}')} \geq 0$ exists and

$$\begin{aligned} P_{\mathbf{K}', L(\mathbf{K}')} - P_{\mathbf{K}, L(\mathbf{K})} &\leq (\mathbf{K}' - \mathbf{K})^\top \mathbf{F}_{\mathbf{K}, L(\mathbf{K})} + \mathbf{F}_{\mathbf{K}, L(\mathbf{K})}^\top (\mathbf{K}' - \mathbf{K}) + (\mathbf{K}' - \mathbf{K})^\top \mathbf{G}_{\mathbf{K}, L(\mathbf{K})} (\mathbf{K}' - \mathbf{K}) \\ &= -4\alpha_{\mathbf{K}} \mathbf{F}_{\mathbf{K}, L(\mathbf{K})}^\top \mathbf{G}_{\mathbf{K}, L(\mathbf{K})}^{-1} \mathbf{F}_{\mathbf{K}, L(\mathbf{K})} + 4\alpha_{\mathbf{K}}^2 \mathbf{F}_{\mathbf{K}, L(\mathbf{K})}^\top \mathbf{G}_{\mathbf{K}, L(\mathbf{K})}^{-1} \mathbf{F}_{\mathbf{K}, L(\mathbf{K})} \\ &= -4\alpha_{\mathbf{K}} \mathbf{F}_{\mathbf{K}, L(\mathbf{K})}^\top ((1 - \alpha_{\mathbf{K}}) \cdot \mathbf{G}_{\mathbf{K}, L(\mathbf{K})}^{-1}) \mathbf{F}_{\mathbf{K}, L(\mathbf{K})} \leq 0. \end{aligned} \quad (\text{B.16})$$

This proves that $\mathbf{K}' \in \mathcal{K}$. Now, we can apply the above analysis iteratively to prove that the sequence $\{P_{\mathbf{K}_k, L(\mathbf{K}_k)}\}_{k \geq 0}$ following the GN update (3.13) with a given $\mathbf{K}_0 \in \mathcal{K}$ and a constant stepsize $\alpha \in [0, 1]$ is monotonically non-increasing in the p.s.d. sense. Thus, we have $\mathbf{K}_k \in \mathcal{K}$ for all $k \geq 0$ given a $\mathbf{K}_0 \in \mathcal{K}$. This completes the proof. ■

B.11. Proof of Theorem 3.8

Proof For the NPG update (3.11), let $\alpha \leq \frac{1}{2\|\mathbf{G}_{\mathbf{K}_0, L(\mathbf{K}_0)}\|}$ and suppose $\mathbf{K}_0 \in \mathcal{K}$. Theorem 3.7 suggests that $P_{\mathbf{K}_k, L(\mathbf{K}_k)} \geq P_{\mathbf{K}^*, L(\mathbf{K}^*)} \geq 0$ exists for all $k \geq 0$ and the sequence $\{P_{\mathbf{K}_k, L(\mathbf{K}_k)}\}_{k \geq 0}$ is non-increasing in the p.s.d. sense. As a result, we have by (B.15) that

$$\text{Tr}((P_{\mathbf{K}_{k+1}, L(\mathbf{K}_{k+1})} - P_{\mathbf{K}_k, L(\mathbf{K}_k)})\Sigma_0) \leq -\frac{\phi}{\|\mathbf{G}_{\mathbf{K}_0, L(\mathbf{K}_0)}\|} \text{Tr}(\mathbf{F}_{\mathbf{K}_k, L(\mathbf{K}_k)}^\top \mathbf{F}_{\mathbf{K}_k, L(\mathbf{K}_k)}) \leq -2\alpha\phi \text{Tr}(\mathbf{F}_{\mathbf{K}_k, L(\mathbf{K}_k)}^\top \mathbf{F}_{\mathbf{K}_k, L(\mathbf{K}_k)}). \quad (\text{B.17})$$

Therefore, from iterations $k = 0$ to $K - 1$, summing over both sides of (B.17) and dividing by K yield,

$$\frac{1}{K} \sum_{k=0}^{K-1} \text{Tr}(\mathbf{F}_{\mathbf{K}_k, L(\mathbf{K}_k)}^\top \mathbf{F}_{\mathbf{K}_k, L(\mathbf{K}_k)}) \leq \frac{\text{Tr}(P_{\mathbf{K}_0, L(\mathbf{K}_0)} - P_{\mathbf{K}^*, L(\mathbf{K}^*)})}{2\alpha\phi \cdot K}. \quad (\text{B.18})$$

Namely, the sequence of natural gradient norm square $\{\text{Tr}(\mathbf{F}_{\mathbf{K}_k, L(\mathbf{K}_k)}^\top \mathbf{F}_{\mathbf{K}_k, L(\mathbf{K}_k)})\}_{k \geq 0}$ converges on average with a globally $\mathcal{O}(1/K)$ rate. The convergent stationary point is also the unique Nash equilibrium of the game, by Lemma 3.5. This completes the convergence proof of the NPG update. Similarly, for the GN update with $\alpha \leq 1/2$ and suppose $\mathbf{K}_0 \in \mathcal{K}$, we can obtain from (B.16) that

$$\text{Tr}((P_{\mathbf{K}_{k+1}, L(\mathbf{K}_{k+1})} - P_{\mathbf{K}_k, L(\mathbf{K}_k)})\Sigma_0) \leq -\frac{2\alpha\phi}{\|\mathbf{G}_{\mathbf{K}_0, L(\mathbf{K}_0)}\|} \text{Tr}(\mathbf{F}_{\mathbf{K}_k, L(\mathbf{K}_k)}^\top \mathbf{F}_{\mathbf{K}_k, L(\mathbf{K}_k)}). \quad (\text{B.19})$$

As before, we sum up (B.19) from $k = 0$ to $K - 1$ and divide both sides by K to obtain

$$\frac{1}{K} \sum_{k=0}^{K-1} \text{Tr}(\mathbf{F}_{\mathbf{K}_k, L(\mathbf{K}_k)}^\top \mathbf{F}_{\mathbf{K}_k, L(\mathbf{K}_k)}) \leq \frac{\|\mathbf{G}_{\mathbf{K}_0, L(\mathbf{K}_0)}\| \text{Tr}(P_{\mathbf{K}_0, L(\mathbf{K}_0)} - P_{\mathbf{K}^*, L(\mathbf{K}^*)})}{2\alpha\phi \cdot K}.$$

This proves that the sequence of natural gradient norm squares, $\{\text{Tr}(\mathbf{F}_{\mathbf{K}_k, L(\mathbf{K}_k)}^\top \mathbf{F}_{\mathbf{K}_k, L(\mathbf{K}_k)})\}_{k \geq 0}$, converges on average with a globally $\mathcal{O}(1/K)$ rate. Also, The convergent stationary point is the unique Nash equilibrium. Finally, faster local rates can be shown by following the techniques presented in Theorem 4.6 of (Zhang et al., 2021a). ■

B.12. Proof of Theorem 4.1 (PG)

Proof We first provide a complete version of Theorem 4.1 for the PG update (4.1) and introduce a helper lemma, whose proof is deferred to §C.2.

Theorem B.2 (Inner-Loop Sample Complexity for PG) For a fixed $\mathbf{K} \in \mathcal{K}$ and an arbitrary \mathbf{L}_0 that induces a finite $\mathcal{G}(\mathbf{K}, \mathbf{L}_0)$, define a superlevel set $\mathcal{L}_{\mathbf{K}}(a)$ as in (3.8), where $a \leq \mathcal{G}(\mathbf{K}, \mathbf{L}_0)$ is an arbitrary constant. Let $\epsilon_1, \delta_1 \in (0, 1)$, and

$M_1, r_1, \eta > 0$ in Algorithm 1 satisfy

$$M_1 \geq \left(\frac{d_1}{r_1} \left(\mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) + \frac{l_{\mathbf{K},a}}{\rho_{\mathbf{K},a}} \right) \sqrt{\log \left(\frac{2d_1 L}{\delta_1} \right)} \right)^2 \frac{1024}{\mu_{\mathbf{K}} \epsilon_1},$$

$$r_1 \leq \min \left\{ \frac{\theta_{\mathbf{K},a} \mu_{\mathbf{K}}}{8\psi_{\mathbf{K},a}} \sqrt{\frac{\epsilon_1}{240}}, \frac{1}{8\psi_{\mathbf{K},a}^2} \sqrt{\frac{\epsilon_1 \mu_{\mathbf{K}}}{30}}, \rho_{\mathbf{K},a} \right\}, \quad \eta \leq \min \left\{ 1, \frac{1}{8\psi_{\mathbf{K},a}}, \rho_{\mathbf{K},a} \cdot \left[\frac{\sqrt{\mu_{\mathbf{K}}}}{32} + \psi_{\mathbf{K},a} + l_{\mathbf{K},a} \right]^{-1} \right\},$$

where $\theta_{\mathbf{K},a} = \min \{1/[2\psi_{\mathbf{K},a}], \rho_{\mathbf{K},a}/l_{\mathbf{K},a}\}$; $l_{\mathbf{K},a}, \psi_{\mathbf{K},a}, \mu_{\mathbf{K}}$ are defined in Lemma 3.4; $\rho_{\mathbf{K},a} := \min_{\mathbf{L} \in \mathcal{L}_{\mathbf{K}}(a)} \rho_{\mathbf{K},\mathbf{L}} > 0$; and $d_1 = nmN$. Then, with probability at least $1 - \delta_1$ and a total number of iterations $L = \frac{8}{\eta \mu_{\mathbf{K}}} \log(\frac{2}{\epsilon_1})$, the inner-loop ZO-PG update (4.1) outputs some $\bar{\mathbf{L}}(\mathbf{K}) := \mathbf{L}_L$ such that $\mathcal{G}(\mathbf{K}, \bar{\mathbf{L}}(\mathbf{K})) \geq \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) - \epsilon_1$, and $\|\mathbf{L}(\mathbf{K}) - \bar{\mathbf{L}}(\mathbf{K})\|_F \leq \sqrt{\lambda_{\min}^{-1}(\mathbf{H}_{\mathbf{K},\mathbf{L}(\mathbf{K}))} \cdot \epsilon_1}$.

Lemma B.3 For a given $\mathbf{K} \in \mathcal{K}$ and a given \mathbf{L} , let $\epsilon_1, \delta_1 \in (0, 1)$ and let the batchsize $M_{1,\mathbf{L}} > 0$ and the smoothing radius $r_{1,\mathbf{L}} > 0$ of a $M_{1,\mathbf{L}}$ -sample one-point minibatch gradient estimator satisfy

$$M_{1,\mathbf{L}} \geq \left(\frac{d_1}{r_{1,\mathbf{L}}} \left(\mathcal{G}(\mathbf{K}, \mathbf{L}) + \frac{l_{\mathbf{K},\mathbf{L}}}{\rho_{\mathbf{K},\mathbf{L}}} \right) \sqrt{\log \left(\frac{2d_1}{\delta_1} \right)} \right)^2 \frac{1024}{\mu_{\mathbf{K}} \epsilon_1}, \quad r_{1,\mathbf{L}} \leq \min \left\{ \frac{\theta_{\mathbf{K},\mathbf{L}} \mu_{\mathbf{K}}}{8\psi_{\mathbf{K},\mathbf{L}}} \sqrt{\frac{\epsilon_1}{240}}, \frac{1}{8\psi_{\mathbf{K},\mathbf{L}}^2} \sqrt{\frac{\epsilon_1 \mu_{\mathbf{K}}}{30}}, \rho_{\mathbf{K},\mathbf{L}} \right\},$$

where $\theta_{\mathbf{K},\mathbf{L}} = \min \left\{ \frac{1}{2\psi_{\mathbf{K},\mathbf{L}}}, \frac{\rho_{\mathbf{K},\mathbf{L}}}{l_{\mathbf{K},\mathbf{L}}} \right\}$; $l_{\mathbf{K},\mathbf{L}}, \psi_{\mathbf{K},\mathbf{L}}, \rho_{\mathbf{K},\mathbf{L}}$ are the local curvature parameters in Lemma 3.4; and $d_1 = nmN$. Also, let the stepsize $\eta_{\mathbf{L}} > 0$ satisfy

$$\eta_{\mathbf{L}} \leq \min \left\{ 1, \frac{1}{8\psi_{\mathbf{K},\mathbf{L}}}, \rho_{\mathbf{K},\mathbf{L}} \cdot \left[\frac{\sqrt{\mu_{\mathbf{K}}}}{32} + \psi_{\mathbf{K},\mathbf{L}} + l_{\mathbf{K},\mathbf{L}} \right] \right\}.$$

Then, the gain matrix after applying one step of (4.1) on \mathbf{L} , denoted as \mathbf{L}' , satisfies with probability at least $1 - \delta_1$ that $\mathcal{G}(\mathbf{K}, \mathbf{L}) \leq \mathcal{G}(\mathbf{K}, \mathbf{L}')$ and

$$\mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) - \mathcal{G}(\mathbf{K}, \mathbf{L}') \leq \left(1 - \frac{\eta_{\mathbf{L}} \mu_{\mathbf{K}}}{8} \right) \cdot (\mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) - \mathcal{G}(\mathbf{K}, \mathbf{L})) + \frac{\eta_{\mathbf{L}} \mu_{\mathbf{K}} \epsilon_1}{16}.$$

We now prove the sample complexity result for a given $\mathbf{L}_0 \in \mathcal{L}_{\mathbf{K}}(a)$. First, we use $\Delta_l := \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) - \mathcal{G}(\mathbf{K}, \mathbf{L}_l)$ to denote the optimality gap at iteration l , where $l \in \{0, \dots, L\}$. By Lemma B.3, if we require $\epsilon_1, \delta_1 \in (0, 1)$, the parameters $M_{1,\mathbf{L}_0}, r_{1,\mathbf{L}_0}, \eta_{\mathbf{L}_0} > 0$ in Algorithm 1 to satisfy

$$M_{1,\mathbf{L}_0} \geq \left(\frac{d_1}{r_{1,\mathbf{L}_0}} \left(\mathcal{G}(\mathbf{K}, \mathbf{L}_0) + \frac{l_{\mathbf{K},\mathbf{L}_0}}{\rho_{\mathbf{K},\mathbf{L}_0}} \right) \sqrt{\log \left(\frac{2d_1 L}{\delta_1} \right)} \right)^2 \frac{1024}{\mu_{\mathbf{K}} \epsilon_1},$$

$$r_{1,\mathbf{L}_0} \leq \min \left\{ \frac{\theta_{\mathbf{K},\mathbf{L}_0} \mu_{\mathbf{K}}}{8\psi_{\mathbf{K},\mathbf{L}_0}} \sqrt{\frac{\epsilon_1}{240}}, \frac{1}{8\psi_{\mathbf{K},\mathbf{L}_0}^2} \sqrt{\frac{\epsilon_1 \mu_{\mathbf{K}}}{30}}, \rho_{\mathbf{K},\mathbf{L}_0} \right\}, \quad \eta_{\mathbf{L}_0} \leq \min \left\{ 1, \frac{1}{8\psi_{\mathbf{K},\mathbf{L}_0}}, \frac{\rho_{\mathbf{K},\mathbf{L}_0}}{\frac{\sqrt{\mu_{\mathbf{K}}}}{32} + \psi_{\mathbf{K},\mathbf{L}_0} + l_{\mathbf{K},\mathbf{L}_0}} \right\},$$

then we ensure with probability at least $1 - \delta_1/L$ that $\mathcal{G}(\mathbf{K}, \mathbf{L}_0) \leq \mathcal{G}(\mathbf{K}, \mathbf{L}_1)$, i.e., $\mathbf{L}_1 \in \mathcal{L}_{\mathbf{K}}(a)$. Moreover, for any \mathbf{L}_l , where $l \in \{0, \dots, L-1\}$, there exist $M_{1,\mathbf{L}_l}, r_{1,\mathbf{L}_l}, \eta_{\mathbf{L}_l} > 0$ as defined in Lemma B.3 that guarantee $\mathcal{G}(\mathbf{K}, \mathbf{L}_l) \leq \mathcal{G}(\mathbf{K}, \mathbf{L}_{l+1})$ with probability at least $1 - \delta_1/L$. Now, we choose uniform constants $M_1, r_1, \eta > 0$ such that

$$M_1 \geq \left(\frac{d_1}{r_1} \left(\mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) + \frac{l_{\mathbf{K},a}}{\rho_{\mathbf{K},a}} \right) \sqrt{\log \left(\frac{2d_1 L}{\delta_1} \right)} \right)^2 \frac{1024}{\mu_{\mathbf{K}} \epsilon_1} \geq \max_{l \in \{0, \dots, L-1\}} M_{1,\mathbf{L}_l},$$

$$r_1 \leq \min \left\{ \frac{\theta_{\mathbf{K},a} \mu_{\mathbf{K}}}{8\psi_{\mathbf{K},a}} \sqrt{\frac{\epsilon_1}{240}}, \frac{1}{8\psi_{\mathbf{K},a}^2} \sqrt{\frac{\epsilon_1 \mu_{\mathbf{K}}}{30}}, \rho_{\mathbf{K},a} \right\} \leq \min_{l \in \{0, \dots, L-1\}} r_{1,\mathbf{L}_l},$$

$$\eta \leq \min \left\{ 1, \frac{1}{8\psi_{\mathbf{K},a}}, \frac{\rho_{\mathbf{K},a}}{\frac{\sqrt{\mu_{\mathbf{K}}}}{32} + \psi_{\mathbf{K},a} + l_{\mathbf{K},a}} \right\} \leq \min_{l \in \{0, \dots, L-1\}} \eta_{\mathbf{L}_l},$$

where $\theta_{\mathbf{K},a} = \min \{1/[2\psi_{\mathbf{K},a}], \rho_{\mathbf{K},a}/l_{\mathbf{K},a}\}$, $\rho_{\mathbf{K},a} = \min_{\mathbf{L} \in \mathcal{L}_{\mathbf{K}}(a)} \rho_{\mathbf{K},\mathbf{L}} > 0$ is due to the compactness of $\mathcal{L}_{\mathbf{K}}(a)$, and $l_{\mathbf{K},a}, \psi_{\mathbf{K},a} < \infty$ are defined in Lemma 3.4 that satisfy

$$l_{\mathbf{K},a} \geq \max_{\mathbf{L} \in \mathcal{L}_{\mathbf{K}}(a)} l_{\mathbf{K},\mathbf{L}}, \quad \psi_{\mathbf{K},a} \geq \max_{\mathbf{L} \in \mathcal{L}_{\mathbf{K}}(a)} \psi_{\mathbf{K},\mathbf{L}}.$$

Then, we can guarantee with probability at least $1 - \delta_1$ that the value of the objective function, following the ZO-PG update (4.1), is monotonically non-decreasing. That is, we have with probability at least $1 - \delta_1$ that $\mathbf{L}_l \in \mathcal{L}_{\mathbf{K}}(a)$, for all $l \in \{0, \dots, L\}$, when a $\mathbf{L}_0 \in \mathcal{L}_{\mathbf{K}}(a)$ is given. By Lemma B.3 and the above choices of M_1, r_1, η , we also ensure with probability at least $1 - \delta_1$ that $\Delta_l \leq (1 - \frac{\eta\mu_{\mathbf{K}}}{8}) \cdot \Delta_{l-1} + \frac{\eta\mu_{\mathbf{K}}\epsilon_1}{16}$, for all $l \in \{1, \dots, L-1\}$. Thus, we can show with probability at least $1 - \delta_1$ that

$$\Delta_L \leq \left(1 - \frac{\eta\mu_{\mathbf{K}}}{8}\right) \cdot \Delta_{L-1} + \eta \frac{\mu_{\mathbf{K}}\epsilon_1}{16} \leq \left(1 - \frac{\eta\mu_{\mathbf{K}}}{8}\right)^L \cdot \Delta_0 + \sum_{i=1}^{L-1} \left(1 - \frac{\eta\mu_{\mathbf{K}}}{8}\right)^i \eta \frac{\mu_{\mathbf{K}}\epsilon_1}{16} \leq \left(1 - \frac{\eta\mu_{\mathbf{K}}}{8}\right)^L \Delta_0 + \frac{\epsilon_1}{2}.$$

As a result, when $L = \frac{8}{\eta\mu_{\mathbf{K}}} \log(\frac{2}{\epsilon_1})$, the inequality $\mathcal{G}(\mathbf{K}, \mathbf{L}_L) \leq \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) - \epsilon_1$ holds with probability at least $1 - \delta_1$. This proves the convergence of the generated values of the objective function for the ZO-PG update (4.1). Lastly, we demonstrate the convergence of the gain matrix to $\mathbf{L}(\mathbf{K})$ by our results in §B.9. Based on (B.10) and the convergence of the objective function values, we have with probability at least $1 - \delta_1$ that

$$\|\mathbf{L}(\mathbf{K}) - \mathbf{L}_L\|_F \leq \sqrt{\lambda_{\min}^{-1}(\mathbf{H}_{\mathbf{K},\mathbf{L}(\mathbf{K}))} \cdot (\mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) - \mathcal{G}(\mathbf{K}, \mathbf{L}_L))} \leq \sqrt{\lambda_{\min}^{-1}(\mathbf{H}_{\mathbf{K},\mathbf{L}(\mathbf{K}))} \cdot \epsilon_1}. \quad (\text{B.20})$$

This completes the proof. \blacksquare

B.13. Proof of Theorem 4.1 (NPG)

Proof We first provide a complete version of Theorem 4.1 for the NPG update (4.2) and introduce a helper lemma, whose proof is deferred to §C.3.

Theorem B.4 (Inner-Loop Sample Complexity for NPG) For a fixed $\mathbf{K} \in \mathcal{K}$ and an arbitrary \mathbf{L}_0 that induces a finite $\mathcal{G}(\mathbf{K}, \mathbf{L}_0)$, define a superlevel set $\mathcal{L}_{\mathbf{K}}(a)$ as in (3.8), where $a \leq \mathcal{G}(\mathbf{K}, \mathbf{L}_0)$ is an arbitrary constant. Let $\epsilon_1, \delta_1 \in (0, 1)$, and $M_1, r_1, \eta > 0$ in Algorithm 1 satisfy

$$\begin{aligned} M_1 &\geq \max \left\{ \left(\mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) + \frac{l_{\mathbf{K},a}}{\rho_{\mathbf{K},a}} \right)^2 \cdot \frac{64d_1^2(\bar{\kappa}_a + 1)^2}{\phi^2 r_1^2 \mu_{\mathbf{K}} \epsilon_1}, \frac{2\bar{\kappa}_a^2}{\phi^2} \right\} \cdot \log \left(\frac{4L \max\{d_1, d_\Sigma\}}{\delta_1} \right), \\ r_1 &\leq \min \left\{ \frac{\phi \sqrt{\mu_{\mathbf{K}} \epsilon_1}}{16\bar{\kappa}_a \psi_{\mathbf{K},a}}, \frac{\phi \mu_{\mathbf{K}} \theta_{\mathbf{K},a} \sqrt{\epsilon_1/2}}{32\bar{\kappa}_a \psi_{\mathbf{K},a}}, \rho_{\mathbf{K},a} \right\}, \\ \eta &\leq \min \left\{ \frac{\phi^2}{32\psi_{\mathbf{K},a} \bar{\kappa}_a}, \frac{1}{2\psi_{\mathbf{K},a}}, \rho_{\mathbf{K},a} \cdot \left[\frac{\sqrt{\mu_{\mathbf{K}}}}{4(\bar{\kappa}_a + 1)} + \frac{2\psi_{\mathbf{K},a}}{\phi} + l_{\mathbf{K},a} + \frac{\phi l_{\mathbf{K},a}}{2} \right]^{-1} \right\}, \end{aligned}$$

where $\theta_{\mathbf{K},a} = \min \{1/[2\psi_{\mathbf{K},a}], \rho_{\mathbf{K},a}/l_{\mathbf{K},a}\}$; $l_{\mathbf{K},a}, \psi_{\mathbf{K},a}, \mu_{\mathbf{K}}$ are defined in Lemma 3.4; $\rho_{\mathbf{K},a}, \bar{\kappa}_a, \bar{\kappa}_a$ are uniform constants over $\mathcal{L}_{\mathbf{K}}(a)$ defined in §B.13; $d_\Sigma = m^2(N+1)$; and $d_1 = nmN$. Then, with probability at least $1 - \delta_1$ and a total number of iterations $L = \frac{8\bar{\kappa}_a}{\eta\mu_{\mathbf{K}}} \log(\frac{2}{\epsilon_1})$, the inner-loop ZO-NPG update (4.2) outputs some $\bar{\mathbf{L}}(\mathbf{K}) := \mathbf{L}_L$ such that $\mathcal{G}(\mathbf{K}, \bar{\mathbf{L}}(\mathbf{K})) \geq \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) - \epsilon_1$, and $\|\mathbf{L}(\mathbf{K}) - \bar{\mathbf{L}}(\mathbf{K})\|_F \leq \sqrt{\lambda_{\min}^{-1}(\mathbf{H}_{\mathbf{K},\mathbf{L}(\mathbf{K}))} \cdot \epsilon_1}$.

Lemma B.5 For a given $\mathbf{K} \in \mathcal{K}$ and a given \mathbf{L} , let $\epsilon_1, \delta_1 \in (0, 1)$ and let the batchsize $M_{1,\mathbf{L}} > 0$ and the smoothing radius $r_{1,\mathbf{L}} > 0$ of a $M_{1,\mathbf{L}}$ -sample one-point minibatch gradient estimator satisfy

$$\begin{aligned} M_{1,\mathbf{L}} &\geq \max \left\{ \left(\mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) + \frac{l_{\mathbf{K},\mathbf{L}}}{\rho_{\mathbf{K},\mathbf{L}}} \right)^2 \cdot \frac{64d_1^2(\kappa + 1)^2}{\phi^2 r_{1,\mathbf{L}}^2 \mu_{\mathbf{K}} \epsilon_1}, \frac{2\kappa^2}{\phi^2} \right\} \cdot \log \left(\frac{4 \max\{d_1, d_\Sigma\}}{\delta_1} \right), \\ r_{1,\mathbf{L}} &\leq \min \left\{ \frac{\phi \sqrt{\mu_{\mathbf{K}} \epsilon_1}}{16\kappa \psi_{\mathbf{K},\mathbf{L}}}, \frac{\phi \mu_{\mathbf{K}} \theta_{\mathbf{K},\mathbf{L}} \sqrt{\epsilon_1/2}}{32\kappa \psi_{\mathbf{K},\mathbf{L}}}, \rho_{\mathbf{K},\mathbf{L}} \right\}, \end{aligned}$$

where $\theta_{\mathbf{K},\mathbf{L}} = \min \{1/[2\psi_{\mathbf{K},\mathbf{L}}], \rho_{\mathbf{K},\mathbf{L}}/l_{\mathbf{K},\mathbf{L}}\}$; $l_{\mathbf{K},\mathbf{L}}, \psi_{\mathbf{K},\mathbf{L}}, \rho_{\mathbf{K},\mathbf{L}}$ are the local curvature parameters in Lemma 3.4; $d_1 = nmN$; $d_\Sigma = m^2(N+1)$; $\varkappa := c_{\Sigma_{\mathbf{K},\mathbf{L}}} + \frac{\phi}{2}$; and $c_{\Sigma_{\mathbf{K},\mathbf{L}}}$ is a polynomial of $\|\mathbf{A}\|_F, \|\mathbf{B}\|_F, \|\mathbf{D}\|_F, \|\mathbf{K}\|_F, \|\mathbf{L}\|_F$ that is linear in c_0 , defined in Lemma C.1. Also, let the stepsize $\eta_L > 0$ satisfy

$$\eta_L \leq \min \left\{ \frac{\phi^2}{32\psi_{\mathbf{K},\mathbf{L}}\varkappa}, \frac{1}{2\psi_{\mathbf{K},\mathbf{L}}}, \rho_{\mathbf{K},\mathbf{L}} \cdot \left[\frac{\sqrt{\mu\mathbf{K}}}{4(\varkappa+1)} + \frac{2\psi_{\mathbf{K},\mathbf{L}}}{\phi} + l_{\mathbf{K},\mathbf{L}} + \frac{\phi l_{\mathbf{K},\mathbf{L}}}{2} \right]^{-1} \right\}.$$

Then, the gain matrix after applying one step of (4.2) on \mathbf{L} , denoted as \mathbf{L}' , satisfies with probability at least $1 - \delta_1$ that $\mathcal{G}(\mathbf{K}, \mathbf{L}) \leq \mathcal{G}(\mathbf{K}, \mathbf{L}')$ and

$$\mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) - \mathcal{G}(\mathbf{K}, \mathbf{L}') \leq \left(1 - \frac{\eta_L \mu \mathbf{K}}{8\varkappa}\right) \cdot (\mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) - \mathcal{G}(\mathbf{K}, \mathbf{L})) + \frac{\eta_L \mu \mathbf{K} \epsilon_1}{16\varkappa}.$$

Based on Lemma B.5, the rest of the proof mostly follows the proof of Theorem 4.1 in §B.12. Specifically, we can choose uniform constants $M_1, r_1, \eta > 0$ such that

$$\begin{aligned} M_1 &\geq \max \left\{ \left(\mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) + \frac{l_{\mathbf{K},a}}{\rho_{\mathbf{K},a}} \right)^2 \cdot \frac{64d_1^2(\bar{\varkappa}_a+1)^2}{\phi^2 r_1^2 \mu \mathbf{K} \epsilon_1}, \frac{2\bar{\varkappa}_a^2}{\phi^2} \right\} \cdot \log \left(\frac{4L \max\{d_1, d_\Sigma\}}{\delta_1} \right) \geq \max_{l \in \{0, \dots, L-1\}} M_{1,L_l}, \\ r_1 &\leq \min \left\{ \frac{\phi \sqrt{\mu \mathbf{K} \epsilon_1}}{16\bar{\varkappa}_a \psi_{\mathbf{K},a}}, \frac{\phi \mu \mathbf{K} \theta_{\mathbf{K},a} \sqrt{\epsilon_1/2}}{32\bar{\varkappa}_a \psi_{\mathbf{K},a}}, \rho_{\mathbf{K},a} \right\} \leq \min_{l \in \{0, \dots, L-1\}} r_{1,L_l}, \\ \eta &\leq \min \left\{ \frac{\phi^2}{32\psi_{\mathbf{K},a}\bar{\varkappa}_a}, \frac{1}{2\psi_{\mathbf{K},a}}, \rho_{\mathbf{K},a} \cdot \left[\frac{\sqrt{\mu \mathbf{K}}}{4(\bar{\varkappa}_a+1)} + \frac{2\psi_{\mathbf{K},a}}{\phi} + l_{\mathbf{K},a} + \frac{\phi l_{\mathbf{K},a}}{2} \right]^{-1} \right\} \leq \min_{l \in \{0, \dots, L-1\}} \eta_{L_l}, \end{aligned}$$

where $\theta_{\mathbf{K},a} = \min \{1/[2\psi_{\mathbf{K},a}], \rho_{\mathbf{K},a}/l_{\mathbf{K},a}\}$ and

$$\bar{\varkappa}_a := \max_{\mathbf{L} \in \mathcal{L}_{\mathbf{K}(a)}} \varkappa < \infty, \bar{\varkappa}_a := \min_{\mathbf{L} \in \mathcal{L}_{\mathbf{K}(a)}} \varkappa > 0, \rho_{\mathbf{K},a} = \min_{\mathbf{L} \in \mathcal{L}_{\mathbf{K}(a)}} \rho_{\mathbf{K},\mathbf{L}} > 0, l_{\mathbf{K},a} \geq \max_{\mathbf{L} \in \mathcal{L}_{\mathbf{K}(a)}} l_{\mathbf{K},\mathbf{L}}, \psi_{\mathbf{K},a} \geq \max_{\mathbf{L} \in \mathcal{L}_{\mathbf{K}(a)}} \psi_{\mathbf{K},\mathbf{L}}.$$

Then, we can guarantee with probability at least $1 - \delta_1$ that the value of the objective function, following the ZO-NPG update (4.2), is monotonically non-decreasing. That is, we have with probability at least $1 - \delta_1$ that $\mathbf{L}_l \in \mathcal{L}_{\mathbf{K}(a)}$, for all $l \in \{0, \dots, L\}$, when an $\mathbf{L}_0 \in \mathcal{L}_{\mathbf{K}(a)}$ is given. By Lemma B.3 and the above choices of M_1, r_1, η , we also ensure with probability at least $1 - \delta_1$ that $\Delta_l \leq \left(1 - \frac{\eta \mu \mathbf{K}}{8\bar{\varkappa}_a}\right) \cdot \Delta_{l-1} + \frac{\eta \mu \mathbf{K} \epsilon_1}{16\bar{\varkappa}_a}$, for all $l \in \{1, \dots, L-1\}$. Thus, we can show with probability at least $1 - \delta_1$ that

$$\Delta_L \leq \left(1 - \frac{\eta \mu \mathbf{K}}{8\bar{\varkappa}_a}\right) \cdot \Delta_{L-1} + \frac{\eta \mu \mathbf{K} \epsilon_1}{16\bar{\varkappa}_a} \leq \left(1 - \frac{\eta \mu \mathbf{K}}{8\bar{\varkappa}_a}\right)^L \cdot \Delta_0 + \sum_{i=1}^{L-1} \left(1 - \frac{\eta \mu \mathbf{K}}{8\bar{\varkappa}_a}\right)^i \frac{\eta \mu \mathbf{K} \epsilon_1}{16\bar{\varkappa}_a} \leq \left(1 - \frac{\eta \mu \mathbf{K}}{8\bar{\varkappa}_a}\right)^L \Delta_0 + \frac{\epsilon_1}{2}.$$

As a result, when $L = \frac{8\bar{\varkappa}_a}{\eta \mu \mathbf{K}} \log(\frac{2}{\epsilon_1})$, the inequality $\mathcal{G}(\mathbf{K}, \mathbf{L}_L) \leq \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) - \epsilon_1$ holds with probability at least $1 - \delta_1$. This proves the convergence of the generated values of the objective function for the ZO-NPG update (4.2). The convergence of the gain matrix to $\mathbf{L}(\mathbf{K})$ follows from the proof of Theorem 4.1 in §B.12. This completes the proof. ■

B.14. Proof of Theorem 4.2

Proof We first present a complete version of Theorem 4.2 and a few useful lemmas, whose proofs are deferred to §C.4-§C.8.

Theorem B.6 (IR: Derivative-Free Setting) For any $\mathbf{K}_0 \in \mathcal{K}$ and defining $\zeta := \lambda_{\min}(\mathbf{H}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)}) > 0$, introduce the following set $\hat{\mathcal{K}}$:

$$\hat{\mathcal{K}} := \left\{ \mathbf{K} \mid (3.6) \text{ admits a solution } \mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})} \geq 0, \text{ and } \mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})} \leq \mathbf{P}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)} + \frac{\zeta}{2\|\mathbf{D}\|^2} \cdot \mathbf{I} \right\} \subset \mathcal{K}. \quad (\text{B.21})$$

Let $\delta_2 \in (0, 1)$ and other parameters in Algorithm 2 satisfy $\delta_1 \leq \delta_2/[6M_2K]$ and

$$\begin{aligned}
 M_2 \geq & \max \left\{ \frac{128d_2^2\alpha^2c_0^2\widehat{\mathcal{G}}(\mathbf{K}, \mathbf{L}(\mathbf{K}))^2}{r_2^2\phi^4\varpi^2}, \frac{512d_2^2\alpha^2(\widehat{\mathcal{G}}(\mathbf{K}, \mathbf{L}(\mathbf{K})) + r_2\mathcal{B}_{P,C_0})^2}{r_2^2\phi^2\varpi^2}, \frac{32\alpha^2(\widehat{c}_{5,\mathbf{K}})^2(\widehat{c}_{\Sigma,\mathbf{K},\mathbf{L}(\mathbf{K}))}^2}{\phi^2\varpi^2}, \frac{8(\widehat{c}_{\Sigma,\mathbf{K},\mathbf{L}(\mathbf{K}))}^2}{\phi^2} \right\} \\
 & \cdot \log \left(\frac{12K \max\{d_2, d_\Sigma\}}{\delta_2} \right), \quad \epsilon_1 \leq \min \left\{ \frac{\phi\varpi r_2}{16\alpha d_2}, \frac{\widehat{\mathcal{B}}_{1,\mathbf{L}(\mathbf{K})}^2\zeta}{2}, \frac{\phi^2\varpi^2\zeta}{128\alpha^2(\widehat{c}_{5,\mathbf{K}})^2\widehat{\mathcal{B}}_{\Sigma,\mathbf{L}(\mathbf{K})}^2}, \frac{\phi^2\zeta}{32\widehat{\mathcal{B}}_{\Sigma,\mathbf{L}(\mathbf{K})}^2} \right\}, \\
 r_2 \leq & \min \left\{ \varpi, \frac{\phi\varpi}{64\alpha\widehat{c}_{5,\mathbf{K}}(\widehat{c}_{\Sigma,\mathbf{K},\mathbf{L}(\mathbf{K}))} + \widehat{\mathcal{B}}_{\Sigma,\mathbf{K}}}, \frac{\phi\varpi}{64\alpha\widehat{c}_{2,\mathbf{K}}\widehat{\mathcal{B}}_{\Sigma,\mathbf{K}}} \right\}, \\
 \alpha \leq & \frac{1}{2} \cdot \|\mathbf{R}^u + \mathbf{B}(\overline{\mathbf{P}} + \overline{\mathbf{P}}\mathbf{D}(\mathbf{R}^w - \mathbf{D}^\top\overline{\mathbf{P}}\mathbf{D})^{-1}\mathbf{D}^\top\overline{\mathbf{P}})\mathbf{B}\|^{-1}, \quad \text{where } \overline{\mathbf{P}} := \mathbf{P}_{\mathbf{K}_0,\mathbf{L}(\mathbf{K}_0)} + \frac{\zeta}{2\|\mathbf{D}\|^2} \cdot \mathbf{I}, \quad (\text{B.22})
 \end{aligned}$$

where $\widehat{\mathcal{G}}(\mathbf{K}, \mathbf{L}(\mathbf{K})), \widehat{c}_{2,\mathbf{K}}, \widehat{c}_{5,\mathbf{K}}, \widehat{c}_{\Sigma,\mathbf{K},\mathbf{L}(\mathbf{K})}, \widehat{\mathcal{B}}_{\Sigma,\mathbf{K}}, \widehat{\mathcal{B}}_{1,\mathbf{L}(\mathbf{K})}, \widehat{\mathcal{B}}_{\Sigma,\mathbf{L}(\mathbf{K})}, \varpi > 0$ are uniform constants over $\widehat{\mathcal{K}}$ defined in §B.14. Then, it holds with probability at least $1 - \delta_2$ that $\mathbf{K}_k \in \widehat{\mathcal{K}} \subset \mathcal{K}$ for all $k \in \{1, \dots, K\}$.

Lemma B.7 For any $\mathbf{K}, \mathbf{K}' \in \mathcal{K}$, there exist some $\mathcal{B}_{1,\mathbf{K}}, \mathcal{B}_{P,\mathbf{K}}, \mathcal{B}_{L(\mathbf{K}),\mathbf{K}}, \mathcal{B}_{\Sigma,\mathbf{K}} > 0$ that are continuous functions of \mathbf{K} such that all \mathbf{K}' satisfying $\|\mathbf{K}' - \mathbf{K}\|_F \leq \mathcal{B}_{1,\mathbf{K}}$ satisfy

$$\|\mathbf{P}_{\mathbf{K}',\mathbf{L}(\mathbf{K}')} - \mathbf{P}_{\mathbf{K},\mathbf{L}(\mathbf{K})}\|_F \leq \mathcal{B}_{P,\mathbf{K}} \cdot \|\mathbf{K}' - \mathbf{K}\|_F \quad (\text{B.23})$$

$$\|\mathbf{L}(\mathbf{K}') - \mathbf{L}(\mathbf{K})\|_F \leq \mathcal{B}_{L(\mathbf{K}),\mathbf{K}} \cdot \|\mathbf{K}' - \mathbf{K}\|_F \quad (\text{B.24})$$

$$\|\Sigma_{\mathbf{K}',\mathbf{L}(\mathbf{K}')} - \Sigma_{\mathbf{K},\mathbf{L}(\mathbf{K})}\|_F \leq \mathcal{B}_{\Sigma,\mathbf{K}} \cdot \|\mathbf{K}' - \mathbf{K}\|_F. \quad (\text{B.25})$$

Lemma B.8 For any $\mathbf{K}, \mathbf{K}' \in \mathcal{K}$, there exist some $\mathcal{B}_{1,\mathbf{K}}, \mathcal{B}_{\Sigma,\mathbf{K}} > 0$ as defined in Lemma B.7 such that if \mathbf{K}' satisfies

$$\|\mathbf{K}' - \mathbf{K}\|_F \leq \left\{ \mathcal{B}_{1,\mathbf{K}}, \frac{\epsilon_2}{4c_{5,\mathbf{K}}(c_{\Sigma,\mathbf{K},\mathbf{L}(\mathbf{K}))} + \mathcal{B}_{\Sigma,\mathbf{K}}}, \frac{\epsilon_2}{4c_{2,\mathbf{K}}\mathcal{B}_{\Sigma,\mathbf{K}}} \right\},$$

where $c_{\Sigma,\mathbf{K},\mathbf{L}(\mathbf{K})}$ is a polynomial of $\|\mathbf{A}\|_F, \|\mathbf{B}\|_F, \|\mathbf{D}\|_F, \|\mathbf{K}\|_F$, and c_0 , and $c_{2,\mathbf{K}}, c_{5,\mathbf{K}}$ are defined in §C.9, then it holds that $\|\nabla_{\mathbf{K}}\mathcal{G}(\mathbf{K}', \mathbf{L}(\mathbf{K}')) - \nabla_{\mathbf{K}}\mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K}))\|_F \leq \epsilon_2$.

Lemma B.9 For any $\mathbf{K} \in \mathcal{K}$, there exists some $\mathcal{B}_{2,\mathbf{K}} > 0$ such that all \mathbf{K}' satisfying $\|\mathbf{K}' - \mathbf{K}\|_F \leq \mathcal{B}_{2,\mathbf{K}}$ satisfy $\mathbf{K}' \in \mathcal{K}$.

Lemma B.10 For any $\mathbf{K} \in \mathcal{K}$, let the batchsize M_2 , smoothing radius r_2 , inner-loop parameters ϵ_1 and $\delta_1 \in (0, 1)$ satisfy

$$\begin{aligned}
 M_2 \geq & \max \left\{ \frac{8d_2^2c_0^2\mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K}))^2}{r_2^2\phi^2\epsilon_2^2}, \frac{32d_2^2(\mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) + r_2\mathcal{B}_{P,C_0})^2}{r_2^2\epsilon_2^2} \right\} \cdot \log \left(\frac{6d_2}{\delta_2} \right), \quad \epsilon_1 \leq \frac{\epsilon_2 r_2}{4d_2} \\
 r_2 \leq & \min \left\{ \mathcal{B}_{1,\mathbf{K}}, \mathcal{B}_{2,\mathbf{K}}, \frac{\epsilon_2}{16c_{5,\mathbf{K}}(c_{\Sigma,\mathbf{K},\mathbf{L}(\mathbf{K}))} + \mathcal{B}_{\Sigma,\mathbf{K}}}, \frac{\epsilon_2}{16c_{2,\mathbf{K}}\mathcal{B}_{\Sigma,\mathbf{K}}} \right\}, \quad \delta_1 \leq \frac{\delta_2}{3M_2}.
 \end{aligned}$$

where $\mathcal{B}_{1,\mathbf{K}}, \mathcal{B}_{2,\mathbf{K}}, \mathcal{B}_{\Sigma,\mathbf{K}}, c_{\Sigma,\mathbf{K},\mathbf{L}(\mathbf{K})} > 0$ are defined in Lemmas B.7, B.8, and B.9, and $d_2 = dmN$. Then, we have with probability at least $1 - \delta_2$ that $\|\overline{\nabla}_{\mathbf{K}}\mathcal{G}(\mathbf{K}, \overline{\mathbf{L}}(\mathbf{K})) - \nabla_{\mathbf{K}}\mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K}))\|_F \leq \epsilon_2$, where $\overline{\nabla}_{\mathbf{K}}\mathcal{G}(\mathbf{K}, \overline{\mathbf{L}}(\mathbf{K}))$ is the estimated PG from Algorithm 2.

Lemma B.11 For any $\mathbf{K} \in \mathcal{K}$, let the batchsize M_2 , inner-loop parameters ϵ_1 and $\delta_1 \in (0, 1)$ satisfy

$$M_2 \geq \frac{2c_{\Sigma,\mathbf{K},\mathbf{L}(\mathbf{K})}^2}{\epsilon_2^2} \cdot \log \left(\frac{4d_\Sigma}{\delta_2} \right), \quad \epsilon_1 \leq \min \left\{ \mathcal{B}_{1,\mathbf{L}(\mathbf{K})}^2 \lambda_{\min}(\mathbf{H}_{\mathbf{K},\mathbf{L}(\mathbf{K})}), \frac{\epsilon_2^2 \lambda_{\min}(\mathbf{H}_{\mathbf{K},\mathbf{L}(\mathbf{K})})}{4\mathcal{B}_{\Sigma,\mathbf{L}(\mathbf{K})}^2} \right\}, \quad \delta_1 \leq \frac{\delta_2}{2},$$

where $\mathcal{B}_{1,\mathbf{L}(\mathbf{K})}, \mathcal{B}_{\Sigma,\mathbf{L}(\mathbf{K})}$ are defined in Lemma C.2 and $c_{\Sigma,\mathbf{K},\mathbf{L}(\mathbf{K})}$ is a polynomial of $\|\mathbf{A}\|_F, \|\mathbf{B}\|_F, \|\mathbf{D}\|_F, \|\mathbf{K}\|_F$, and c_0 . Then, we have with probability at least $1 - \delta_2$ that $\|\overline{\Sigma}_{\mathbf{K},\overline{\mathbf{L}}(\mathbf{K})} - \Sigma_{\mathbf{K},\mathbf{L}(\mathbf{K})}\|_F \leq \epsilon_2$, with $\overline{\Sigma}_{\mathbf{K},\overline{\mathbf{L}}(\mathbf{K})}$ being the estimated correlation matrix from Algorithm 2. Moreover, it holds with probability at least $1 - \delta_2$ that $\lambda_{\min}(\overline{\Sigma}_{\mathbf{K},\overline{\mathbf{L}}(\mathbf{K})}) \geq \phi/2$ if $\epsilon_2 \leq \phi/2$.

Based on the above lemmas, we prove the implicit regularization property of the outer-loop ZO-NPG update (4.3). Define $\zeta := \lambda_{\min}(\mathbf{H}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)}) > 0$ and the following set $\widehat{\mathcal{K}}$:

$$\widehat{\mathcal{K}} := \left\{ \mathbf{K} \mid (3.6) \text{ admits a solution } \mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})} \geq 0, \text{ and } \mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})} \leq \mathbf{P}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)} + \frac{\zeta}{2\|\mathbf{D}\|^2} \cdot \mathbf{I} \right\} \subset \mathcal{K}.$$

It is a strict subset of \mathcal{K} as all $\mathbf{K} \in \widehat{\mathcal{K}}$ satisfy $\lambda_{\min}(\mathbf{H}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}) \geq \frac{\zeta}{2} > 0$. Then, we prove the compactness of $\widehat{\mathcal{K}}$ by first proving its boundedness. Specifically, for any $\mathbf{K} \in \widehat{\mathcal{K}} \subset \mathcal{K}$, we have by (3.6) that $\mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})} \geq 0$ solves

$$\mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})} = \mathbf{Q} + \mathbf{K}^\top \mathbf{R}^u \mathbf{K} + (\mathbf{A} - \mathbf{B}\mathbf{K})^\top (\mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})} + \mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})} \mathbf{D} (\mathbf{R}^w - \mathbf{D}^\top \mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})} \mathbf{D})^{-1} \mathbf{D}^\top \mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}) (\mathbf{A} - \mathbf{B}\mathbf{K}), \quad (\text{B.26})$$

where the second term on the RHS of (B.26) is p.s.d. and thus $\mathbf{Q} + \mathbf{K}^\top \mathbf{R}^u \mathbf{K} \leq \mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}$ with $\mathbf{Q} \geq 0$ and $\mathbf{R}^u > 0$. Since $\mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}, \mathbf{P}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)}$ are symmetric and p.s.d., all $\mathbf{K} \in \widehat{\mathcal{K}}$ satisfy $\|\mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F \leq \|\mathbf{P}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)} + \frac{\zeta}{2\|\mathbf{D}\|^2} \cdot \mathbf{I}\|_F$. These arguments together imply that for all $\mathbf{K} \in \widehat{\mathcal{K}}$, $\|\mathbf{K}\|_F \leq \sqrt{\|\mathbf{P}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)} + \frac{\zeta}{2\|\mathbf{D}\|^2} \cdot \mathbf{I}\|_F / \lambda_{\min}(\mathbf{R}^u)}$, proving the boundedness of $\widehat{\mathcal{K}}$.

Next, we take an arbitrary sequence $\{\mathbf{K}_n\} \in \widehat{\mathcal{K}}$ and note that $\|\mathbf{K}_n\|_F$ is bounded for all n . Applying Bolzano-Weierstrass theorem implies that the set of limit points of $\{\mathbf{K}_n\}$, denoted as $\widehat{\mathcal{K}}_{\text{lim}}$, is nonempty. Then, for any $\mathbf{K}_{\text{lim}} \in \widehat{\mathcal{K}}_{\text{lim}}$, we can find a subsequence $\{\mathbf{K}_{\tau_n}\} \in \widehat{\mathcal{K}}$ that converges to \mathbf{K}_{lim} . We denote the corresponding sequence of solutions to (B.26) as $\{\mathbf{P}_{\mathbf{K}_{\tau_n}, \mathbf{L}(\mathbf{K}_{\tau_n})}\}$, where $0 \leq \mathbf{P}_{\mathbf{K}_{\tau_n}, \mathbf{L}(\mathbf{K}_{\tau_n})} \leq \mathbf{P}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)} + \frac{\zeta}{2\|\mathbf{D}\|^2} \cdot \mathbf{I}$ for all n . By Bolzano-Weierstrass theorem, the boundedness of $\{\mathbf{P}_{\mathbf{K}_{\tau_n}, \mathbf{L}(\mathbf{K}_{\tau_n})}\}$, and the continuity of (B.26) with respect to \mathbf{K} , we have the set of limit points of $\{\mathbf{P}_{\mathbf{K}_{\tau_n}, \mathbf{L}(\mathbf{K}_{\tau_n})}\}$, denoted as $\widehat{\mathcal{P}}_{\text{lim}}$, is nonempty. Then, for any $\mathbf{P}_{\text{lim}} \in \widehat{\mathcal{P}}_{\text{lim}}$, we can again find a subsequence $\{\mathbf{P}_{\mathbf{K}_{\kappa\tau_n}, \mathbf{L}(\mathbf{K}_{\kappa\tau_n})}\}$ that converges to \mathbf{P}_{lim} . Since $\mathbf{P}_{\mathbf{K}_{\kappa\tau_n}, \mathbf{L}(\mathbf{K}_{\kappa\tau_n})}$ is a p.s.d. solution to (B.26) satisfying $0 \leq \mathbf{P}_{\mathbf{K}_{\kappa\tau_n}, \mathbf{L}(\mathbf{K}_{\kappa\tau_n})} \leq \mathbf{P}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)} + \frac{\zeta}{2\|\mathbf{D}\|^2} \cdot \mathbf{I}$ for all n and (B.26) is continuous in \mathbf{K} , \mathbf{P}_{lim} must solve (B.26) and satisfy $0 \leq \mathbf{P}_{\text{lim}} \leq \mathbf{P}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)} + \frac{\zeta}{2\|\mathbf{D}\|^2} \cdot \mathbf{I}$, which implies $\mathbf{K}_{\text{lim}} \in \widehat{\mathcal{K}}$. Note that the above arguments work for any sequence $\{\mathbf{K}_n\} \in \widehat{\mathcal{K}}$ and any limit points \mathbf{K}_{lim} and \mathbf{P}_{lim} , which proves the closedness of $\widehat{\mathcal{K}}$. Together with the boundedness, $\widehat{\mathcal{K}}$ is thus compact.

Now, denote the iterates after one-step of the outer-loop ZO-NPG update (4.3) and the exact NPG update (3.11) from \mathbf{K}_k as \mathbf{K}_{k+1} and $\widetilde{\mathbf{K}}_{k+1}$, respectively, for $k \in \{0, \dots, K-1\}$. Clearly, it holds that $\mathbf{K}_0 \in \widehat{\mathcal{K}}$. Additionally, we require the stepsize $\alpha > 0$ to satisfy

$$\alpha \leq \frac{1}{2} \cdot \|\mathbf{R}^u + \mathbf{B}(\overline{\mathbf{P}} + \overline{\mathbf{P}}\mathbf{D}(\mathbf{R}^w - \mathbf{D}^\top \overline{\mathbf{P}}\mathbf{D})^{-1} \mathbf{D}^\top \overline{\mathbf{P}})\mathbf{B}\|^{-1}, \quad \text{where } \overline{\mathbf{P}} := \mathbf{P}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)} + \frac{\zeta}{2\|\mathbf{D}\|^2} \cdot \mathbf{I}, \quad (\text{B.27})$$

which is stricter than the requirement in Theorem 3.7. Given our stepsize choice, Theorem 3.7 guarantees that $\mathbf{P}_{\widetilde{\mathbf{K}}_1, \mathbf{L}(\widetilde{\mathbf{K}}_1)} \geq 0$ exists and satisfies $\mathbf{P}_{\widetilde{\mathbf{K}}_1, \mathbf{L}(\widetilde{\mathbf{K}}_1)} \leq \mathbf{P}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)}$ almost surely. By requiring $\|\mathbf{K}_1 - \widetilde{\mathbf{K}}_1\|_F \leq \min\{\mathcal{B}_{1, \widetilde{\mathbf{K}}_1}, \mathcal{B}_{2, \widetilde{\mathbf{K}}_1}, \frac{\zeta}{2\mathcal{B}_{\mathbf{P}, \widetilde{\mathbf{K}}_1} K \|\mathbf{D}\|^2}\}$ with probability at least $1 - \frac{\delta_2}{K}$, we have by Lemma B.9 that $\mathbf{K}_1 \in \mathcal{K}$ with probability at least $1 - \frac{\delta_2}{K}$. Thus, we leverage the recursive arguments in the proof of Theorem 3.7 (see §B.10) to show that $\mathbf{P}_{\mathbf{K}_1, \mathbf{L}(\mathbf{K}_1)} \geq 0$ exists and invoke Lemma B.7 to obtain that

$$\|\mathbf{P}_{\mathbf{K}_1, \mathbf{L}(\mathbf{K}_1)} - \mathbf{P}_{\widetilde{\mathbf{K}}_1, \mathbf{L}(\widetilde{\mathbf{K}}_1)}\| \leq \frac{\zeta}{2K\|\mathbf{D}\|^2} \implies \mathbf{P}_{\mathbf{K}_1, \mathbf{L}(\mathbf{K}_1)} \leq \mathbf{P}_{\widetilde{\mathbf{K}}_1, \mathbf{L}(\widetilde{\mathbf{K}}_1)} + \frac{\zeta}{2K\|\mathbf{D}\|^2} \cdot \mathbf{I} \leq \mathbf{P}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)} + \frac{\zeta}{2K\|\mathbf{D}\|^2} \cdot \mathbf{I}. \quad (\text{B.28})$$

In other words, it holds that $\mathbf{K}_1 \in \widehat{\mathcal{K}}$ with probability at least $1 - \frac{\delta_2}{K}$. Subsequently, we can invoke Theorem 3.7 again to get that taking another step of the exact outer-loop NPG update starting from \mathbf{K}_1 , the updated gain matrix, denoted as $\widetilde{\mathbf{K}}_2$, satisfies $\mathbf{P}_{\widetilde{\mathbf{K}}_2, \mathbf{L}(\widetilde{\mathbf{K}}_2)} \leq \mathbf{P}_{\mathbf{K}_1, \mathbf{L}(\mathbf{K}_1)}$. Similarly, by requiring $\|\mathbf{K}_2 - \widetilde{\mathbf{K}}_2\|_F \leq \min\{\mathcal{B}_{1, \widetilde{\mathbf{K}}_2}, \mathcal{B}_{2, \widetilde{\mathbf{K}}_2}, \frac{\zeta}{2\mathcal{B}_{\mathbf{P}, \widetilde{\mathbf{K}}_2} K \|\mathbf{D}\|^2}\}$, we can guarantee with probability at least $1 - \frac{\delta_2}{K}$ that $\mathbf{K}_2 \in \mathcal{K}$, conditioned on $\mathbf{K}_1 \in \widehat{\mathcal{K}}$. Conditioning on (B.28), we apply Theorem 3.7 and Lemma B.7 again to get with probability at least $1 - \frac{\delta_2}{K}$ that $\mathbf{P}_{\mathbf{K}_2, \mathbf{L}(\mathbf{K}_2)} \geq 0$ exists and satisfies

$$\mathbf{P}_{\mathbf{K}_2, \mathbf{L}(\mathbf{K}_2)} \leq \mathbf{P}_{\widetilde{\mathbf{K}}_2, \mathbf{L}(\widetilde{\mathbf{K}}_2)} + \frac{\zeta}{2K\|\mathbf{D}\|^2} \cdot \mathbf{I} \leq \mathbf{P}_{\mathbf{K}_1, \mathbf{L}(\mathbf{K}_1)} + \frac{\zeta}{2K\|\mathbf{D}\|^2} \cdot \mathbf{I} \leq \mathbf{P}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)} + \frac{\zeta}{K\|\mathbf{D}\|^2} \cdot \mathbf{I}.$$

That is, $\mathbf{K}_2 \in \widehat{\mathcal{K}}$ with probability at least $1 - \frac{\delta_2}{K}$. Applying above arguments iteratively for all iterations and taking a union bound yield that if the estimation accuracy, for all $k \in \{1, \dots, K\}$, satisfies with probability at least $1 - \frac{\delta_2}{K}$ that

$$\|\mathbf{K}_k - \widetilde{\mathbf{K}}_k\|_F \leq \min \left\{ \mathcal{B}_{1, \widetilde{\mathcal{K}}_k}, \mathcal{B}_{2, \widetilde{\mathcal{K}}_k}, \frac{\zeta}{2\mathcal{B}_{P, \widetilde{\mathcal{K}}_k} K \|\mathbf{D}\|^2} \right\}, \quad (\text{B.29})$$

then it holds with probability at least $1 - \delta_2$ that $\mathbf{K}_K \in \widehat{\mathcal{K}} \subset \mathcal{K}$, where \mathbf{K}_K is the gain matrix after K steps of the ZO-NPG update (4.3) starting from \mathbf{K}_0 and with a constant stepsize α satisfying (B.27).

Lastly, we provide precise choices of the parameters for Algorithm 2 such that (B.29) can be satisfied. Since $\widehat{\mathcal{K}}$ is compact, there exist some uniform constants over $\widehat{\mathcal{K}}$

$$\mathcal{B}_1 := \min_{\mathbf{K} \in \widehat{\mathcal{K}}} \mathcal{B}_{1, \mathbf{K}} > 0, \quad \mathcal{B}_P := \max_{\mathbf{K} \in \widehat{\mathcal{K}}} \mathcal{B}_{P, \mathbf{K}} < \infty.$$

Moreover, let us denote the closure of the complement of \mathcal{K} as $\overline{\mathcal{K}^c}$. Since $\widehat{\mathcal{K}} \subset \mathcal{K}$ is compact and all $\mathbf{K} \in \widehat{\mathcal{K}}$ satisfy: i) $\mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})} \geq 0$ exists; and ii) $\lambda_{\min}(\mathbf{H}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}) \geq \frac{\zeta}{2} > 0$, $\widehat{\mathcal{K}}$ is disjoint from $\overline{\mathcal{K}^c}$, i.e., $\widehat{\mathcal{K}} \cap \overline{\mathcal{K}^c} = \emptyset$. Hence, one can deduce (see for example Lemma A.1 of (Berg, 2001)) that, there exists a distance $\mathcal{B}_2 > 0$ between $\widehat{\mathcal{K}}$ and $\overline{\mathcal{K}^c}$ such that $\mathcal{B}_2 \leq \min_{\mathbf{K} \in \widehat{\mathcal{K}}} \mathcal{B}_{2, \mathbf{K}}$. As a result, we now aim to enforce for all $k \in \{1, \dots, K\}$, that the following inequality holds with probability at least $1 - \frac{\delta_2}{K}$:

$$\|\mathbf{K}_k - \widetilde{\mathbf{K}}_k\|_F \leq \varpi := \min \left\{ \mathcal{B}_1, \mathcal{B}_2, \frac{\zeta}{2\mathcal{B}_P K \|\mathbf{D}\|^2} \right\}. \quad (\text{B.30})$$

Note that compared to (B.29), (B.30) is independent of the iteration index k , and is more stringent. Thus, (B.30) also ensures $\mathbf{K}_k \in \widehat{\mathcal{K}}$ for all $k \in \{1, \dots, K\}$ with probability at least $1 - \delta_2$, as (B.29) does. Condition (B.30) can be achieved by Lemmas B.10 and B.11. Specifically, we start from any gain matrix \mathbf{K} from the set of $\{\mathbf{K}_0, \dots, \mathbf{K}_{K-1}\}$ and use \mathbf{K}' and $\widetilde{\mathbf{K}}'$ to denote the gain matrices after one step of the ZO-NPG (4.3) and the exact NPG (3.11) updates, respectively. Now, suppose that the parameters of Algorithm 2 satisfy $\delta_1 \leq \delta_2/[6M_2K]$ and

$$\begin{aligned} M_2 \geq \max & \left\{ \frac{128d_2^2\alpha^2c_0^2\widehat{\mathcal{G}}(\mathbf{K}, \mathbf{L}(\mathbf{K}))^2}{r_2^2\phi^4\varpi^2}, \frac{512d_2^2\alpha^2(\widehat{\mathcal{G}}(\mathbf{K}, \mathbf{L}(\mathbf{K})) + r_2\mathcal{B}_P c_0)^2}{r_2^2\phi^2\varpi^2}, \frac{32\alpha^2(\widehat{c}_{5, \mathbf{K}})^2(\widehat{c}_{\Sigma, \mathbf{K}, \mathbf{L}(\mathbf{K})})^2}{\phi^2\varpi^2}, \frac{8(\widehat{c}_{\Sigma, \mathbf{K}, \mathbf{L}(\mathbf{K})})^2}{\phi^2} \right\} \\ & \cdot \log \left(\frac{12K \max\{d_2, d_\Sigma\}}{\delta_2} \right), \quad \epsilon_1 \leq \min \left\{ \frac{\phi\varpi r_2}{16\alpha d_2}, \frac{\widehat{\mathcal{B}}_{1, \mathbf{L}(\mathbf{K})}\zeta}{2}, \frac{\phi^2\varpi^2\zeta}{128\alpha^2(\widehat{c}_{5, \mathbf{K}})^2\widehat{\mathcal{B}}_{\Sigma, \mathbf{L}(\mathbf{K})}^2}, \frac{\phi^2\zeta}{32\widehat{\mathcal{B}}_{\Sigma, \mathbf{L}(\mathbf{K})}^2} \right\}, \\ r_2 \leq \min & \left\{ \varpi, \frac{\phi\varpi}{64\alpha\widehat{c}_{5, \mathbf{K}}(\widehat{c}_{\Sigma, \mathbf{K}, \mathbf{L}(\mathbf{K})} + \widehat{\mathcal{B}}_{\Sigma, \mathbf{K}})}, \frac{\phi\varpi}{64\alpha\widehat{c}_{2, \mathbf{K}}\widehat{\mathcal{B}}_{\Sigma, \mathbf{K}}} \right\}, \end{aligned}$$

with the requirement on ϵ_1 uses the fact that all $\mathbf{K} \in \widehat{\mathcal{K}}$ satisfy $\lambda_{\min}(\mathbf{H}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}) \geq \frac{\zeta}{2}$. Also, $\widehat{\mathcal{G}}(\mathbf{K}, \mathbf{L}(\mathbf{K}))$, $\widehat{c}_{2, \mathbf{K}}$, $\widehat{c}_{5, \mathbf{K}}$, $\widehat{c}_{\Sigma, \mathbf{K}, \mathbf{L}(\mathbf{K})}$, $\widehat{\mathcal{B}}_{\Sigma, \mathbf{K}}$, $\widehat{\mathcal{B}}_{1, \mathbf{L}(\mathbf{K})}$, $\widehat{\mathcal{B}}_{\Sigma, \mathbf{L}(\mathbf{K})} > 0$ are uniform constants over $\widehat{\mathcal{K}}$ such that

$$\begin{aligned} \widehat{\mathcal{G}}(\mathbf{K}, \mathbf{L}(\mathbf{K})) &:= \max_{\mathbf{K} \in \widehat{\mathcal{K}}} \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})), \quad \widehat{\mathcal{B}}_{\Sigma, \mathbf{K}} := \max_{\mathbf{K} \in \widehat{\mathcal{K}}} \mathcal{B}_{\Sigma, \mathbf{K}}, \quad \widehat{\mathcal{B}}_{1, \mathbf{L}(\mathbf{K})} := \min_{\mathbf{K} \in \widehat{\mathcal{K}}} \mathcal{B}_{1, \mathbf{L}(\mathbf{K})}, \\ \widehat{\mathcal{B}}_{\Sigma, \mathbf{L}(\mathbf{K})} &:= \max_{\mathbf{K} \in \widehat{\mathcal{K}}} \mathcal{B}_{\Sigma, \mathbf{L}(\mathbf{K})}, \quad \widehat{c}_{2, \mathbf{K}} := \max_{\mathbf{K} \in \widehat{\mathcal{K}}} c_{2, \mathbf{K}}, \quad \widehat{c}_{5, \mathbf{K}} := \max_{\mathbf{K} \in \widehat{\mathcal{K}}} c_{5, \mathbf{K}}, \quad \widehat{c}_{\Sigma, \mathbf{K}, \mathbf{L}(\mathbf{K})} := \max_{\mathbf{K} \in \widehat{\mathcal{K}}} c_{\Sigma, \mathbf{K}, \mathbf{L}(\mathbf{K})}, \end{aligned}$$

where $\mathcal{B}_{\Sigma, \mathbf{K}}$ is from Lemma B.7, $\mathcal{B}_{1, \mathbf{L}(\mathbf{K})}$, $\mathcal{B}_{\Sigma, \mathbf{L}(\mathbf{K})}$ are from Lemma C.2, $c_{2, \mathbf{K}}$, $c_{5, \mathbf{K}}$ are defined in §C.9, and $c_{\Sigma, \mathbf{K}, \mathbf{L}(\mathbf{K})}$ is from Lemma B.8. Then, Lemma B.10 proves that the following inequality holds with probability at least $1 - \frac{\delta_2}{2K}$:

$$\|\bar{\nabla}_{\mathbf{K}} \mathcal{G}(\mathbf{K}, \bar{\mathbf{L}}(\mathbf{K})) - \nabla_{\mathbf{K}} \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K}))\|_F \leq \frac{\phi\varpi}{4\alpha}. \quad (\text{B.31})$$

Moreover, we invoke Lemma B.11 to get with probability at least $1 - \frac{\delta_2}{2K}$ that

$$\|\bar{\Sigma}_{\mathbf{K}, \bar{\mathbf{L}}(\mathbf{K})} - \Sigma_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F \leq \min \left\{ \frac{\phi\varpi}{4\alpha\widehat{c}_{5, \mathbf{K}}}, \frac{\phi}{2} \right\} \leq \min \left\{ \frac{\phi\varpi}{4\alpha c_{5, \mathbf{K}}}, \frac{\phi}{2} \right\}.$$

By matrix perturbation theory (see for example Theorem 35 of (Fazel et al., 2018)) and noting that $\bar{\Sigma}_{\mathbf{K}, \bar{\mathbf{L}}(\mathbf{K})} = \Sigma_{\mathbf{K}, \mathbf{L}(\mathbf{K})} + (\bar{\Sigma}_{\mathbf{K}, \bar{\mathbf{L}}(\mathbf{K})} - \Sigma_{\mathbf{K}, \mathbf{L}(\mathbf{K})})$, since $\Sigma_{\mathbf{K}, \mathbf{L}(\mathbf{K})} \geq \Sigma_0$ and $\|\bar{\Sigma}_{\mathbf{K}, \bar{\mathbf{L}}(\mathbf{K})} - \Sigma_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F \leq \frac{\phi}{2}$, we have with probability at least $1 - \frac{\delta_2}{2K}$ that

$$\|\bar{\Sigma}_{\mathbf{K}, \bar{\mathbf{L}}(\mathbf{K})}^{-1} - \Sigma_{\mathbf{K}, \mathbf{L}(\mathbf{K})}^{-1}\|_F \leq \frac{2\|\bar{\Sigma}_{\mathbf{K}, \bar{\mathbf{L}}(\mathbf{K})} - \Sigma_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F}{\phi} \leq \frac{\varpi}{2\alpha c_{5, \mathbf{K}}}. \quad (\text{B.32})$$

Lastly, by Lemma B.11 and $\|\bar{\Sigma}_{\mathbf{K}, \bar{\mathbf{L}}(\mathbf{K})} - \Sigma_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F \leq \frac{\phi}{2}$, we can ensure that $\lambda_{\min}(\bar{\Sigma}_{\mathbf{K}, \bar{\mathbf{L}}(\mathbf{K})}) \geq \frac{\phi}{2}$ and thus $\|\bar{\Sigma}_{\mathbf{K}, \bar{\mathbf{L}}(\mathbf{K})}^{-1}\|_F \leq \frac{2}{\phi}$, all with probability at least $1 - \frac{\delta_2}{2K}$. Then, we combine (B.31) and (B.32) to show that with probability at least $1 - \frac{\delta_2}{K}$,

$$\begin{aligned} \|\widetilde{\mathbf{K}}' - \mathbf{K}'\|_F &= \alpha \|\bar{\nabla}_{\mathbf{K}} \mathcal{G}(\mathbf{K}, \bar{\mathbf{L}}(\mathbf{K})) \bar{\Sigma}_{\mathbf{K}, \bar{\mathbf{L}}(\mathbf{K})}^{-1} - \nabla_{\mathbf{K}} \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) \Sigma_{\mathbf{K}, \mathbf{L}(\mathbf{K})}^{-1}\|_F \\ &\leq \alpha \|\bar{\nabla}_{\mathbf{K}} \mathcal{G}(\mathbf{K}, \bar{\mathbf{L}}(\mathbf{K})) - \nabla_{\mathbf{K}} \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K}))\|_F \|\bar{\Sigma}_{\mathbf{K}, \bar{\mathbf{L}}(\mathbf{K})}^{-1}\|_F + \alpha \|\nabla_{\mathbf{K}} \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K}))\|_F \|\bar{\Sigma}_{\mathbf{K}, \bar{\mathbf{L}}(\mathbf{K})}^{-1} - \Sigma_{\mathbf{K}, \mathbf{L}(\mathbf{K})}^{-1}\|_F \\ &\leq \frac{2\alpha}{\phi} \cdot \|\bar{\nabla}_{\mathbf{K}} \mathcal{G}(\mathbf{K}, \bar{\mathbf{L}}(\mathbf{K})) - \nabla_{\mathbf{K}} \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K}))\|_F + \alpha c_{5, \mathbf{K}} \cdot \|\bar{\Sigma}_{\mathbf{K}, \bar{\mathbf{L}}(\mathbf{K})}^{-1} - \Sigma_{\mathbf{K}, \mathbf{L}(\mathbf{K})}^{-1}\|_F \leq \varpi, \end{aligned}$$

where the second inequality uses (C.16). Therefore, the above choices of parameters fulfill the requirements in (B.30), and thus guarantee $\mathbf{K}_k \in \hat{\mathcal{K}}$ for all $k \in \{1, \dots, K\}$ with probability at least $1 - \delta_2$. This completes the proof. \blacksquare

B.15. Proof of Theorem 4.3

Proof We provide a complete version of Theorem 4.3 as follows.

Theorem B.12 (Outer-Loop Sample Complexity) For any $\mathbf{K}_0 \in \mathcal{K}$, let $\epsilon_2, \delta_2, \epsilon_1, \delta_1, M_2, r_2, \alpha$ in Algorithm 2 satisfy the requirements in Theorem B.6, but with ϖ therein replaced by $\varphi := \min\{\varpi, \alpha\phi\epsilon_2/[\sqrt{sc_0}\mathcal{B}_{\mathbf{P}}]\}$. Then, it holds with probability at least $1 - \delta_2$ that the sequence $\{\mathbf{K}_k\}$, $k \in \{0, \dots, K\}$, converges with $\mathcal{O}(1/K)$ rate in the sense that $K^{-1} \sum_{k=0}^{K-1} \|\mathbf{F}_{\mathbf{K}_k, \mathbf{L}(\mathbf{K}_k)}\|_F^2 \leq \epsilon_2$ with $K = \text{Tr}(\mathbf{P}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)} - \mathbf{P}_{\mathbf{K}^*, \mathbf{L}(\mathbf{K}^*)})/[\alpha\phi\epsilon_2]$, which is also towards the unique Nash equilibrium.

Then, we focus on proving Theorem B.12. First, we require the parameters in Algorithm 2 to satisfy (B.22). Then, Theorem 4.2 proves with probability at least $1 - \delta_2$ that $\mathbf{K}_k \in \hat{\mathcal{K}}$, for all $k \in \{1, \dots, K\}$. Subsequently, we characterize the convergence rate of the outer-loop ZO-NPG update (4.3), conditioned on $\mathbf{K}_k \in \hat{\mathcal{K}}$, for all $k \in \{1, \dots, K\}$. Once again, we start from \mathbf{K}_0 and characterize the one-step progress of (4.3) from \mathbf{K}_0 to \mathbf{K}_1 . In addition to the requirement in (B.29), we now require $\|\mathbf{K}_1 - \widetilde{\mathbf{K}}_1\|_F \leq \frac{\alpha\phi\epsilon_2}{\sqrt{sc_0}\mathcal{B}_{\mathbf{P}, \widetilde{\mathbf{K}}_1}}$ with probability at least $1 - \frac{\delta_2}{K}$, where $\widetilde{\mathbf{K}}_1$ is the gain matrix after one step of (3.11) starting from \mathbf{K}_0 and with a constant stepsize α satisfying (B.27). Invoking Lemma B.7 yields with probability at least $1 - \frac{\delta_2}{K}$ that

$$\text{Tr}(\mathbf{P}_{\mathbf{K}_1, \mathbf{L}(\mathbf{K}_1)} - \mathbf{P}_{\widetilde{\mathbf{K}}_1, \mathbf{L}(\widetilde{\mathbf{K}}_1)}) \leq \sqrt{s} \|\mathbf{P}_{\mathbf{K}_1, \mathbf{L}(\mathbf{K}_1)} - \mathbf{P}_{\widetilde{\mathbf{K}}_1, \mathbf{L}(\widetilde{\mathbf{K}}_1)}\|_F \leq \sqrt{s} \mathcal{B}_{\mathbf{P}, \widetilde{\mathbf{K}}_1} \cdot \|\mathbf{K}_1 - \widetilde{\mathbf{K}}_1\|_F \leq \frac{\alpha\phi\epsilon_2}{c_0}, \quad (\text{B.33})$$

where $\mathbf{P}_{\mathbf{K}_1, \mathbf{L}(\mathbf{K}_1)} - \mathbf{P}_{\widetilde{\mathbf{K}}_1, \mathbf{L}(\widetilde{\mathbf{K}}_1)} \in \mathbb{R}^{s \times s}$ is symmetric. By (B.17), we have with probability at least $1 - \frac{\delta_2}{K}$ that

$$\begin{aligned} \text{Tr}((\mathbf{P}_{\mathbf{K}_1, \mathbf{L}(\mathbf{K}_1)} - \mathbf{P}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)})\Sigma_0) &= \text{Tr}((\mathbf{P}_{\mathbf{K}_1, \mathbf{L}(\mathbf{K}_1)} - \mathbf{P}_{\widetilde{\mathbf{K}}_1, \mathbf{L}(\widetilde{\mathbf{K}}_1)})\Sigma_0) + \text{Tr}((\mathbf{P}_{\widetilde{\mathbf{K}}_1, \mathbf{L}(\widetilde{\mathbf{K}}_1)} - \mathbf{P}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)})\Sigma_0) \\ &\leq \frac{\alpha\phi\epsilon_2}{c_0} \cdot c_0 - 2\alpha\phi \text{Tr}(\mathbf{F}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)}^\top \mathbf{F}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)}) \leq -\alpha\phi \text{Tr}(\mathbf{F}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)}^\top \mathbf{F}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)}), \end{aligned}$$

where the first inequality is due to $\|\Sigma_0\|_F \leq c_0$ almost surely and the last inequality follows from our assumption that $\text{Tr}(\mathbf{F}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)}^\top \mathbf{F}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)}) \geq \epsilon_2$. Similarly, for all future iterations $k \in \{2, \dots, K\}$, we can require $\|\mathbf{K}_k - \widetilde{\mathbf{K}}_k\|_F \leq \frac{\alpha\phi\epsilon_2}{\sqrt{sc_0}\mathcal{B}_{\mathbf{P}, \widetilde{\mathbf{K}}_k}}$ with probability at least $1 - \frac{\delta_2}{K}$ to obtain $\text{Tr}((\mathbf{P}_{\mathbf{K}_k, \mathbf{L}(\mathbf{K}_k)} - \mathbf{P}_{\mathbf{K}_{k-1}, \mathbf{L}(\mathbf{K}_{k-1})})\Sigma_0) \leq -\alpha\phi \text{Tr}(\mathbf{F}_{\mathbf{K}_{k-1}, \mathbf{L}(\mathbf{K}_{k-1})}^\top \mathbf{F}_{\mathbf{K}_{k-1}, \mathbf{L}(\mathbf{K}_{k-1})})$ until the update converges to the ϵ_2 -stationary point of the outer-loop, which

is also the unique Nash equilibrium. Summing up all iterations and taking a union bound yield with probability at least $1 - \delta_2$ the sublinear convergence rate, in that

$$\frac{1}{K} \sum_{k=0}^{K-1} \text{Tr}(\mathbf{F}_{\mathbf{K}_k, \mathbf{L}(\mathbf{K}_k)}^\top \mathbf{F}_{\mathbf{K}_k, \mathbf{L}(\mathbf{K}_k)}) \leq \frac{\text{Tr}(\mathbf{P}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)} - \mathbf{P}_{\mathbf{K}^*, \mathbf{L}(\mathbf{K}^*)})}{\alpha\phi \cdot K}. \quad (\text{B.34})$$

That is, when $K = \text{Tr}(\mathbf{P}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)} - \mathbf{P}_{\mathbf{K}^*, \mathbf{L}(\mathbf{K}^*)})/[\alpha\phi\epsilon_2]$, it satisfies with probability at least $1 - \delta_2$ that $K^{-1} \sum_{k=0}^{K-1} \|\mathbf{F}_{\mathbf{K}_k, \mathbf{L}(\mathbf{K}_k)}\|_F^2 \leq \epsilon_2$.

Lastly, the precise choices of the parameters for Algorithm 2 such that the above requirements on the estimation accuracy can be satisfied follow from the proof of Theorem 4.2 in §B.14, with additionally $\epsilon_2 \leq \phi/2$, and ϖ in (B.30) replaced by

$$\varphi := \min \left\{ \varpi, \frac{\alpha\phi\epsilon_2}{\sqrt{sc_0}\mathcal{B}_P} \right\}, \quad (\text{B.35})$$

where $\mathcal{B}_P = \max_{\mathbf{K} \in \hat{\mathcal{K}}} \mathcal{B}_{P, \mathbf{K}}$ is a uniform constant. Note that (B.35) is again independent of the iteration index k , and is more stringent than (B.30). Thus, (B.35) also ensures $\mathbf{K}_k \in \hat{\mathcal{K}}$ for all $k \in \{1, \dots, K\}$ with probability at least $1 - \delta_2$, as (B.30) does. By Lemmas B.10 and B.11, requirements on the parameters in (B.22) with ϖ replaced by φ suffice to guarantee that $\mathbf{K}_k \in \hat{\mathcal{K}}$ for all $k \in \{1, \dots, K\}$ with probability at least $1 - \delta_2$. Also, it ensures with probability at least $1 - \delta_2$ that the convergence in (B.34) holds. This completes the proof. ■

C. Supplementary Results

C.1. Proof of Lemma B.1

Proof Firstly, we have $R_t^w - D_t^\top P_{K_{t+1}, L(K_{t+1})} D_t$, $R_t^w - D_t^\top P_{K'_{t+1}, L(K'_{t+1})} D_t$ invertible since the conditions

$$R_t^w - D_t^\top P_{K_{t+1}, L(K_{t+1})} D_t > 0, \quad R_t^w - D_t^\top P_{K'_{t+1}, L(K'_{t+1})} D_t > 0$$

are satisfied. Then, by the definition of $P_{K, L}$ in (3.1) and the definition of $L(K)$, we can derive

$$\begin{aligned} & P_{K'_t, L(K'_t)} - P_{K_t, L(K_t)} - A_{K'_t, L(K'_t)}^\top (P_{K'_{t+1}, L(K'_{t+1})} - P_{K_{t+1}, L(K_{t+1})}) A_{K'_t, L(K'_t)} \\ &= A_{K'_t, L(K'_t)}^\top P_{K_{t+1}, L(K_{t+1})} A_{K'_t, L(K'_t)} - P_{K_t, L(K_t)} + P_{K'_t, L(K'_t)} - A_{K'_t, L(K'_t)}^\top P_{K'_{t+1}, L(K'_{t+1})} A_{K'_t, L(K'_t)} \\ &= (A_t - B_t K'_t)^\top P_{K_{t+1}, L(K_{t+1})} (A_t - B_t K'_t) - L(K'_t)^\top D_t^\top P_{K_{t+1}, L(K_{t+1})} (A_t - B_t K'_t) - (A_t - B_t K'_t)^\top P_{K_{t+1}, L(K_{t+1})} D_t L(K'_t) \\ &\quad - P_{K_t, L(K_t)} + Q_t + (K'_t)^\top R_t^u(K'_t) - L(K'_t)^\top (R_t^w - D_t^\top P_{K_{t+1}, L(K_{t+1})} D_t) L(K'_t) \\ &= (A_t - B_t K'_t)^\top \tilde{P}_{K_{t+1}, L(K_{t+1})} (A_t - B_t K'_t) \\ &\quad - (A_t - B_t K'_t)^\top P_{K_{t+1}, L(K_{t+1})} D_t (R_t^w - D_t^\top P_{K_{t+1}, L(K_{t+1})} D_t)^{-1} D_t^\top P_{K_{t+1}, L(K_{t+1})} (A_t - B_t K'_t) - P_{K_t, L(K_t)} + Q_t \\ &\quad + (K'_t)^\top R_t^u(K'_t) - L(K'_t)^\top D_t^\top P_{K_{t+1}, L(K_{t+1})} (A_t - B_t K'_t) - (A_t - B_t K'_t)^\top P_{K_{t+1}, L(K_{t+1})} D_t L(K'_t) \\ &\quad - L(K'_t)^\top (R_t^w - D_t^\top P_{K_{t+1}, L(K_{t+1})} D_t) L(K'_t) \\ &= \mathcal{R}_{K_t, K'_t} - (A_t - B_t K'_t)^\top P_{K_{t+1}, L(K_{t+1})} D_t (R_t^w - D_t^\top P_{K_{t+1}, L(K_{t+1})} D_t)^{-1} D_t^\top P_{K_{t+1}, L(K_{t+1})} (A_t - B_t K'_t) \\ &\quad - L(K'_t)^\top D_t^\top P_{K_{t+1}, L(K_{t+1})} (A_t - B_t K'_t) - (A_t - B_t K'_t)^\top P_{K_{t+1}, L(K_{t+1})} D_t L(K'_t) \\ &\quad - L(K'_t)^\top (R_t^w - D_t^\top P_{K_{t+1}, L(K_{t+1})} D_t) L(K'_t) \\ &= \mathcal{R}_{K_t, K'_t} - \left[- (R_t^w - D_t^\top P_{K_{t+1}, L(K_{t+1})} D_t) L(K'_t) - D_t^\top P_{K_{t+1}, L(K_{t+1})} (A_t - B_t K'_t) \right]^\top (R_t^w - D_t^\top P_{K_{t+1}, L(K_{t+1})} D_t)^{-1} \\ &\quad \cdot \left[- (R_t^w - D_t^\top P_{K_{t+1}, L(K_{t+1})} D_t) L(K'_t) - D_t^\top P_{K_{t+1}, L(K_{t+1})} (A_t - B_t K'_t) \right] \\ &= \mathcal{R}_{K_t, K'_t} - \Xi_{K_t, K'_t}^\top (R_t^w - D_t^\top P_{K_{t+1}, L(K_{t+1})} D_t)^{-1} \Xi_{K_t, K'_t}. \end{aligned}$$

This completes the proof of (B.11). Further, we invoke (3.1) to represent \mathcal{R}_{K_t, K'_t} as

$$\begin{aligned}
 \mathcal{R}_{K_t, K'_t} &= (A_t - B_t K'_t)^\top \tilde{P}_{K_{t+1}, L(K_{t+1})} (A_t - B_t K'_t) - P_{K_t, L(K_t)} + Q_t + (K'_t)^\top R_t^u (K'_t) \\
 &= (K'_t - K_t)^\top \left((R_t^u + B_t^\top \tilde{P}_{K_{t+1}, L(K_{t+1})} B_t) K_t - B_t^\top \tilde{P}_{K_{t+1}, L(K_{t+1})} A_t \right) \\
 &\quad + (K'_t - K_t)^\top (R_t^u + B_t^\top \tilde{P}_{K_{t+1}, L(K_{t+1})} B_t) (K'_t - K_t) \\
 &\quad + \left((R_t^u + B_t^\top \tilde{P}_{K_{t+1}, L(K_{t+1})} B_t) K_t - B_t^\top \tilde{P}_{K_{t+1}, L(K_{t+1})} A_t \right)^\top (K'_t - K_t) \\
 &= (K'_t - K_t)^\top F_{K_t, L(K_t)} + F_{K_t, L(K_t)}^\top (K'_t - K_t) + (K'_t - K_t)^\top (R_t^u + B_t^\top \tilde{P}_{K_{t+1}, L(K_{t+1})} B_t) (K'_t - K_t).
 \end{aligned}$$

This completes the proof. \blacksquare

C.2. Proof of Lemma B.3

Proof Firstly, for a fixed $\mathbf{K} \in \mathcal{K}$ and a given \mathbf{L} , we require $r_{1,L} \leq \rho_{\mathbf{K},L}$ to ensure that the “size” of the perturbation added to the gain matrix \mathbf{L} is smaller than the radius within which the local Lipschitz continuity and the local smoothness properties in Lemma 3.4 hold. In other words, we have $\|r_{1,L} \mathbf{U}\|_F \leq \rho_{\mathbf{K},L}$, where \mathbf{U} is drawn uniformly from $\mathcal{S}(n, m, N)$ and has $\|\mathbf{U}\|_F = 1$. Then, we invoke Lemma 23 of (Malik et al., 2020) to obtain that with probability at least $1 - \delta_1$, the concentration of gradient estimates around its mean is $\|\bar{\nabla}_L \mathcal{G}(\mathbf{K}, \mathbf{L}) - \nabla_L \mathcal{G}_{r_{1,L}}(\mathbf{K}, \mathbf{L})\|_F \leq \frac{d_1}{r_{1,L} \sqrt{M_{1,L}}} \left(\mathcal{G}(\mathbf{K}, \mathbf{L}) + \frac{l_{\mathbf{K},L}}{\rho_{\mathbf{K},L}} \right) \sqrt{\log \left(\frac{2d_1}{\delta_1} \right)}$, where $\bar{\nabla}_L \mathcal{G}(\mathbf{K}, \mathbf{L})$ is the gradient estimate obtained from an $M_{1,L}$ -sample one-point minibatch estimator, $\nabla_L \mathcal{G}_{r_{1,L}}(\mathbf{K}, \mathbf{L}) := \nabla_L \mathbb{E}[\mathcal{G}(\mathbf{K}, \mathbf{L} + r_{1,L} \mathbf{U})]$ is the gradient of the smoothed version of $\mathcal{G}(\mathbf{K}, \mathbf{L})$, and $d_1 = nmN$ is the degrees of freedom of $\mathcal{S}(n, m, N)$. Therefore, it suffices to require the batchsize $M_{1,L} > 0$ to satisfy $M_{1,L} \geq \left(\frac{d_1}{r_{1,L}} \left(\mathcal{G}(\mathbf{K}, \mathbf{L}) + \frac{l_{\mathbf{K},L}}{\rho_{\mathbf{K},L}} \right) \sqrt{\log \left(\frac{2d_1}{\delta_1} \right)} \right)^2 \frac{1024}{\mu_{\mathbf{K}} \epsilon_1}$ in order to ensure that with probability at least $1 - \delta_1$ that

$$\|\bar{\nabla}_L \mathcal{G}(\mathbf{K}, \mathbf{L}) - \nabla_L \mathcal{G}_{r_{1,L}}(\mathbf{K}, \mathbf{L})\|_F \leq \frac{\sqrt{\mu_{\mathbf{K}} \epsilon_1}}{32}. \quad (\text{C.1})$$

Then, we can bound, with probability at least $1 - \delta_1$, the “size” of the one-step ZO-PG update (4.1) by

$$\begin{aligned}
 \|\mathbf{L}' - \mathbf{L}\|_F &= \|\eta_L \bar{\nabla}_L \mathcal{G}(\mathbf{K}, \mathbf{L})\|_F \\
 &\leq \eta_L \|\bar{\nabla}_L \mathcal{G}(\mathbf{K}, \mathbf{L}) - \nabla_L \mathcal{G}_{r_{1,L}}(\mathbf{K}, \mathbf{L})\|_F + \eta_L \|\nabla_L \mathcal{G}_{r_{1,L}}(\mathbf{K}, \mathbf{L}) - \nabla_L \mathcal{G}(\mathbf{K}, \mathbf{L})\|_F + \eta_L \|\nabla_L \mathcal{G}(\mathbf{K}, \mathbf{L})\|_F \\
 &\leq \eta_L \left(\frac{\sqrt{\mu_{\mathbf{K}} \epsilon_1}}{32} + \psi_{\mathbf{K},L} \sqrt{\epsilon_1} + l_{\mathbf{K},L} \right) \leq \eta_L \left(\frac{\sqrt{\mu_{\mathbf{K}}}}{32} + \psi_{\mathbf{K},L} + l_{\mathbf{K},L} \right),
 \end{aligned}$$

where the second to last inequality follows from (C.1), Lemma 14(b) of (Malik et al., 2020), Lemma 3.4, and $r_{1,L} \sim \Theta(\sqrt{\epsilon_1})$. The last inequality follows from our requirement that $\epsilon_1 < 1$. Therefore, it suffices to require $\eta_L \leq \rho_{\mathbf{K},L} \cdot \left(\frac{\sqrt{\mu_{\mathbf{K}}}}{32} + \psi_{\mathbf{K},L} + l_{\mathbf{K},L} \right)^{-1}$ to ensure that \mathbf{L}' lies within the radius that the local Lipschitz continuity and the local smoothness properties hold with probability at least $1 - \delta_1$. Now, we exploit the local smoothness property to derive that

$$\begin{aligned}
 \mathcal{G}(\mathbf{K}, \mathbf{L}) - \mathcal{G}(\mathbf{K}, \mathbf{L}') &\leq -\eta_L \langle \nabla_L \mathcal{G}(\mathbf{K}, \mathbf{L}), \bar{\nabla}_L \mathcal{G}(\mathbf{K}, \mathbf{L}) \rangle + \frac{\psi_{\mathbf{K},L} \eta_L^2}{2} \|\bar{\nabla}_L \mathcal{G}(\mathbf{K}, \mathbf{L})\|_F^2 \\
 &= -\eta_L \langle \nabla_L \mathcal{G}(\mathbf{K}, \mathbf{L}), \bar{\nabla}_L \mathcal{G}(\mathbf{K}, \mathbf{L}) - \nabla_L \mathcal{G}_{r_{1,L}}(\mathbf{K}, \mathbf{L}) \rangle - \eta_L \langle \nabla_L \mathcal{G}(\mathbf{K}, \mathbf{L}), \nabla_L \mathcal{G}_{r_{1,L}}(\mathbf{K}, \mathbf{L}) \rangle + \frac{\psi_{\mathbf{K},L} \eta_L^2}{2} \|\bar{\nabla}_L \mathcal{G}(\mathbf{K}, \mathbf{L})\|_F^2 \\
 &\leq \eta_L \|\nabla_L \mathcal{G}(\mathbf{K}, \mathbf{L})\|_F \|\bar{\nabla}_L \mathcal{G}(\mathbf{K}, \mathbf{L}) - \nabla_L \mathcal{G}_{r_{1,L}}(\mathbf{K}, \mathbf{L})\|_F - \eta_L \|\nabla_L \mathcal{G}(\mathbf{K}, \mathbf{L})\|_F^2 \\
 &\quad + \eta_L \psi_{\mathbf{K},L} r_{1,L} \|\nabla_L \mathcal{G}(\mathbf{K}, \mathbf{L})\|_F + \frac{\psi_{\mathbf{K},L} \eta_L^2}{2} \|\bar{\nabla}_L \mathcal{G}(\mathbf{K}, \mathbf{L})\|_F^2 \\
 &\leq -\frac{\eta_L}{2} \|\nabla_L \mathcal{G}(\mathbf{K}, \mathbf{L})\|_F^2 + \frac{\eta_L}{2} \|\bar{\nabla}_L \mathcal{G}(\mathbf{K}, \mathbf{L}) - \nabla_L \mathcal{G}_{r_{1,L}}(\mathbf{K}, \mathbf{L})\|_F^2 \\
 &\quad + \eta_L \psi_{\mathbf{K},L} r_{1,L} \|\nabla_L \mathcal{G}(\mathbf{K}, \mathbf{L})\|_F + \frac{\psi_{\mathbf{K},L} \eta_L^2}{2} \|\bar{\nabla}_L \mathcal{G}(\mathbf{K}, \mathbf{L})\|_F^2.
 \end{aligned} \quad (\text{C.2})$$

Moreover, we can bound the last term of (C.2) as

$$\frac{\psi_{\mathbf{K},\mathbf{L}}\eta_{\mathbf{L}}^2}{2}\|\bar{\nabla}_{\mathbf{L}}\mathcal{G}(\mathbf{K},\mathbf{L})\|_F^2 \leq \psi_{\mathbf{K},\mathbf{L}}\eta_{\mathbf{L}}^2\|\bar{\nabla}_{\mathbf{L}}\mathcal{G}(\mathbf{K},\mathbf{L}) - \nabla_{\mathbf{L}}\mathcal{G}_{r_{1,\mathbf{L}}}(\mathbf{K},\mathbf{L})\|_F^2 + 2\psi_{\mathbf{K},\mathbf{L}}\eta_{\mathbf{L}}^2(\psi_{\mathbf{K},\mathbf{L}}^2r_1^2 + \|\nabla_{\mathbf{L}}\mathcal{G}(\mathbf{K},\mathbf{L})\|_F^2). \quad (\text{C.3})$$

Substituting (C.3) into (C.2) yields that the following inequality holds almost surely:

$$\begin{aligned} \mathcal{G}(\mathbf{K},\mathbf{L}) - \mathcal{G}(\mathbf{K},\mathbf{L}') &\leq \left(-\frac{\eta_{\mathbf{L}}}{2} + 2\psi_{\mathbf{K},\mathbf{L}}\eta_{\mathbf{L}}^2\right)\|\nabla_{\mathbf{L}}\mathcal{G}(\mathbf{K},\mathbf{L})\|_F^2 + \left(\frac{\eta_{\mathbf{L}}}{2} + \psi_{\mathbf{K},\mathbf{L}}\eta_{\mathbf{L}}^2\right)\|\bar{\nabla}_{\mathbf{L}}\mathcal{G}(\mathbf{K},\mathbf{L}) - \nabla_{\mathbf{L}}\mathcal{G}_{r_{1,\mathbf{L}}}(\mathbf{K},\mathbf{L})\|_F^2 \\ &\quad + \eta_{\mathbf{L}}\psi_{\mathbf{K},\mathbf{L}}r_{1,\mathbf{L}}\|\nabla_{\mathbf{L}}\mathcal{G}(\mathbf{K},\mathbf{L})\|_F + 2\psi_{\mathbf{K},\mathbf{L}}^3\eta_{\mathbf{L}}^2r_{1,\mathbf{L}}^2. \end{aligned}$$

Next, recalling the definition that $\theta_{\mathbf{K},\mathbf{L}} = \min\left\{\frac{1}{2\psi_{\mathbf{K},\mathbf{L}}}, \frac{\rho_{\mathbf{K},\mathbf{L}}}{l_{\mathbf{K},\mathbf{L}}}\right\}$, we invoke the local smoothness property to obtain

$$\left(\theta_{\mathbf{K},\mathbf{L}} - \frac{\theta_{\mathbf{K},\mathbf{L}}^2\psi_{\mathbf{K},\mathbf{L}}}{2}\right)\|\nabla_{\mathbf{L}}\mathcal{G}(\mathbf{K},\mathbf{L})\|_F^2 \leq \mathcal{G}(\mathbf{K},\mathbf{L} + \theta_{\mathbf{K},\mathbf{L}}\nabla_{\mathbf{L}}\mathcal{G}(\mathbf{K},\mathbf{L})) - \mathcal{G}(\mathbf{K},\mathbf{L}) \leq \mathcal{G}(\mathbf{K},\mathbf{L}(\mathbf{K})) - \mathcal{G}(\mathbf{K},\mathbf{L}).$$

Define $\Delta := \mathcal{G}(\mathbf{K},\mathbf{L}(\mathbf{K})) - \mathcal{G}(\mathbf{K},\mathbf{L})$ and $\Delta' := \mathcal{G}(\mathbf{K},\mathbf{L}(\mathbf{K})) - \mathcal{G}(\mathbf{K},\mathbf{L}')$, then it holds almost surely that

$$\begin{aligned} \Delta' - \Delta &\leq \left(-\frac{\eta_{\mathbf{L}}}{2} + 2\psi_{\mathbf{K},\mathbf{L}}\eta_{\mathbf{L}}^2\right)\|\nabla_{\mathbf{L}}\mathcal{G}(\mathbf{K},\mathbf{L})\|_F^2 + \frac{2\eta_{\mathbf{L}}\psi_{\mathbf{K},\mathbf{L}}r_{1,\mathbf{L}}}{\theta_{\mathbf{K},\mathbf{L}}}\Delta^{1/2} \\ &\quad + \left(\frac{\eta_{\mathbf{L}}}{2} + \psi_{\mathbf{K},\mathbf{L}}\eta_{\mathbf{L}}^2\right)\|\bar{\nabla}_{\mathbf{L}}\mathcal{G}(\mathbf{K},\mathbf{L}) - \nabla_{\mathbf{L}}\mathcal{G}_{r_{1,\mathbf{L}}}(\mathbf{K},\mathbf{L})\|_F^2 + 2\psi_{\mathbf{K},\mathbf{L}}^3\eta_{\mathbf{L}}^2r_{1,\mathbf{L}}^2 \\ &\leq -\frac{\eta_{\mathbf{L}}\mu_{\mathbf{K}}}{4}\Delta + \frac{2\eta_{\mathbf{L}}\psi_{\mathbf{K},\mathbf{L}}r_{1,\mathbf{L}}}{\theta_{\mathbf{K},\mathbf{L}}}\Delta^{1/2} + \eta_{\mathbf{L}}\|\bar{\nabla}_{\mathbf{L}}\mathcal{G}(\mathbf{K},\mathbf{L}) - \nabla_{\mathbf{L}}\mathcal{G}_{r_{1,\mathbf{L}}}(\mathbf{K},\mathbf{L})\|_F^2 + 2\psi_{\mathbf{K},\mathbf{L}}^3\eta_{\mathbf{L}}^2r_{1,\mathbf{L}}^2 \\ &\leq -\frac{\eta_{\mathbf{L}}\mu_{\mathbf{K}}}{8}\Delta + \frac{8\eta_{\mathbf{L}}\psi_{\mathbf{K},\mathbf{L}}^2r_{1,\mathbf{L}}^2}{\mu_{\mathbf{K}}\theta_{\mathbf{K},\mathbf{L}}^2} + \eta_{\mathbf{L}}\|\bar{\nabla}_{\mathbf{L}}\mathcal{G}(\mathbf{K},\mathbf{L}) - \nabla_{\mathbf{L}}\mathcal{G}_{r_{1,\mathbf{L}}}(\mathbf{K},\mathbf{L})\|_F^2 + 2\psi_{\mathbf{K},\mathbf{L}}^3\eta_{\mathbf{L}}^2r_{1,\mathbf{L}}^2, \end{aligned}$$

where the last two inequalities utilize $\eta_{\mathbf{L}} \leq \frac{1}{8\psi_{\mathbf{K},\mathbf{L}}}$ and the $\mu_{\mathbf{K}}$ -PL condition in Lemma 3.4. Thus, if the stepsize $\eta_{\mathbf{L}}$ of the one-step ZO-PG update (4.1) and the smoothing radius $r_{1,\mathbf{L}}$ for the minibatch estimator further satisfy

$$\eta_{\mathbf{L}} \leq \min\left\{1, \frac{1}{8\psi_{\mathbf{K},\mathbf{L}}}\right\}, \quad r_{1,\mathbf{L}} \leq \frac{1}{8\psi_{\mathbf{K},\mathbf{L}}} \min\left\{\theta_{\mathbf{K},\mathbf{L}}\mu_{\mathbf{K}}\sqrt{\frac{\epsilon_1}{240}}, \frac{1}{\psi_{\mathbf{K},\mathbf{L}}}\sqrt{\frac{\epsilon_1\mu_{\mathbf{K}}}{30}}\right\},$$

then with probability at least $1 - \delta_1$, we can bound the one-step ascent as

$$\Delta' - \Delta \leq -\frac{\eta_{\mathbf{L}}\mu_{\mathbf{K}}}{8}\Delta + \eta_{\mathbf{L}}\|\bar{\nabla}_{\mathbf{L}}\mathcal{G}(\mathbf{K},\mathbf{L}) - \nabla_{\mathbf{L}}\mathcal{G}_{r_{1,\mathbf{L}}}(\mathbf{K},\mathbf{L})\|_F^2 + \frac{\eta_{\mathbf{L}}\mu_{\mathbf{K}}\epsilon_1}{60} \Rightarrow \Delta' \leq \left(1 - \frac{\eta_{\mathbf{L}}\mu_{\mathbf{K}}}{8}\right)\Delta + \eta_{\mathbf{L}}\frac{\mu_{\mathbf{K}}\epsilon_1}{16}, \quad (\text{C.4})$$

where the last inequality follows from (C.1). Lastly, by (C.4), we have $\mathcal{G}(\mathbf{K},\mathbf{L}) - \mathcal{G}(\mathbf{K},\mathbf{L}') = \Delta' - \Delta \leq -\frac{\eta_{\mathbf{L}}\mu_{\mathbf{K}}}{8}\Delta + \eta_{\mathbf{L}}\frac{\mu_{\mathbf{K}}\epsilon_1}{16}$. Therefore, it holds that $\mathcal{G}(\mathbf{K},\mathbf{L}) - \mathcal{G}(\mathbf{K},\mathbf{L}') \leq 0$ with probability at least $1 - \delta_1$ since by the implicit assumption that $\epsilon_1 \leq \Delta$. This completes the proof. \blacksquare

C.3. Proof of Lemma B.5

Proof We first introduce the following lemma and defer its proof to the end of this subsection.

Lemma C.1 ($\Sigma_{\mathbf{K},\mathbf{L}}$ Estimation) For any $\mathbf{K} \in \mathcal{K}$ and any $\mathbf{L} \in \mathcal{S}(n, m, N)$, let the batchsize M_1 in Algorithm 1 satisfy

$$M_1 \geq \frac{1}{2}c_{\Sigma_{\mathbf{K},\mathbf{L}}}^2\epsilon_1^{-2}\log\left(\frac{2d_{\Sigma}}{\delta_1}\right),$$

where $c_{\Sigma_{\mathbf{K},\mathbf{L}}}$ is a polynomial of $\|\mathbf{A}\|_F, \|\mathbf{B}\|_F, \|\mathbf{D}\|_F, \|\mathbf{K}\|_F, \|\mathbf{L}\|_F$, and is linear in c_0 . Then, it holds with probability at least $1 - \delta_1$ that $\|\bar{\Sigma}_{\mathbf{K},\mathbf{L}} - \Sigma_{\mathbf{K},\mathbf{L}}\|_F \leq \epsilon_1$, where $\bar{\Sigma}_{\mathbf{K},\mathbf{L}}$ is the estimated state correlation matrix obtained from Algorithm 1. Moreover, if $\epsilon_1 \leq \phi/2$, then it satisfies with probability at least $1 - \delta_1$ that $\lambda_{\min}(\bar{\Sigma}_{\mathbf{K},\mathbf{L}}) \geq \phi/2$.

Similar to the proof of Lemma B.3, we first require $r_{1,L} \leq \rho_{K,L}$ to ensure that perturbing L preserves the local Lipschitz and smoothness properties. Then, Lemma 23 of (Malik et al., 2020) suggests that if $M_{1,L} \geq \left(\frac{d_1}{r_{1,L}} \left(\mathcal{G}(K, L(K)) + \frac{l_{K,L}}{\rho_{K,L}} \right) \sqrt{\log \left(\frac{4d_1}{\delta_1} \right)} \right)^2 \frac{64(\kappa+1)^2}{\phi^2 \mu_K \epsilon_1}$, where $\kappa := c_{\Sigma_{K,L}} + \frac{\phi}{2}$ and $c_{\Sigma_{K,L}}$ follows the definition in Lemma C.1, then it holds with probability at least $1 - \frac{\delta_1}{2}$ that

$$\|\bar{\nabla}_L \mathcal{G}(K, L) - \nabla_L \mathcal{G}_{r_1}(K, L)\|_F \leq \frac{\phi \sqrt{\mu_K \epsilon_1}}{8(\kappa+1)}. \quad (\text{C.5})$$

Moreover, by Lemma C.1 and if we require $M_{1,L} \geq 2\kappa^2/\phi^2 \cdot \log(4d_\Sigma/\delta_1)$, then it holds with probability at least $1 - \frac{\delta_1}{2}$ that $\|\bar{\Sigma}_{K,L} - \Sigma_{K,L}\|_F \leq \frac{\phi}{2}$. By the standard matrix perturbation theory (see for example Theorem 35 of (Fazel et al., 2018)), we have

$$\|\bar{\Sigma}_{K,L}^{-1} - \Sigma_{K,L}^{-1}\|_F \leq \frac{2\|\bar{\Sigma}_{K,L} - \Sigma_{K,L}\|_F}{\phi} \leq 1. \quad (\text{C.6})$$

Furthermore, by Lemma C.1 and the condition that $\|\bar{\Sigma}_{K,L} - \Sigma_{K,L}\| \leq \frac{\phi}{2}$, we can ensure that $\lambda_{\min}(\bar{\Sigma}_{K,L}) \geq \frac{\phi}{2}$ and thus $\|\bar{\Sigma}_{K,L}^{-1}\| \leq \frac{2}{\phi}$. Combining (C.5) and (C.6) yields that with probability at least $1 - \delta_1$, the “size” of the one-step inner-loop ZO-NPG update (4.2) can be bounded by

$$\begin{aligned} \|L' - L\|_F &= \eta_L \|\bar{\nabla}_L \mathcal{G}(K, L) \bar{\Sigma}_{K,L}^{-1}\|_F \leq \eta_L \|\bar{\nabla}_L \mathcal{G}(K, L) \bar{\Sigma}_{K,L}^{-1} - \nabla_L \mathcal{G}(K, L) \Sigma_{K,L}^{-1}\|_F + \eta_L \|\nabla_L \mathcal{G}(K, L) \Sigma_{K,L}^{-1}\|_F \\ &\leq \eta_L \left[\|\bar{\nabla}_L \mathcal{G}(K, L) - \nabla_L \mathcal{G}(K, L)\|_F \|\bar{\Sigma}_{K,L}^{-1}\|_F + \|\nabla_L \mathcal{G}(K, L)\|_F \|\bar{\Sigma}_{K,L}^{-1} - \Sigma_{K,L}^{-1}\|_F + \|\nabla_L \mathcal{G}(K, L)\|_F \|\Sigma_{K,L}^{-1}\|_F \right] \\ &\leq \eta_L \left[\frac{\sqrt{\mu_K}}{4(\kappa+1)} + \frac{2\psi_{K,L}}{\phi} + l_{K,L} + \frac{\phi l_{K,L}}{2} \right], \end{aligned}$$

where the last inequality follows from (C.5), Lemma 14 of (Malik et al., 2020), $\epsilon_1 < 1$ and $r_{1,L} \sim \Theta(\sqrt{\epsilon_1})$. Therefore, it suffices to require $\eta_L \leq \rho_{K,L} \cdot \left[\frac{\sqrt{\mu_K}}{4(\kappa+1)} + \frac{2\psi_{K,L}}{\phi} + l_{K,L} + \frac{\phi l_{K,L}}{2} \right]^{-1}$ to ensure that L' lies within the radius that the Lipschitz and smoothness properties hold. Now, we exploit the smoothness property to derive that

$$\begin{aligned} \mathcal{G}(K, L) - \mathcal{G}(K, L') &\leq -\eta_L \langle \nabla_L \mathcal{G}(K, L), \bar{\nabla}_L \mathcal{G}(K, L) \bar{\Sigma}_{K,L}^{-1} \rangle + \frac{\psi_{K,L} \eta_L^2}{2} \|\bar{\nabla}_L \mathcal{G}(K, L) \bar{\Sigma}_{K,L}^{-1}\|_F^2 \\ &\leq -\eta_L \langle \nabla_L \mathcal{G}(K, L), \nabla_L \mathcal{G}(K, L) \bar{\Sigma}_{K,L}^{-1} \rangle - \eta_L \langle \nabla_L \mathcal{G}(K, L), \bar{\nabla}_L \mathcal{G}(K, L) \bar{\Sigma}_{K,L}^{-1} - \nabla_L \mathcal{G}(K, L) \bar{\Sigma}_{K,L}^{-1} \rangle \\ &\quad + \frac{\psi_{K,L} \eta_L^2}{2} \|\bar{\nabla}_L \mathcal{G}(K, L) \bar{\Sigma}_{K,L}^{-1}\|_F^2 \\ &\leq -\eta_L \lambda_{\min}(\bar{\Sigma}_{K,L}) \|\nabla_L \mathcal{G}(K, L)\|_F^2 + \eta_L \|\nabla_L \mathcal{G}(K, L)\|_F \|\bar{\nabla}_L \mathcal{G}(K, L) - \nabla_L \mathcal{G}(K, L)\|_F \|\bar{\Sigma}_{K,L}^{-1}\|_F \\ &\quad + \frac{\psi_{K,L} \eta_L^2}{2} \|\bar{\Sigma}_{K,L}^{-1}\|_F^2 \|\bar{\nabla}_L \mathcal{G}(K, L)\|_F^2 \\ &\leq -\eta_L \lambda_{\min}(\bar{\Sigma}_{K,L}) \|\nabla_L \mathcal{G}(K, L)\|_F^2 + \frac{\eta_L}{2} \lambda_{\min}(\bar{\Sigma}_{K,L}) \|\nabla_L \mathcal{G}(K, L)\|_F^2 + \frac{\eta_L \|\bar{\Sigma}_{K,L}^{-1}\|_F^2}{2\lambda_{\min}(\bar{\Sigma}_{K,L})} \|\bar{\nabla}_L \mathcal{G}(K, L) - \nabla_L \mathcal{G}_{r_1}(K, L)\|_F^2 \\ &\quad + \psi_{K,L} r_{1,L} \|\nabla_L \mathcal{G}(K, L)\|_F \|\bar{\Sigma}_{K,L}^{-1}\|_F + \frac{\psi_{K,L} \eta_L^2 \|\bar{\Sigma}_{K,L}^{-1}\|_F^2}{2} \|\bar{\nabla}_L \mathcal{G}(K, L)\|_F^2 \\ &\leq -\frac{\eta_L}{2\kappa} \|\nabla_L \mathcal{G}(K, L)\|_F^2 + \frac{2\eta_L \kappa}{\phi^2} \|\bar{\nabla}_L \mathcal{G}(K, L) - \nabla_L \mathcal{G}_{r_1}(K, L)\|_F^2 + \frac{2\eta_L \psi_{K,L} r_{1,L}}{\phi} \|\nabla_L \mathcal{G}(K, L)\|_F \\ &\quad + \frac{2\psi_{K,L} \eta_L^2}{\phi^2} \|\bar{\nabla}_L \mathcal{G}(K, L)\|_F^2, \end{aligned} \quad (\text{C.7})$$

where the second to last inequality utilizes Lemma 14 of (Malik et al., 2020) and

$$\begin{aligned} \eta_L \|\nabla_L \mathcal{G}(K, L)\|_F \|\bar{\nabla}_L \mathcal{G}(K, L) - \nabla_L \mathcal{G}(K, L)\|_F \|\bar{\Sigma}_{K,L}^{-1}\|_F &\leq \frac{\eta_L \lambda_{\min}(\bar{\Sigma}_{K,L})}{2} \|\nabla_L \mathcal{G}(K, L)\|_F^2 \\ &\quad + \frac{\eta_L \|\bar{\Sigma}_{K,L}^{-1}\|_F^2}{2\lambda_{\min}(\bar{\Sigma}_{K,L})} \|\bar{\nabla}_L \mathcal{G}(K, L) - \nabla_L \mathcal{G}_{r_1}(K, L)\|_F^2. \end{aligned}$$

The last inequality in (C.7) uses $\lambda_{\min}(\bar{\Sigma}_{K,L}^{-1}) = \|\bar{\Sigma}_{K,L}\|^{-1} \geq [\|\Sigma_{K,L}\|_F + \|\bar{\Sigma}_{K,L} - \Sigma_{K,L}\|_F]^{-1}$ and $\|\bar{\Sigma}_{K,L}^{-1}\| \leq \frac{2}{\phi}$. Moreover, we can bound the last term in (C.7) as

$$\frac{2\psi_{K,L}\eta_L^2}{\phi^2} \|\bar{\nabla}_L \mathcal{G}(K, L)\|_F^2 \leq \frac{4\psi_{K,L}\eta_L^2}{\phi^2} \|\bar{\nabla}_L \mathcal{G}(K, L) - \nabla_L \mathcal{G}_{r_1}(K, L)\|_F^2 + \frac{8\psi_{K,L}\eta_L^2}{\phi^2} (\psi_{K,L}^2 r_{1,L}^2 + \|\nabla_L \mathcal{G}(K, L)\|_F^2). \quad (\text{C.8})$$

Further, recalling the definition that $\theta_{K,L} = \min \left\{ \frac{1}{2\psi_{K,L}}, \frac{\rho_{K,L}}{l_{K,L}} \right\}$, we invoke the local smoothness property to obtain

$$\left(\theta_{K,L} - \frac{\theta_{K,L}^2 \psi_{K,L}}{2} \right) \|\nabla_L \mathcal{G}(K, L)\|_F^2 \leq \mathcal{G}(K, L + \theta_{K,L} \nabla_L \mathcal{G}(K, L)) - \mathcal{G}(K, L) \leq \mathcal{G}(K, L(K)) - \mathcal{G}(K, L).$$

Let $\Delta := \mathcal{G}(K, L(K)) - \mathcal{G}(K, L)$ and $\Delta' := \mathcal{G}(K, L(K)) - \mathcal{G}(K, L')$. Substituting (C.8) into (C.7) implies that the following inequality holds almost surely:

$$\begin{aligned} \Delta' - \Delta &\leq \left(-\frac{\eta_L}{2\kappa} + \frac{8\psi_{K,L}\eta_L^2}{\phi^2} \right) \|\nabla_L \mathcal{G}(K, L)\|_F^2 + \frac{2\eta_L \psi_{K,L} r_{1,L}}{\phi} \|\nabla_L \mathcal{G}(K, L)\|_F \\ &\quad + \left(\frac{2\eta_L \kappa}{\phi^2} + \frac{4\psi_{K,L}\eta_L^2}{\phi^2} \right) \|\bar{\nabla}_L \mathcal{G}(K, L) - \nabla_L \mathcal{G}_{r_1,L}(K, L)\|_F^2 + \frac{8\psi_{K,L}^3 \eta_L^2 r_{1,L}^2}{\phi^2} \\ &\leq -\frac{\eta_L \mu_K}{4\kappa} \Delta + \frac{4\eta_L \psi_{K,L} r_{1,L}}{\phi \theta_{K,L}} \Delta^{1/2} + \left(\frac{2\eta_L \kappa}{\phi^2} + \frac{4\psi_{K,L}\eta_L^2}{\phi^2} \right) \|\bar{\nabla}_L \mathcal{G}(K, L) - \nabla_L \mathcal{G}_{r_1,L}(K, L)\|_F^2 + \frac{8\psi_{K,L}^3 \eta_L^2 r_{1,L}^2}{\phi^2} \\ &\leq -\frac{\eta_L \mu_K}{8\kappa} \Delta + \frac{32\eta_L \kappa \psi_{K,L}^2 r_{1,L}^2}{\mu_K \phi^2 \theta_{K,L}^2} + \left(\frac{2\eta_L \kappa}{\phi^2} + \frac{4\psi_{K,L}\eta_L^2}{\phi^2} \right) \|\bar{\nabla}_L \mathcal{G}(K, L) - \nabla_L \mathcal{G}_{r_1,L}(K, L)\|_F^2 + \frac{8\psi_{K,L}^3 \eta_L^2 r_{1,L}^2}{\phi^2} \end{aligned}$$

where the second to last inequality uses $\eta_L \leq \phi^2 / [32\psi_{K,L}\kappa]$ and the PL condition. The last inequality uses $2ab \leq a^2 + b^2$. Thus, if the stepsize η_L of the one-step ZO-NPG update (4.2) and the smoothing radius $r_{1,L}$ for the minibatch estimator satisfy

$$\eta_L \leq \min \left\{ \frac{\phi^2}{32\psi_{K,L}\kappa}, \frac{1}{2\psi_{K,L}} \right\}, \quad r_{1,L} \leq \frac{\phi}{16\kappa\psi_{K,L}} \min \left\{ \sqrt{\mu_K \epsilon_1}, \mu_K \theta_{K,L} \sqrt{\frac{\epsilon_1}{8}} \right\},$$

then with probability at least $1 - \delta_1$, we can bound the one-step ascent as

$$\begin{aligned} \Delta' - \Delta &\leq -\frac{\eta_L \mu_K}{8\kappa} \Delta + \frac{2\eta_L(\kappa + 1)}{\phi^2} \|\bar{\nabla}_L \mathcal{G}(K, L) - \nabla_L \mathcal{G}_{r_1,L}(K, L)\|_F^2 + \frac{32\eta_L \kappa \psi_{K,L}^2 r_{1,L}^2}{\mu_K \phi^2 \theta_{K,L}^2} + \frac{4\eta_L \kappa \psi_{K,L}^2 r_{1,L}^2}{\phi^2} \\ &\leq -\frac{\eta_L \mu_K}{8\kappa} \Delta + \frac{\eta_L \mu_K \epsilon_1}{16\kappa} \implies \Delta' \leq \left(1 - \frac{\eta_L \mu_K}{8\kappa} \right) \Delta + \frac{\eta_L \mu_K \epsilon_1}{16\kappa}, \end{aligned} \quad (\text{C.9})$$

where the last inequality uses (C.5). By (C.9), we have $\mathcal{G}(K, L) - \mathcal{G}(K, L') = \Delta' - \Delta \leq -\frac{\eta_L \mu_K}{8\kappa} \Delta + \frac{\eta_L \mu_K \epsilon_1}{16\kappa}$. Therefore, it holds that $\mathcal{G}(K, L) - \mathcal{G}(K, L') \leq 0$ with probability at least $1 - \delta_1$ because $\epsilon_1 \leq \Delta$.

Lastly, we prove Lemma C.1. Due to the modification to sample state trajectories using unperturbed gain matrix, we avoid the bias induced by the perturbation on the gain matrix compared to the results presented in (Fazel et al., 2018). As a result, the only bias between the estimated correlation matrix and the exact one, denoted as $\|\bar{\Sigma}_{K,L} - \Sigma_{K,L}\|_F$, where $\bar{\Sigma}_{K,L(K)} = \frac{1}{M_1} \sum_{i=0}^{M_1-1} \text{diag}(x_0^i(x_0^i)^\top, \dots, x_N^i(x_N^i)^\top)$ and $\Sigma_{K,L} = \mathbb{E}_\xi[\bar{\Sigma}_{K,L}]$, is induced by ξ . Since

$$\|\text{diag}(x_0^i(x_0^i)^\top, \xi_0^i(\xi_0^i)^\top, \dots, \xi_{N-1}^i(\xi_{N-1}^i)^\top)\|_F \leq c_0$$

holds almost surely, and

$$\text{diag}(x_0^i(x_0^i)^\top, \dots, x_N^i(x_N^i)^\top) = \sum_{t=0}^{N-1} (\mathbf{A}_{K,L})^t \cdot \text{diag}(x_0^i(x_0^i)^\top, \xi_0^i(\xi_0^i)^\top, \dots, \xi_{N-1}^i(\xi_{N-1}^i)^\top) \cdot (\mathbf{A}_{K,L}^\top)^t,$$

there exists a $c_{\Sigma_{K,L}}$ that is a polynomial of $\|\mathbf{A}\|_F$, $\|\mathbf{B}\|_F$, $\|\mathbf{D}\|_F$, $\|\mathbf{K}\|_F$, $\|\mathbf{L}\|_F$, and is linear in c_0 such that $\|\text{diag}(x_0^i(x_0^i)^\top, \dots, x_N^i(x_N^i)^\top)\|_F \leq c_{\Sigma_{K,L}}$ almost surely. Then, we apply Hoeffding's inequality to get with probability at least $1 - 2\delta_1$ that $\|\bar{\Sigma}_{K,L} - \Sigma_{K,L}\|_F \leq \sqrt{c_{\Sigma_{K,L}}^2 \log(d_\Sigma/\delta_1)/[2M_1]}$. Therefore, it suffices to choose $M_1 \geq \frac{1}{2} c_{\Sigma_{K,L}}^2 \epsilon_1^{-2} \log(2d_\Sigma/\delta_1)$ to ensure that $\|\bar{\Sigma}_{K,L} - \Sigma_{K,L}\|_F \leq \epsilon_1$ with probability at least $1 - \delta_1$. Lastly, by Weyl's theorem, we can bound $\lambda_{\min}(\bar{\Sigma}_{K,L}) \geq \phi/2$ by requiring $\epsilon_1 \leq \phi/2$. This completes the proof. \blacksquare

C.4. Proof of Lemma B.7

Proof We note that (B.23) follows from Proposition B.1 of (Zhang et al., 2019), but with the inner- and outer-loop players being interchanged. That is, our B, D, K, L, R^u, R^w correspond to C, B, L, K, R^v, R^u in (Zhang et al., 2019). Following (B.31) and (B.32) of (Zhang et al., 2019), we can prove (B.23). We note that the constants in Proposition B.1 of (Zhang et al., 2019) are uniform over the compact set Ω therein. However, it is hard to construct such a compact set in our setting without posing additional constraints. Therefore, our $\mathcal{B}_{1,K}$ and $\mathcal{B}_{P,K}$ are only continuous functions of K rather than being absolute constants, due to the quotient of two continuous functions being a continuous function provided that the denominator is not zero. Equations (B.24) and (B.25) follow from Lemmas C.2 and B.8, respectively, in the supplementary material of (Zhang et al., 2019), with the same notational correspondences introduced above. Similarly, we do not have uniform constants here but instead $\mathcal{B}_{L(K),K}, \mathcal{B}_{\Sigma,K}$ are continuous functions of K . This completes the proof. ■

C.5. Proof of Lemma B.8

Proof By the definition that $\nabla_K \mathcal{G}(K, L(K)) = 2F_{K,L(K)} \Sigma_{K,L(K)}$, we have

$$\begin{aligned} \|\nabla_K \mathcal{G}(K', L(K')) - \nabla_K \mathcal{G}(K, L(K))\|_F &\leq 2\|F_{K',L(K')} - F_{K,L(K)}\|_F \|\Sigma_{K',L(K')}\|_F \\ &\quad + 2\|F_{K,L(K)}\|_F \|\Sigma_{K',L(K')} - \Sigma_{K,L(K)}\|_F. \end{aligned}$$

From Lemma C.4 and let $K, K' \in \mathcal{K}$ be sufficiently close to each other such that $\|K' - K\|_F \leq \mathcal{B}_{1,K}$, we have $\|F_{K',L(K')} - F_{K,L(K)}\|_F \leq c_{5,K} \cdot \|K' - K\|_F$, where $c_{5,K}$ is defined in §C.9. Next, the term $\|\Sigma_{K',L(K')}\|_F$ can be separated into $\|\Sigma_{K',L(K')}\|_F \leq \|\Sigma_{K,L(K)}\|_F + \|\Sigma_{K',L(K')} - \Sigma_{K,L(K)}\|_F$. To bound $\|\Sigma_{K,L(K)}\|_F$, we note that because $\|x_0\|, \|\xi_t\| \leq \vartheta$ almost surely for all t , there exists a finite $c_{\Sigma_{K,L(K)}}$ that is a polynomial of $\|A\|_F, \|B\|_F, \|D\|_F, \|K\|_F$, and c_0 such that $\|\Sigma_{K,L(K)}\|_F \leq c_{\Sigma_{K,L(K)}}$ almost surely. By Lemma B.7 and the condition that $\|K' - K\|_F \leq \mathcal{B}_{1,K}$, we have $\|\Sigma_{K',L(K')}\|_F \leq c_{\Sigma_{K,L(K)}} + \mathcal{B}_{\Sigma,K}$, which leads to

$$2\|F_{K',L(K')} - F_{K,L(K)}\|_F \|\Sigma_{K',L(K')}\|_F \leq 2c_{5,K}(c_{\Sigma_{K,L(K)}} + \mathcal{B}_{\Sigma,K}) \cdot \|K' - K\|_F.$$

Therefore, if we require $\|K' - K\|_F \leq \epsilon_2 / [4c_{5,K}(c_{\Sigma_{K,L(K)}} + \mathcal{B}_{\Sigma,K})]$, then $2\|F_{K',L(K')} - F_{K,L(K)}\|_F \|\Sigma_{K',L(K')}\|_F \leq \frac{\epsilon_2}{2}$ holds almost surely. Subsequently, we combine (C.15) with (B.25) to get

$$2\|F_{K,L(K)}\|_F \|\Sigma_{K',L(K')} - \Sigma_{K,L(K)}\|_F \leq 2c_{2,K} \mathcal{B}_{\Sigma,K} \cdot \|K' - K\|_F.$$

By requiring $\|K' - K\|_F \leq \frac{\epsilon_2}{4c_{2,K} \mathcal{B}_{\Sigma,K}}$, it holds almost surely that $2\|F_{K,L(K)}\|_F \|\Sigma_{K',L(K')} - \Sigma_{K,L(K)}\|_F \leq \frac{\epsilon_2}{2}$. Therefore, we can almost surely bound $\|\nabla_K \mathcal{G}(K', L(K')) - \nabla_K \mathcal{G}(K, L(K))\|_F \leq \epsilon_2$ by enforcing the above requirements on $\|K' - K\|_F$. This completes the proof. ■

C.6. Proof of Lemma B.9

Proof For any $\underline{K} \in \mathcal{K}$, we can define the following set:

$$\underline{\mathcal{K}} := \left\{ K \mid (3.6) \text{ admits a solution } P_{K,L(K)} \geq 0, \text{ and } P_{K,L(K)} \leq P_{\underline{K},L(\underline{K})} \right\}. \quad (\text{C.10})$$

Clearly, it holds that $\underline{K} \in \underline{\mathcal{K}}$. Then, we first prove that $\underline{\mathcal{K}}$ is bounded. Recall that $P_{K,L(K)} \geq 0$ is the solution to (3.6), such that

$$P_{K,L(K)} = Q + K^\top R^u K + (A - BK)^\top (P_{K,L(K)} + P_{K,L(K)} D (R^w - D^\top P_{K,L(K)} D)^{-1} D^\top P_{K,L(K)}) (A - BK). \quad (\text{C.11})$$

Since $\underline{K} \in \underline{\mathcal{K}}$, any $K \in \underline{\mathcal{K}}$ also satisfies $K \in \mathcal{K}$, due to that $P_{K,L(K)} \leq P_{\underline{K},L(\underline{K})}$ implies $R^w - D^\top P_{K,L(K)} D > 0$. As a result, the second term on the RHS of (C.11) is p.s.d. and we can obtain $Q + K^\top R^u K \leq P_{K,L(K)}$, where $Q \geq 0$ and $R^u > 0$. Then, from the definition of $\underline{\mathcal{K}}$ and $P_{\underline{K},L(\underline{K})}, P_{K,L(K)}$ are symmetric and p.s.d., we have $\|P_{K,L(K)}\|_F \leq \|P_{\underline{K},L(\underline{K})}\|_F$. These arguments together imply that for any $K \in \underline{\mathcal{K}}$, $\|K\|_F \leq \sqrt{\|P_{\underline{K},L(\underline{K})}\|_F / \lambda_{\min}(R^u)}$, proving the boundedness of $\underline{\mathcal{K}}$.

Next, we take an arbitrary sequence $\{\mathbf{K}_n\} \in \underline{\mathcal{K}}$ and note that $\|\mathbf{K}_n\|_F$ is bounded for all n . Applying Bolzano-Weierstrass theorem implies that the set of limit points of $\{\mathbf{K}_n\}$, denoted as $\underline{\mathcal{K}}_{\text{lim}}$, is nonempty. Then, for any $\mathbf{K}_{\text{lim}} \in \underline{\mathcal{K}}_{\text{lim}}$, we can find a subsequence $\{\mathbf{K}_{\tau_n}\} \in \underline{\mathcal{K}}$ that converges to \mathbf{K}_{lim} . We denote the corresponding sequence of solutions to (C.11) as $\{\mathbf{P}_{\mathbf{K}_{\tau_n}, \mathbf{L}(\mathbf{K}_{\tau_n})}\}$, where $0 \leq \mathbf{P}_{\mathbf{K}_{\tau_n}, \mathbf{L}(\mathbf{K}_{\tau_n})} \leq \mathbf{P}_{\underline{\mathbf{K}}, \mathbf{L}(\underline{\mathbf{K}})}$ for all n . By Bolzano-Weierstrass theorem, the boundedness of $\{\mathbf{P}_{\mathbf{K}_{\tau_n}, \mathbf{L}(\mathbf{K}_{\tau_n})}\}$, and the continuity of (C.11) with respect to \mathbf{K} , we have the set of limit points of $\{\mathbf{P}_{\mathbf{K}_{\tau_n}, \mathbf{L}(\mathbf{K}_{\tau_n})}\}$, denoted as $\underline{\mathcal{P}}_{\text{lim}}$, is nonempty. Then, for any $\mathbf{P}_{\text{lim}} \in \underline{\mathcal{P}}_{\text{lim}}$, we can again find a subsequence $\{\mathbf{P}_{\mathbf{K}_{\kappa_{\tau_n}}, \mathbf{L}(\mathbf{K}_{\kappa_{\tau_n}})}\}$ that converges to \mathbf{P}_{lim} . Since $\mathbf{P}_{\mathbf{K}_{\kappa_{\tau_n}}, \mathbf{L}(\mathbf{K}_{\kappa_{\tau_n}})}$ is a p.s.d. solution to (C.11) satisfying $0 \leq \mathbf{P}_{\mathbf{K}_{\kappa_{\tau_n}}, \mathbf{L}(\mathbf{K}_{\kappa_{\tau_n}})} \leq \mathbf{P}_{\underline{\mathbf{K}}, \mathbf{L}(\underline{\mathbf{K}})}$ for all n and (C.11) is continuous in \mathbf{K} , \mathbf{P}_{lim} must solve (C.11) and satisfy $0 \leq \mathbf{P}_{\text{lim}} \leq \mathbf{P}_{\underline{\mathbf{K}}, \mathbf{L}(\underline{\mathbf{K}})}$, which implies $\mathbf{K}_{\text{lim}} \in \underline{\mathcal{K}}$. Note that the above arguments work for any sequence $\{\mathbf{K}_n\} \in \underline{\mathcal{K}}$ and any limit points \mathbf{K}_{lim} and \mathbf{P}_{lim} , which proves the closedness of $\underline{\mathcal{K}}$. Together with the boundedness, $\underline{\mathcal{K}}$ is thus compact.

Finally, let us denote the closure of the complement of \mathcal{K} as $\overline{\mathcal{K}}^c$. By $\underline{\mathbf{K}} \in \underline{\mathcal{K}} \subset \mathcal{K}$ and (C.10), any $\mathbf{K} \in \underline{\mathcal{K}}$ satisfies: i) $\mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})} \geq 0$ exists; ii) $\mathbf{R}^w - \mathbf{D}^\top \mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})} \mathbf{D} \geq \mathbf{R}^w - \mathbf{D}^\top \mathbf{P}_{\underline{\mathbf{K}}, \mathbf{L}(\underline{\mathbf{K}})} \mathbf{D} > 0$. This implies that $\underline{\mathcal{K}}$ is disjoint with $\overline{\mathcal{K}}^c$, i.e., $\underline{\mathcal{K}} \cap \overline{\mathcal{K}}^c = \emptyset$. Then, there exists a distance $\mathcal{B}_{2, \mathbf{K}} > 0$ between $\underline{\mathcal{K}}$ and $\overline{\mathcal{K}}^c$ such that for a given $\underline{\mathbf{K}} \in \underline{\mathcal{K}}$, all \mathbf{K}' satisfying $\|\mathbf{K}' - \underline{\mathbf{K}}\| \leq \mathcal{B}_{2, \mathbf{K}}$ also satisfy $\mathbf{K}' \in \mathcal{K}$ (see for example Lemma A.1 of (Berg, 2001)). This completes the proof. ■

C.7. Proof of Lemma B.10

Proof To simplify the notations, we first define

$$\begin{aligned} \check{\nabla} &:= \overline{\nabla}_{\mathbf{K}} \mathcal{G}(\mathbf{K}, \overline{\mathbf{L}}(\mathbf{K})) = \frac{1}{M_2} \sum_{j=0}^{M_2-1} \frac{d_2}{r_2} \left[\sum_{t=0}^N c_t^j \right] \mathbf{V}^j, \quad \overline{\nabla} := \overline{\nabla}_{\mathbf{K}} \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})), \quad \widehat{\nabla} := \mathbb{E}_{\xi}[\overline{\nabla}], \\ \widetilde{\nabla} &:= \nabla_{\mathbf{K}} \mathcal{G}_{r_2}(\mathbf{K}, \mathbf{L}(\mathbf{K})), \quad \nabla := \nabla_{\mathbf{K}} \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})), \end{aligned}$$

where $\overline{\mathbf{L}}(\mathbf{K})$ is the approximate inner-loop solution obtained from Algorithm 1. Our goal is to quantify $\|\check{\nabla} - \nabla\|_F$, which can be separated into four terms $\|\check{\nabla} - \nabla\|_F \leq \|\check{\nabla} - \overline{\nabla}\|_F + \|\overline{\nabla} - \widehat{\nabla}\|_F + \|\widehat{\nabla} - \widetilde{\nabla}\|_F + \|\widetilde{\nabla} - \nabla\|_F$.

According to Theorem 4.1, $\mathcal{G}(\mathbf{K} + r_2 \mathbf{V}, \overline{\mathbf{L}}(\mathbf{K} + r_2 \mathbf{V})) \geq \mathcal{G}(\mathbf{K} + r_2 \mathbf{V}, \mathbf{L}(\mathbf{K} + r_2 \mathbf{V})) - \epsilon_1$ holds with probability at least $1 - \delta_1$. Applying union bound and triangle inequality yield that $\|\check{\nabla} - \overline{\nabla}\|_F \leq d_2 \epsilon_1 / r_2$ holds with probability at least $1 - M_2 \delta_1$. Therefore, we can bound $\|\check{\nabla} - \overline{\nabla}\|_F$ by $\epsilon_2 / 4$ with probability at least $1 - \delta_2 / 3$ by requiring $\epsilon_1 \leq \epsilon_2 r_2 / [4d_2]$, and $\delta_1 \leq \delta_2 / [3M_2]$.

Subsequently, we consider the bias induced by ξ . By definition, we have $\sum_{t=0}^N c_t^j$ driven by $(\mathbf{K}, \mathbf{L}(\mathbf{K}))$ and $x_0, \xi_t \sim \mathcal{D}$ for all t can also be represented as $\text{Tr}[\mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})} \cdot \text{diag}(x_0 x_0^\top, \xi_0 \xi_0^\top, \dots, \xi_{N-1} \xi_{N-1}^\top)]$, where $\|\text{diag}(x_0 x_0^\top, \xi_0 \xi_0^\top, \dots, \xi_{N-1} \xi_{N-1}^\top)\|_F \leq c_0$ almost surely. Then, we apply Hoeffding's inequality and (C.13) to get with probability at least $1 - 2\delta_2$ that $\|\overline{\nabla} - \widehat{\nabla}\|_F \leq d_2 c_0 \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) / [r_2 \phi] \cdot \sqrt{\log(d_2 / \delta_2) / [2M_2]}$. Then, it suffices to choose $M_2 \geq 8d_2^2 c_0^2 \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K}))^2 / [r_2^2 \phi^2 \epsilon_2^2] \cdot \log(6d_2 / \delta_2)$ to ensure that $\|\overline{\nabla} - \widehat{\nabla}\|_F \leq \epsilon_2 / 4$ with probability at least $1 - \delta_2 / 3$.

For the third term, we can apply standard results in zeroth-order optimization (Flaxman et al., 2005), which leads to $\widetilde{\nabla} = d_2 / r_2 \cdot \mathbb{E}_{\mathbf{V}}[\mathcal{G}(\mathbf{K} + r_2 \mathbf{V}, \mathbf{L}(\mathbf{K} + r_2 \mathbf{V})) \mathbf{V}]$. Then, we invoke Lemma B.7 and Hoeffding's inequality to get with probability at least $1 - 2\delta_2$ that $\|\widehat{\nabla} - \widetilde{\nabla}\|_F \leq d_2 / r_2 \cdot (\mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) + r_2 \mathcal{B}_{\mathbf{P}, \mathbf{K}} c_0) \cdot \sqrt{2 \log(d_2 / \delta_2) / M_2}$, where we require $r_2 \leq \mathcal{B}_{1, \mathbf{K}}$. Thus, it suffices to choose $M_2 \geq 32d_2^2 (\mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) + r_2 \mathcal{B}_{\mathbf{P}, \mathbf{K}} c_0)^2 / [r_2^2 \epsilon_2^2] \cdot \log(6d_2 / \delta_2)$ to ensure that $\|\widehat{\nabla} - \widetilde{\nabla}\|_F \leq \epsilon_2 / 4$ with probability at least $1 - \delta_2 / 3$.

Lastly, by Lemma B.8, we can bound the last term $\|\widetilde{\nabla} - \nabla\|_F$ almost surely by $\epsilon_2 / 4$ if requiring

$$r_2 = \|\mathbf{K}' - \mathbf{K}\|_F \leq \min \left\{ \mathcal{B}_{1, \mathbf{K}}, \mathcal{B}_{2, \mathbf{K}}, \frac{\epsilon_2}{16c_{5, \mathbf{K}}(c_{\Sigma_{\mathbf{K}, \mathbf{L}(\mathbf{K})}} + \mathcal{B}_{\Sigma, \mathbf{K}})}, \frac{\epsilon_2}{16c_{2, \mathbf{K}} \mathcal{B}_{\Sigma, \mathbf{K}}} \right\}$$

where $\mathcal{B}_{1, \mathbf{K}}, \mathcal{B}_{\Sigma, \mathbf{K}}$ are defined in Lemma B.7, $c_{\Sigma_{\mathbf{K}, \mathbf{L}(\mathbf{K})}}$ is defined in §C.5, and $r_2 \leq \mathcal{B}_{2, \mathbf{K}}$ ensures the perturbed gain matrices $\mathbf{K} + r_2 \mathbf{V}^j$, for all j , lie within \mathcal{K} , as proved in Lemma B.9. Thus, we have $\|\check{\nabla} - \nabla\|_F \leq \epsilon_2$ with probability at least $1 - \delta_2$, which completes the proof. ■

C.8. Proof of Lemma B.11

Proof We first introduce a lemma similar to Lemma 16 of (Fazel et al., 2018) and defer its proof to the end of this subsection.

Lemma C.2 ($\Sigma_{K,L}$ Perturbation) *For a fixed $K \in \mathcal{K}$, there exist some continuous functions $\mathcal{B}_{1,L}, \mathcal{B}_{\Sigma,L} > 0$ such that if $\|L' - L\|_F \leq \mathcal{B}_{1,L}$, then it holds that $\|\Sigma_{K,L'} - \Sigma_{K,L}\|_F \leq \mathcal{B}_{\Sigma,L} \cdot \|L' - L\|_F$.*

The estimation bias, $\|\bar{\Sigma}_{K,\bar{L}(K)} - \Sigma_{K,L(K)}\|_F$, can be separated into two terms, such that

$$\|\bar{\Sigma}_{K,\bar{L}(K)} - \Sigma_{K,L(K)}\|_F \leq \|\bar{\Sigma}_{K,\bar{L}(K)} - \Sigma_{K,\bar{L}(K)}\|_F + \|\Sigma_{K,\bar{L}(K)} - \Sigma_{K,L(K)}\|_F, \quad (\text{C.12})$$

where $\bar{\Sigma}_{K,\bar{L}(K)} = \frac{1}{M_2} \sum_{j=0}^{M_2-1} \text{diag}(x_0^j(x_0^j)^\top, \dots, x_N^j(x_N^j)^\top)$, $\Sigma_{K,\bar{L}(K)} = \mathbb{E}_\xi [\bar{\Sigma}_{K,\bar{L}(K)}]$, and $\{x_t\}$ is the sequence of noisy states driven by the independently sampled noises $x_0, \xi_t \sim \mathcal{D}$, for all t , and the pair of control gain matrix $(K, \bar{L}(K))$, where $\bar{L}(K)$ is the approximate inner-loop solution. Recall that $\|\text{diag}(x_0^j(x_0^j)^\top, \xi_0^j(\xi_0^j)^\top, \dots, \xi_{N-1}^j(\xi_{N-1}^j)^\top)\|_F \leq c_0$ almost surely and

$$\text{diag}(x_0^j(x_0^j)^\top, \dots, x_N^j(x_N^j)^\top) = \sum_{t=0}^{N-1} (A_{K,\bar{L}(K)})^t \cdot \text{diag}(x_0^j(x_0^j)^\top, \xi_0^j(\xi_0^j)^\top, \dots, \xi_{N-1}^j(\xi_{N-1}^j)^\top) \cdot (A_{K,\bar{L}(K)}^\top)^t,$$

then there exists a $c_{\Sigma_{K,L(K)}}$ that is a polynomial of $\|A\|_F, \|B\|_F, \|D\|_F, \|K\|_F$, and is linear in c_0 such that $\|\text{diag}(x_0^j(x_0^j)^\top, \dots, x_N^j(x_N^j)^\top)\|_F \leq c_{\Sigma_{K,L(K)}}$ almost surely. Then, we apply Hoeffding's inequality to get with probability at least $1 - 2\delta_2$ that $\|\bar{\Sigma}_{K,\bar{L}(K)} - \Sigma_{K,\bar{L}(K)}\|_F \leq \sqrt{c_{\Sigma_{K,L(K)}}^2 \log(d_\Sigma/\delta_2)/[2M_2]}$. Therefore, it suffices to choose $M_2 \geq 2c_{\Sigma_{K,L(K)}}^2 \epsilon_2^{-2} \log(\frac{4d_\Sigma}{\delta_2})$ to ensure that $\|\bar{\Sigma}_{K,\bar{L}(K)} - \Sigma_{K,\bar{L}(K)}\|_F \leq \epsilon_2/2$ with probability at least $1 - \delta_2/2$.

For the second term on the RHS of (C.12), we have by Theorems 4.1 or B.4, depending on whether the inner-loop oracle is implemented with the ZO-PG (4.1) or the ZO-NPG (4.2) updates, that with probability at least $1 - \delta_1$: $\|\bar{L}(K) - L(K)\|_F^2 \leq \lambda_{\min}^{-1}(H_{K,L(K)}) \cdot \epsilon_1$. Then, we apply Lemma C.2 to get when $\|\bar{L}(K) - L(K)\|_F \leq \mathcal{B}_{1,L(K)}$, it holds that $\|\Sigma_{K,\bar{L}(K)} - \Sigma_{K,L(K)}\|_F \leq \mathcal{B}_{\Sigma,L(K)} \|\bar{L}(K) - L(K)\|_F$ with probability at least $1 - \delta_1$. Thus, if

$$\epsilon_1 \leq \min \left\{ \mathcal{B}_{1,L(K)}^2 \lambda_{\min}(H_{K,L(K)}), \frac{\epsilon_2^2 \lambda_{\min}(H_{K,L(K)})}{4\mathcal{B}_{\Sigma,L(K)}^2} \right\}, \quad \delta_1 \leq \frac{\delta_2}{2},$$

then we can bound $\|\Sigma_{K,\bar{L}(K)} - \Sigma_{K,L(K)}\|_F \leq \epsilon_2/2$ with probability at least $1 - \delta_2/2$. Combining two terms together, we can conclude that $\|\bar{\Sigma}_{K,\bar{L}(K)} - \Sigma_{K,L(K)}\|_F \leq \epsilon_2$ with probability at least $1 - \delta_2$. Then, by the Weyl's Theorem, we can bound $\lambda_{\min}(\bar{\Sigma}_{K,\bar{L}(K)}) \geq \phi/2$ if $\epsilon_2 \leq \phi/2$. Lastly, we prove Lemma C.2, which mostly follows from Lemma 16 of (Fazel et al., 2018), with our $A - BK, D, Q + K^\top R^u K, L$ matrices being replaced by the A, B, Q, K therein, respectively, except that the upper bound for $\|\Sigma_K\|$ therein does not hold in our setting since we only require $Q \geq 0$ and thus $Q + K^\top R^u K$ may not be full-rank. Instead, we utilize that the value of the objective function following (4.1) or (4.2) is monotonically non-decreasing, as shown in Theorems 4.1 and B.4. Moreover, the superlevel set $\mathcal{L}_K(\mathcal{G}(K, L_0))$ following the definition in (3.8) is compact. Thus, there exists a constant $c_{\Sigma_{\mathcal{G}(K, L_0)}} := \max_{L \in \mathcal{L}_K(\mathcal{G}(K, L_0))} \|\Sigma_{K,L}\|$ depending on K such that for a fixed $K \in \mathcal{K}$, $\|\Sigma_{K,L}\|_F \leq c_{\Sigma_{\mathcal{G}(K, L_0)}}$ holds for all iterates of L following (4.1) or (4.2) until convergence of the inner loop. This implies that for a fixed $K \in \mathcal{K}$, if $\|L' - L\|_F \leq \mathcal{B}_{1,L} := \phi/[4c_{\Sigma_{\mathcal{G}(K, L_0)}}\|D\|_F(\|A - BK - DL\|_F + 1)]$, then it holds that $\|\Sigma_{K,L'} - \Sigma_{K,L}\|_F \leq \mathcal{B}_{\Sigma,L} \cdot \|L' - L\|_F$, where $\mathcal{B}_{\Sigma,L} := 4c_{\Sigma_{\mathcal{G}(K, L_0)}}^2\|D\|_F(\|A - BK - DL\|_F + 1)/\phi$. This completes the proof. \blacksquare

C.9. Auxiliary Bounds

Define the following polynomials of $\mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K}))$:

$$\begin{aligned} c_{1,\mathbf{K}} &:= \frac{\mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K}))}{\phi} + \frac{\|\mathbf{D}\|_F^2 \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K}))^2}{\phi^2 \cdot \lambda_{\min}(\mathbf{H}_{\mathbf{K}, \mathbf{L}(\mathbf{K}))}}, & c_{2,\mathbf{K}} &:= \sqrt{\frac{\|\mathbf{G}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|}{\phi} \cdot (\mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) - \mathcal{G}(\mathbf{K}^*, \mathbf{L}^*))}, \\ c_{3,\mathbf{K}} &:= \frac{c_{2,\mathbf{K}} + \|\mathbf{B}\|_F \|\mathbf{A}\|_F c_{1,\mathbf{K}}}{\lambda_{\min}(\mathbf{R}^u)}, & c_{4,\mathbf{K}} &:= \frac{\|\mathbf{D}\|_F \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K}))}{\phi \cdot \lambda_{\min}(\mathbf{H}_{\mathbf{K}, \mathbf{L}(\mathbf{K}))}} (\|\mathbf{A}\|_F + \|\mathbf{B}\|_F c_{3,\mathbf{K}}), \\ c_{5,\mathbf{K}} &:= 2 \left[\|\mathbf{R}^u\|_F + \|\mathbf{B}\|_F^2 \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) + \|\mathbf{B}\|_F \|\mathbf{D}\|_F \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) \mathcal{B}_{\mathbf{L}(\mathbf{K}), \mathbf{K}} \right. \\ &\quad \left. + \mathcal{B}_{\mathbf{P}, \mathbf{K}} \|\mathbf{B}\|_F \left(\|\mathbf{B}\|_F (c_{3,\mathbf{K}} + \mathcal{B}_{1,\mathbf{K}}) + \|\mathbf{A}\|_F + \|\mathbf{D}\|_F (c_{4,\mathbf{K}} + \mathcal{B}_{\mathbf{L}(\mathbf{K}), \mathbf{K}}) \right) \right], \end{aligned}$$

where $\mathcal{B}_{1,\mathbf{K}}$, $\mathcal{B}_{\mathbf{P}, \mathbf{K}}$, and $\mathcal{B}_{\mathbf{L}(\mathbf{K}), \mathbf{K}}$ follow from Lemma B.7. Then, we present the following lemmas.

Lemma C.3 For $\mathbf{K}, \mathbf{K}' \in \mathcal{K}$, the following inequalities hold

$$\|\mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F \leq \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K}))/\phi, \quad (\text{C.13}) \quad \|\mathbf{F}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F \leq c_{2,\mathbf{K}} \quad (\text{C.15}) \quad \|\mathbf{K}\|_F \leq c_{3,\mathbf{K}}, \quad (\text{C.17})$$

$$\|\tilde{\mathbf{P}}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F \leq c_{1,\mathbf{K}}, \quad (\text{C.14}) \quad \|\nabla_{\mathbf{K}} \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K}))\|_F \leq 2\|\Sigma_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\| c_{2,\mathbf{K}}, \quad (\text{C.16}) \quad \|\mathbf{L}(\mathbf{K})\|_F \leq c_{4,\mathbf{K}}. \quad (\text{C.18})$$

Proof To start with, we note that for any $\mathbf{K} \in \mathcal{K}$, the solution to the Riccati equation (3.6) satisfies $\mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})} \geq 0$. Thus, (C.13) can be proved by $\mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) = \text{Tr}(\mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})} \Sigma_0) \geq \phi \|\mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F$, where the inequality follows from $\phi = \lambda_{\min}(\Sigma_0) > 0$ and $\text{Tr}(M) \geq \|M\|_F$ for any matrix $M \geq 0$. Then, for any $\mathbf{K} \in \mathcal{K}$, we have by (C.13) and the definition of $\tilde{\mathbf{P}}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}$ in (3.5) that

$$\|\tilde{\mathbf{P}}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F \leq \|\mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F + \|\mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})} \mathbf{D} \mathbf{H}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}^{-1} \mathbf{D}^\top \mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F \leq \frac{\mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K}))}{\phi} + \frac{\|\mathbf{D}\|_F^2 \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K}))^2}{\phi^2 \cdot \lambda_{\min}(\mathbf{H}_{\mathbf{K}, \mathbf{L}(\mathbf{K}))}} = c_{1,\mathbf{K}}.$$

This completes the proof of (C.14). Subsequently, for any $\mathbf{K}, \mathbf{K}' \in \mathcal{K}$, we invoke Lemma 11 of (Fazel et al., 2018) to get

$$\mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) - \mathcal{G}(\mathbf{K}^*, \mathbf{L}^*) \geq \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) - \mathcal{G}(\mathbf{K}', \mathbf{L}(\mathbf{K}')) \geq \frac{\phi}{\|\mathbf{G}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|} \text{Tr}(\mathbf{F}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}^\top \mathbf{F}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}).$$

As a result, we have $\|\mathbf{F}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F^2 = \text{Tr}(\mathbf{F}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}^\top \mathbf{F}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}) \leq \frac{\|\mathbf{G}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|}{\phi} (\mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) - \mathcal{G}(\mathbf{K}^*, \mathbf{L}^*))$, and taking a square root of both sides yields (C.15). Moreover, we can prove (C.16) such that

$$\|\nabla_{\mathbf{K}} \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K}))\|_F^2 \leq 4\|\Sigma_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|^2 \text{Tr}(\mathbf{F}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}^\top \mathbf{F}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}) \leq \frac{4\|\Sigma_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|^2 \|\mathbf{G}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|}{\phi} (\mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) - \mathcal{G}(\mathbf{K}^*, \mathbf{L}(\mathbf{K}^*))).$$

Next, we invoke Lemma 25 of (Fazel et al., 2018) and (C.14) to get

$$\|\mathbf{K}\|_F \leq \frac{c_{2,\mathbf{K}}}{\lambda_{\min}(\mathbf{R}^u)} + \frac{\|\mathbf{B}^\top \tilde{\mathbf{P}}_{\mathbf{K}, \mathbf{L}(\mathbf{K})} \mathbf{A}\|_F}{\lambda_{\min}(\mathbf{R}^u)} \leq \frac{c_{2,\mathbf{K}} + \|\mathbf{B}\|_F \|\mathbf{A}\|_F c_{1,\mathbf{K}}}{\lambda_{\min}(\mathbf{R}^u)} = c_{3,\mathbf{K}},$$

which proves (C.17). Lastly, for any $\mathbf{K} \in \mathcal{K}$ and recall the definition of $\mathbf{L}(\mathbf{K})$ in Lemma 3.3

$$\|\mathbf{L}(\mathbf{K})\|_F = \|\mathbf{H}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}^{-1} \mathbf{D}^\top \mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})} (\mathbf{A} - \mathbf{B}\mathbf{K})\|_F \leq \frac{\|\mathbf{D}\|_F \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K}))}{\phi \cdot \lambda_{\min}(\mathbf{H}_{\mathbf{K}, \mathbf{L}(\mathbf{K}))}} (\|\mathbf{A}\|_F + \|\mathbf{B}\|_F c_{3,\mathbf{K}}) = c_{4,\mathbf{K}},$$

which proves (C.18). ■

Lemma C.4 For $\mathbf{K}, \mathbf{K}' \in \mathcal{K}$ satisfying $\|\mathbf{K}' - \mathbf{K}\|_F \leq \mathcal{B}_{1,\mathbf{K}}$, where $\mathcal{B}_{1,\mathbf{K}}$ is as defined in Lemma B.7, it holds that

$$\|\mathbf{F}_{\mathbf{K}', \mathbf{L}(\mathbf{K}')} - \mathbf{F}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F \leq c_{5,\mathbf{K}} \cdot \|\mathbf{K}' - \mathbf{K}\|_F. \quad (\text{C.19})$$

Algorithm 1 Inner-Loop Zeroth-Order Maximization Oracle

- 1: Input: gain matrices $(\mathbf{K}, \mathbf{L}_0)$, iteration L , batchsize M_1 , problem horizon N , distribution \mathcal{D} , smoothing radius r_1 , dimension $d_1 = mnN$, stepsize η .
- 2: **for** $l = 0, \dots, L - 1$ **do**
- 3: **for** $i = 0, \dots, M_1 - 1$ **do**
- 4: Sample $\mathbf{L}_l^i = \mathbf{L}_l + r_1 \mathbf{U}_l^i$, where \mathbf{U}_l^i is uniformly drawn from $\mathcal{S}(n, m, N)$ with $\|\mathbf{U}_l^i\|_F = 1$.
- 5: Simulate $(\mathbf{K}, \mathbf{L}_l^i)$ and $(\mathbf{K}, \mathbf{L}_l)$ for horizon N starting from $x_{l,0}^{i,0}, x_{l,0}^{i,1} \sim \mathcal{D}$, and collect the empirical estimates $\bar{\mathcal{G}}(\mathbf{K}, \mathbf{L}_l^i) = \sum_{t=0}^N c_{l,t}^{i,0}, \bar{\Sigma}_{\mathbf{K}, \mathbf{L}_l^i}^i = \text{diag}[x_{l,0}^{i,1}(x_{l,0}^{i,1})^\top, \dots, x_{l,N}^{i,1}(x_{l,N}^{i,1})^\top]$, where $\{c_{l,t}^{i,0}\}$ is the sequence of stage costs following the trajectory generated by $(\mathbf{K}, \mathbf{L}_l^i)$ and $\{x_{l,t}^{i,1}\}$ is the sequence of states following the trajectory generated by $(\mathbf{K}, \mathbf{L}_l)$, for $t \in \{0, \dots, N\}$, under independently sampled noises $\xi_{l,t}^{i,0}, \xi_{l,t}^{i,1} \sim \mathcal{D}$ for all $t \in \{0, \dots, N - 1\}$.
- 6: **end for**
- 7: Obtain: $\bar{\nabla}_L \mathcal{G}(\mathbf{K}, \mathbf{L}_l) = \frac{1}{M_1} \sum_{i=0}^{M_1-1} \frac{d_1}{r_1} \bar{\mathcal{G}}(\mathbf{K}, \mathbf{L}_l^i) \mathbf{U}_l^i$, $\bar{\Sigma}_{\mathbf{K}, \mathbf{L}_l} = \frac{1}{M_1} \sum_{i=0}^{M_1-1} \bar{\Sigma}_{\mathbf{K}, \mathbf{L}_l^i}^i$.
- 8: Execute one-step update:

$$\text{PG: } \mathbf{L}_{l+1} = \mathbf{L}_l + \eta \bar{\nabla}_L \mathcal{G}(\mathbf{K}, \mathbf{L}_l), \quad \text{NPG: } \mathbf{L}_{l+1} = \mathbf{L}_l + \eta \bar{\nabla}_L \mathcal{G}(\mathbf{K}, \mathbf{L}_l) \bar{\Sigma}_{\mathbf{K}, \mathbf{L}_l}^{-1}.$$
- 9: **end for**
- 10: Return \mathbf{L}_L .

Proof For $\mathbf{K}, \mathbf{K}' \in \mathcal{K}$ and recalling the definition of $\mathbf{F}_{\mathbf{K}, \mathbf{L}}$ in (3.4), we have

$$\begin{aligned}
 & \|\mathbf{F}_{\mathbf{K}', \mathbf{L}(\mathbf{K}')} - \mathbf{F}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F \\
 &= 2\|(\mathbf{R}^u + \mathbf{B}^\top \mathbf{P}_{\mathbf{K}', \mathbf{L}(\mathbf{K}')} \mathbf{B}) \mathbf{K}' - \mathbf{B}^\top \mathbf{P}_{\mathbf{K}', \mathbf{L}(\mathbf{K}')} (\mathbf{A} - \mathbf{D} \mathbf{L}(\mathbf{K}')) - (\mathbf{R}^u + \mathbf{B}^\top \mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})} \mathbf{B}) \mathbf{K} - \mathbf{B}^\top \mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})} (\mathbf{A} - \mathbf{D} \mathbf{L}(\mathbf{K}))\|_F \\
 &\leq 2\|\mathbf{R}^u\|_F \|\mathbf{K}' - \mathbf{K}\|_F + 2\|\mathbf{B}\|_F^2 \|\mathbf{K}'\|_F \|\mathbf{P}_{\mathbf{K}', \mathbf{L}(\mathbf{K}')} - \mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F + 2\|\mathbf{B}\|_F^2 \|\mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F \|\mathbf{K}' - \mathbf{K}\|_F \\
 &\quad + 2\|\mathbf{B}\|_F \|\mathbf{A}\|_F \|\mathbf{P}_{\mathbf{K}', \mathbf{L}(\mathbf{K}')} - \mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F + 2\|\mathbf{B}\|_F \|\mathbf{D}\|_F \|\mathbf{L}(\mathbf{K}') - \mathbf{L}(\mathbf{K})\|_F \|\mathbf{P}_{\mathbf{K}', \mathbf{L}(\mathbf{K}')} - \mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F \\
 &\quad + 2\|\mathbf{B}\|_F \|\mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F \|\mathbf{D}\|_F \|\mathbf{L}(\mathbf{K}') - \mathbf{L}(\mathbf{K})\|_F \\
 &= 2(\|\mathbf{R}^u\|_F + \|\mathbf{B}\|_F^2 \|\mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F) \cdot \|\mathbf{K}' - \mathbf{K}\|_F + 2(\|\mathbf{B}\|_F \|\mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F \|\mathbf{D}\|_F) \cdot \|\mathbf{L}(\mathbf{K}') - \mathbf{L}(\mathbf{K})\|_F \\
 &\quad + 2\left(\|\mathbf{B}\|_F^2 (\|\mathbf{K}\|_F + \|\mathbf{K}' - \mathbf{K}\|_F) + \|\mathbf{B}\|_F \|\mathbf{A}\|_F + \|\mathbf{B}\|_F \|\mathbf{D}\|_F (\|\mathbf{L}(\mathbf{K})\|_F + \|\mathbf{L}(\mathbf{K}') - \mathbf{L}(\mathbf{K})\|_F)\right) \cdot \|\mathbf{P}_{\mathbf{K}', \mathbf{L}(\mathbf{K}')} - \mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F.
 \end{aligned}$$

Applying (C.13), (C.17), (C.18), (B.23), and (B.24) proves (C.19), such that

$$\begin{aligned}
 & \|\mathbf{F}_{\mathbf{K}', \mathbf{L}(\mathbf{K}')} - \mathbf{F}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F \\
 &\leq 2(\|\mathbf{R}^u\|_F + \|\mathbf{B}\|_F^2 \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K}))) \cdot \|\mathbf{K}' - \mathbf{K}\|_F + 2(\|\mathbf{B}\|_F \|\mathbf{D}\|_F \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K}))) \cdot \|\mathbf{L}(\mathbf{K}') - \mathbf{L}(\mathbf{K})\|_F \\
 &\quad + 2\left(\|\mathbf{B}\|_F^2 (c_{3, \mathbf{K}} + \mathcal{B}_{1, \mathbf{K}}) + \|\mathbf{B}\|_F \|\mathbf{A}\|_F + \|\mathbf{B}\|_F \|\mathbf{D}\|_F (c_{4, \mathbf{K}} + \mathcal{B}_{\mathbf{L}(\mathbf{K}), \mathbf{K}})\right) \cdot \|\mathbf{P}_{\mathbf{K}', \mathbf{L}(\mathbf{K}')} - \mathbf{P}_{\mathbf{K}, \mathbf{L}(\mathbf{K})}\|_F \\
 &\leq 2\left[\|\mathbf{R}^u\|_F + \|\mathbf{B}\|_F^2 \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) + \|\mathbf{B}\|_F \|\mathbf{D}\|_F \mathcal{G}(\mathbf{K}, \mathbf{L}(\mathbf{K})) \mathcal{B}_{\mathbf{L}(\mathbf{K}), \mathbf{K}}\right. \\
 &\quad \left.+ \mathcal{B}_{\mathbf{P}, \mathbf{K}} \|\mathbf{B}\|_F \left(\|\mathbf{B}\|_F (c_{3, \mathbf{K}} + \mathcal{B}_{1, \mathbf{K}}) + \|\mathbf{A}\|_F + \|\mathbf{D}\|_F (c_{4, \mathbf{K}} + \mathcal{B}_{\mathbf{L}(\mathbf{K}), \mathbf{K}})\right)\right] \cdot \|\mathbf{K}' - \mathbf{K}\|_F.
 \end{aligned}$$

This completes the proof. ■

D. Pseudocode and Simulations

In this section, we provide pseudocode and simulations to complement our theories. In particular, we present convergence (divergence) results in the following four scenarios where we apply PG updates in Algorithms 1 and 2 to solve a finite-horizon zero-sum LQ dynamic game with time-invariant system parameters: (i) double-loop update scheme with both the inner- and outer-loop solved exactly; (ii) double-loop scheme with the inner-loop solved in the derivative-free scheme by sampling system trajectories while the outer-loop solved exactly; (iii) double-loop scheme with the inner-loop solved exactly while the outer-loop solved in the derivative-free scheme by sampling system trajectories; (iv) descent-multi-step-ascent/AGDA/ τ -GDA with exact gradient accesses diverge. Subsequently, we extend our numerical experiments to a setting with *time-varying* system parameters in §D.5, and demonstrate the convergence results of double-loop PG updates.

Algorithm 2 Outer-Loop ZO-NPG with Maximization Oracle

- 1: Input: initial gain matrix $\mathbf{K}_0 \in \mathcal{K}$, number of iterates K , batchsize M_2 , problem horizon N , distribution \mathcal{D} , smoothing radius r_2 , dimension $d_2 = mdN$, stepsize α .
- 2: **for** $k = 0, \dots, K - 1$ **do**
- 3: Find $\bar{\mathbf{L}}(\mathbf{K}_k)$ such that $\mathcal{G}(\mathbf{K}_k, \bar{\mathbf{L}}(\mathbf{K}_k)) \geq \mathcal{G}(\mathbf{K}_k, \mathbf{L}(\mathbf{K}_k)) - \epsilon_1$.
- 4: **for** $j = 0, \dots, M_2 - 1$ **do**
- 5: Sample $\mathbf{K}_k^j = \mathbf{K}_k + r_2 \mathbf{V}_k^j$, where \mathbf{V}_k^j is uniformly drawn from $\mathcal{S}(d, m, N)$ with $\|\mathbf{V}_k^j\|_F = 1$.
- 6: Find $\bar{\mathbf{L}}(\mathbf{K}_k^j)$ such that $\mathcal{G}(\mathbf{K}_k^j, \bar{\mathbf{L}}(\mathbf{K}_k^j)) \geq \mathcal{G}(\mathbf{K}_k^j, \mathbf{L}(\mathbf{K}_k^j)) - \epsilon_1$ using Algorithm 1.
- 7: Simulate $(\mathbf{K}_k^j, \bar{\mathbf{L}}(\mathbf{K}_k^j))$ and $(\mathbf{K}_k, \bar{\mathbf{L}}(\mathbf{K}_k))$ for horizon N starting from $x_{k,0}^{j,0}, x_{k,0}^{j,1} \sim \mathcal{D}$, and collect $\bar{\mathcal{G}}(\mathbf{K}_k^j, \bar{\mathbf{L}}(\mathbf{K}_k^j)) = \sum_{t=0}^N c_{k,t}^{j,0}, \bar{\Sigma}_{\mathbf{K}_k, \bar{\mathbf{L}}(\mathbf{K}_k)}^j = \text{diag}[x_{k,0}^{j,1}(x_{k,0}^{j,1})^\top, \dots, x_{k,N}^{j,1}(x_{k,N}^{j,1})^\top]$, where $\{c_{k,t}^{j,0}\}$ is the sequence of stage costs following the trajectory generated by $(\mathbf{K}_k^j, \bar{\mathbf{L}}(\mathbf{K}_k^j))$ and $\{x_{k,t}^{j,1}\}$ is the sequence of states following the trajectory generated by $(\mathbf{K}_k, \bar{\mathbf{L}}(\mathbf{K}_k))$, for $t \in \{0, \dots, N\}$, under independently sampled noises $\xi_{k,t}^{j,0}, \xi_{k,t}^{j,1} \sim \mathcal{D}$ for all $t \in \{0, \dots, N - 1\}$.
- 8: **end for**
- 9: Obtain: $\bar{\nabla}_{\mathbf{K}} \mathcal{G}(\mathbf{K}_k, \bar{\mathbf{L}}(\mathbf{K}_k)) = \frac{1}{M_2} \sum_{j=0}^{M_2-1} \frac{d_2}{r_2} \bar{\mathcal{G}}(\mathbf{K}_k^j, \bar{\mathbf{L}}(\mathbf{K}_k^j)) \mathbf{V}_k^j$, $\bar{\Sigma}_{\mathbf{K}_k, \bar{\mathbf{L}}(\mathbf{K}_k)} = \frac{1}{M_2} \sum_{j=0}^{M_2-1} \bar{\Sigma}_{\mathbf{K}_k, \bar{\mathbf{L}}(\mathbf{K}_k)}^j$.
- 10: Execute one step NPG update: $\mathbf{K}_{k+1} = \mathbf{K}_k - \alpha \bar{\nabla}_{\mathbf{K}} \mathcal{G}(\mathbf{K}_k, \bar{\mathbf{L}}(\mathbf{K}_k)) \bar{\Sigma}_{\mathbf{K}_k, \bar{\mathbf{L}}(\mathbf{K}_k)}^{-1}$.
- 11: **end for**
- 12: Return \mathbf{K}_K .

Simulation Setup. All the experiments are executed on a desktop computer equipped with a 3.7 GHz Hexa-Core Intel Core i7-8700K processor with Matlab R2019b. The device also has two 8GB 3000MHz DDR4 memories and a NVIDIA GeForce GTX 1080 8GB GDDR5X graphic card. In all settings except the one used in §D.5, we set the horizon of the problem to $N = 5$ and test a linear time-invariant system with the set of system matrices being $A_t = A, B_t = B, D_t = D, Q_t = Q, R_t^u = R^u$, and $R_t^w = R^w$, where $R^w = 5 \cdot \mathbf{I}$ and

$$A = \begin{bmatrix} 1 & 0 & -5 \\ -1 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}, \quad B = \begin{bmatrix} 1 & -10 & 0 \\ 0 & 3 & 1 \\ -1 & 0 & 2 \end{bmatrix}, \quad D = \begin{bmatrix} 0.5 & 0 & 0 \\ 0 & 0.2 & 0 \\ 0 & 0 & 0.2 \end{bmatrix},$$

$$Q = \begin{bmatrix} 2 & -1 & 0 \\ -1 & 2 & -1 \\ 0 & -1 & 2 \end{bmatrix}, \quad R^u = \begin{bmatrix} 4 & -1 & 0 \\ -1 & 4 & -2 \\ 0 & -2 & 3 \end{bmatrix}. \quad (\text{D.1})$$

D.1. Exact Double-Loop Updates

We set $\Sigma_0 = \mathbf{I}$. Then, we present the convergence result following two different initializations, denoted as \mathbf{K}_0^1 and \mathbf{K}_0^2 , where \mathbf{K}_0^2 is closer to the boundary of \mathcal{K} . Specifically, $\mathbf{K}_0^i = [\text{diag}((K_0^i)^5) \quad \mathbf{0}_{15 \times 3}]$, for $i \in \{1, 2\}$, where $K_0^1 = \begin{bmatrix} -0.12 & -0.01 & 0.62 \\ -0.21 & 0.14 & 0.15 \\ -0.06 & 0.05 & 0.42 \end{bmatrix}$ and $K_0^2 = \begin{bmatrix} -0.14 & -0.04 & 0.62 \\ -0.21 & 0.14 & 0.15 \\ -0.06 & 0.05 & 0.42 \end{bmatrix}$. Then, we can verify that $\lambda_{\min}(\mathbf{H}_{\mathbf{K}_0^1, \mathbf{L}(\mathbf{K}_0^1)}) = 0.5041$ and $\lambda_{\min}(\mathbf{H}_{\mathbf{K}_0^2, \mathbf{L}(\mathbf{K}_0^2)}) = 0.0199$. Moreover, we initialize $\mathbf{L} = \mathbf{0}$ in both cases.

Case 1: We demonstrate in Figure 4 the convergence of three update combinations, namely PG-NPG, NPG-NPG, and GN-GN, for the inner and outer loop, respectively. The stepsizes are chosen to be $(\eta, \alpha) = (1 \times 10^{-4}, 3 \times 10^{-6})$ for PG-NPG, $(\eta, \alpha) = (0.0635, 3 \times 10^{-6})$ for NPG-NPG, and $(\eta, \alpha) = (0.5, 5 \times 10^{-4})$ for GN-GN. Also, we require the approximate inner-loop solution to have the accuracy of $\epsilon_1 = 0.001$. As shown in the top of Figure 4, the double-loop algorithm with all three update combinations successfully converges to the unique Nash equilibrium of the game. Also, the sequence $\{\lambda_{\min}(\mathbf{H}_{\mathbf{K}_k, \mathbf{L}(\mathbf{K}_k)})\}_{k \geq 0}$ is monotonically non-decreasing, which matches the implicit regularization property we introduced in Theorem 3.7 for the outer-loop NPG and GN updates. As $\mathbf{K}_0^1 \in \mathcal{K}$, we can guarantee that all future iterates will stay inside \mathcal{K} . The convergences of the average gradient norms with respect to both \mathbf{L} and \mathbf{K} are also presented in the bottom of Figure 4.

Case 2: When the initial gain matrix \mathbf{K}_0^2 is closer to the boundary of \mathcal{K} , one non-judicious update could easily drive the gain

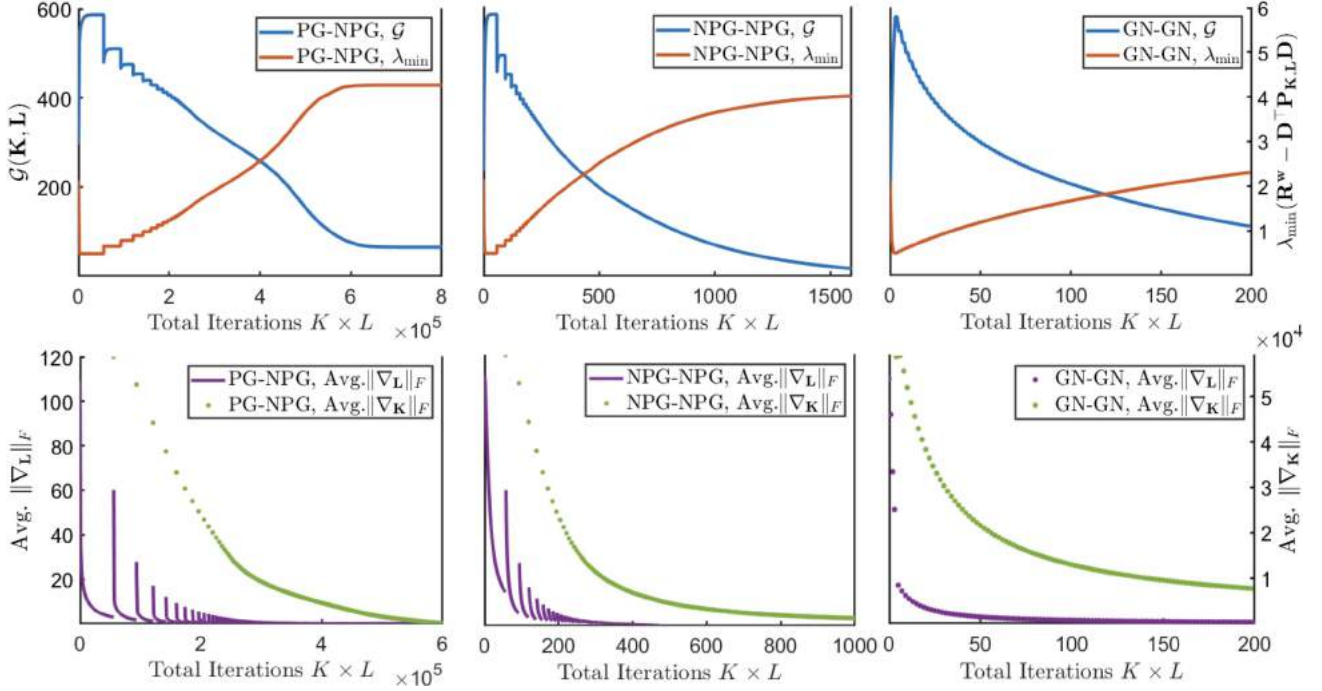


Figure 4. Three exact double-loop PG updates with initial gain matrix being \mathbf{K}_0^1 . *Top*: Convergence of $\mathcal{G}(\mathbf{K}, \mathbf{L})$ and $\lambda_{\min}(\mathbf{H}_{\mathbf{K}, \mathbf{L}})$. *Bottom*: Convergence of $\{l^{-1} \sum_{i=0}^{l-1} \|\nabla_{\mathbf{L}}(\mathbf{K}, \mathbf{L}_i)\|_F\}_{l \geq 1}$ and $\{k^{-1} \sum_{\kappa=0}^{k-1} \|\nabla_{\mathbf{K}} \mathcal{G}(\mathbf{K}_{\kappa}, \bar{\mathbf{L}}(\mathbf{K}_{\kappa}))\|_F\}_{k \geq 1}$. $\bar{\mathbf{L}}(\mathbf{K})$ is the approximate inner-loop solution with $\epsilon_1 = 10^{-3}$.

matrix outside of \mathcal{K} . In this case, we present two combinations of updates, namely NPG-NPG and GN-GN, for the inner and outer loop, respectively. We use the stepsizes $(\eta, \alpha) = (0.0635, 2.48 \times 10^{-7})$ for NPG-NPG and $(\eta, \alpha) = (0.5, 2.5 \times 10^{-4})$ for GN-GN. Similarly, we set $\epsilon_1 = 0.001$. The convergence patterns presented in Figure 5 are similar to those of **Case 1**.

D.2. Derivative-Free Inner-Loop Oracles

We provide two sets of experiments to validate Theorems 4.1 and B.4. We consider the same example as the one in §D.1, but using derivative-free updates (4.1)-(4.2) by sampling system trajectories to approximately solve the inner-loop subproblem, as proposed in §4.1. Note that the inner-loop subproblem is essentially an indefinite LQR, as introduced in §3. The outer-loop problem is solved using the exact NPG update (3.11). Note that according to Theorem 4.3, one can also use the outer-loop ZO-NPG (4.3) together with the inner-loop derivative-free updates to achieve the same convergence pattern. However, we use an exact outer-loop update, as it suffices to illustrate our idea here and also has better computational efficiency. In contrast to the example in §D.1, we set $\Sigma_0 = 0.1 \cdot \mathbf{I}$ and initialize $\mathbf{K}_0^3 = [\text{diag}((\mathbf{K}_0^3)^5) \quad \mathbf{0}_{15 \times 3}]$ with

$$\mathbf{K}_0^3 = \begin{bmatrix} -0.04 & -0.01 & 0.61 \\ -0.21 & 0.15 & 0.15 \\ -0.06 & 0.05 & 0.42 \end{bmatrix}, \text{ where one can verify that } \lambda_{\min}(\mathbf{H}_{\mathbf{K}_0^3, \mathbf{L}(\mathbf{K}_0^3)}) = 1.8673 > 0. \text{ That is, } \mathbf{K}_0^3 \in \mathcal{K}.$$

Inner-Loop ZO-PG: When we use the inner-loop ZO-PG update (4.1) in Algorithm 1, we choose $M_1 = 10^6$ and $r_1 = 1$. Moreover, we set the stepsizes of the inner-loop ZO-PG and the outer-loop exact NPG to be $(\eta, \alpha) = (8 \times 10^{-3}, 4.5756 \times 10^{-4})$, respectively. The desired accuracy level of the double-loop algorithm is picked to be $(\epsilon_1, \epsilon_2) = (0.8, 0.5)$. As shown in the left of Figure 6, the inner-loop ZO-PG (4.2) successfully converges for every fixed outer-loop update, validating our results in Theorem 4.1. Also, the double-loop algorithm converges to the unique Nash equilibrium of the game.

Inner-Loop ZO-NPG: When the inner-loop subproblem is solved via the ZO-NPG update (4.2), we choose the same $M_1, r_1, \alpha, \epsilon_1, \epsilon_2$ as the ones used in the inner-loop ZO-PG update but in contrast we set $\eta = 5 \times 10^{-2}$. Similar convergence

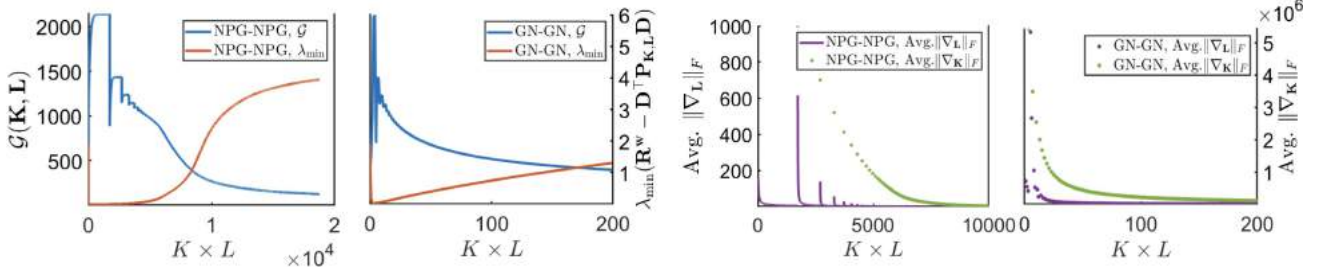


Figure 5. Two exact double-loop PG updates with initial gain matrix being \mathbf{K}_0^2 . *Left*: Convergence of $\mathcal{G}(\mathbf{K}, \mathbf{L})$ and $\lambda_{\min}(\mathbf{H}_{\mathbf{K}, \mathbf{L}(\mathbf{K})})$. *Right*: Convergence of $\{l^{-1} \sum_{\iota=0}^{l-1} \|\nabla_{\mathbf{L}}(\mathbf{K}, \mathbf{L}_{\iota})\|_F\}_{l \geq 1}$ and $\{k^{-1} \sum_{\kappa=0}^{k-1} \|\nabla_{\mathbf{K}} \mathcal{G}(\mathbf{K}_{\kappa}, \bar{\mathbf{L}}(\mathbf{K}_{\kappa}))\|_F\}_{k \geq 1}$. $\bar{\mathbf{L}}(\mathbf{K})$ is the approximate inner-loop solution with $\epsilon_1 = 10^{-3}$.

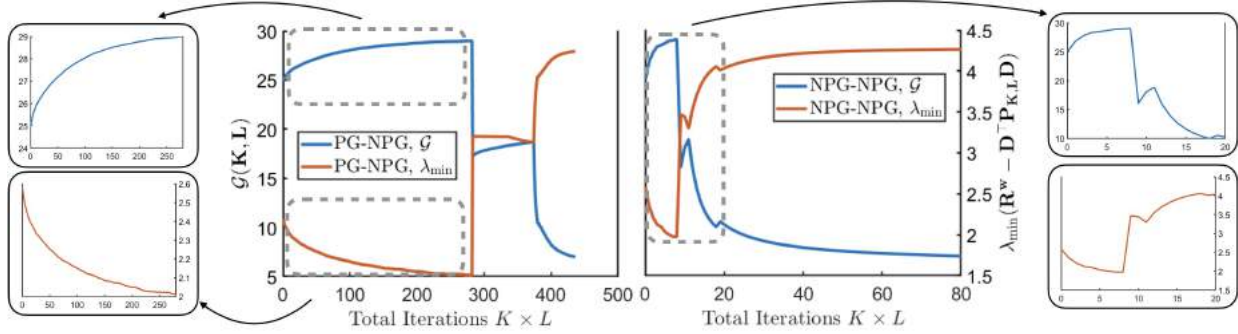


Figure 6. Convergence of $\mathcal{G}(\mathbf{K}, \mathbf{L})$ and $\lambda_{\min}(\mathbf{H}_{\mathbf{K}, \mathbf{L}(\mathbf{K})})$ with the initial gain matrix being \mathbf{K}_0^3 . *Left*: The inner loop follows the ZO-PG update (4.1) and the outer loop follows the exact NPG update (3.11). *Right*: The inner loop follows the ZO-NPG update (4.2) and the outer loop follows the exact NPG update (3.11).

patterns can be observed in the right of Figure 6.

D.3. Derivative-Free Outer-Loop NPG

We set $\Sigma_0 = 0.05 \cdot \mathbf{I}$. We present convergence results following the initial gain matrix $\mathbf{K}_0 = [\text{diag}(\mathbf{K}^5) \quad \mathbf{0}_{15 \times 3}]$, where

$\mathbf{K} = \begin{bmatrix} -0.08 & 0.35 & 0.62 \\ -0.21 & 0.19 & 0.32 \\ -0.06 & 0.10 & 0.41 \end{bmatrix}$ and $\lambda_{\min}(\mathbf{H}_{\mathbf{K}_0, \mathbf{L}(\mathbf{K}_0)}) = 3.2325$. The convergence of Algorithm 2 is implemented for two cases,

with the inner-loop oracle being implemented using (i) the exact solution as computed by $\mathbf{L}(\mathbf{K})$ in Lemma 3.3; and (ii) the approximate inner-loop solution following the exact NPG update (3.10). As proved by Theorems 4.1 and B.4, the inner loop can also be solved approximately using either ZO-PG or ZO-NPG updates, which has been verified via the simulation results presented in §D.2. The parameters of Algorithm 2 in both cases are set to $M_2 = 5 \times 10^5$, $r_2 = 0.08$, $(\eta, \alpha) = (0.1, 4.67 \times 10^{-5})$, and $(\epsilon_1, \epsilon_2) = (10^{-4}, 0.8)$. Figure 7 illustrates the behaviors of the objective function and the smallest eigenvalue of $\mathbf{H}_{\mathbf{K}, \mathbf{L}}$ following the double-loop update scheme. It is shown that iterates of the outer-loop gain matrix $\{\mathbf{K}_k\}_{k \geq 0}$ stay within the feasible set \mathcal{K} along with the double-loop update, which preserves a certain disturbance attenuation level in the view of Remark A.9. Moreover, in both cases, we have Algorithm 2 converging sublinearly to the unique Nash equilibrium.

D.4. Divergent Cases

As noted in Remark A.7, it is unclear yet if descent-multi-step-ascent, AGDA, or τ -GDA, where $\tau = \eta/\alpha$, can converge globally to the Nash equilibrium in our setting. In this section, we present some scenarios where descent-multi-step-ascent, AGDA, and τ -GDA diverges even with infinitesimal stepsizes. We use the same example as §D.1 and initialize $\mathbf{K}_0^4, \mathbf{L}_0^4$ to be

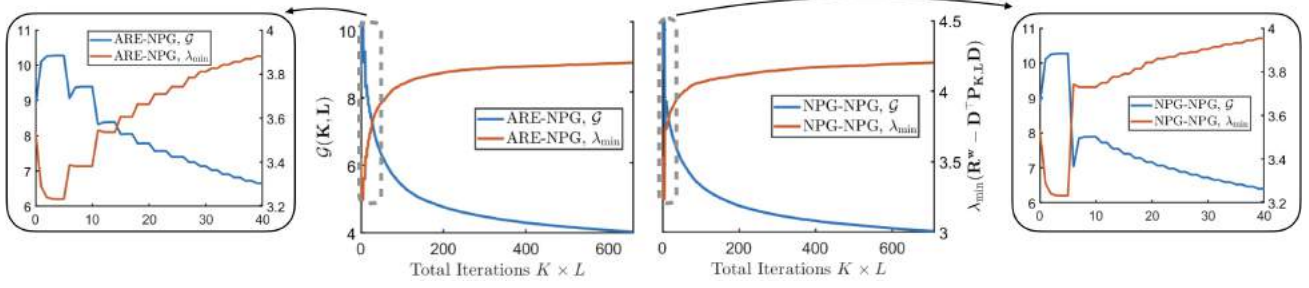


Figure 7. Convergence of $\mathcal{G}(\mathbf{K}, \mathbf{L})$ and $\lambda_{\min}(\mathbf{H}_{\mathbf{K}, \mathbf{L}})$ when the outer loop follows the ZO-NPG update (4.3). *Left:* The inner-loop solutions are computed by the expression of $\mathbf{L}(\mathbf{K})$ in Lemma 3.3. *Right:* The inner-loop solutions are approximated by the exact NPG update (3.10) with an accuracy level of $\epsilon_1 = 10^{-4}$.

time-invariant such that $K_{0,t}^4 = K_0^4 = \begin{bmatrix} -0.1362 & 0.0934 & 0.6458 \\ -0.2717 & -0.1134 & -0.4534 \\ -0.6961 & -0.9279 & -0.6620 \end{bmatrix}$ and $L_{0,t}^4 = L_0^4 = \begin{bmatrix} 0.2887 & -0.2286 & 0.4588 \\ -0.7849 & -0.1089 & -0.3755 \\ -0.2935 & 0.9541 & 0.7895 \end{bmatrix}$

for all t . Note that in descent-multi-step-ascent/AGDA/ τ -GDA, the maximizing problem with respect to \mathbf{L} for a fixed \mathbf{K} is no longer solved to a high accuracy for each iteration of the updates on \mathbf{K} . Thus, we can relax the constraint $\mathbf{K} \in \mathcal{K}$ since the maximizing player will not drive the cost to ∞ in such iterates under descent-multi-step-ascent/AGDA/ τ -GDA. Figure 8 illustrates the behaviors of the objective and the gradient norms with respect to \mathbf{K} and \mathbf{L} when applying alternating natural GDA (ANGDA), τ -natural GDA (τ -NGDA), and descent-multi-step-ascent with updates (3.10) and (3.11) to our problem. The stepsizes for the ANGDA are chosen to be infinitesimal such that $\eta = \alpha = 1.7319 \times 10^{-11}$. However, we can still observe the diverging patterns from the top row of Figure 8, even with such tiny stepsizes. Further, we test the same example with the same initialization but use τ -NGDA with $\tau = 10^3$. That is, $\eta = 10^3 \cdot \alpha$. In such case, it is shown in the middle row of Figure 8 that diverging behaviors still appear. Whether there exists a finite timescale separation τ^* , similar to the one proved in (Fiez & Ratliff, 2020), such that τ -GDA type algorithms with $\tau \in (\tau^*, \infty)$ provably converge to the unique Nash equilibrium in our setting requires further investigation, and is left as our future work. Lastly, we present a case where descent-multi-step-ascent with updates following (3.10) and (3.11) diverges. The stepsizes of the inner-loop and the outer-loop updates are chosen to be $\eta = \alpha = 1.7319 \times 10^{-9}$ and we run 10 iterations of (3.10) for each iteration of (3.11). The bottom row of Figure 8 demonstrates the diverging behaviors of this update scheme.

D.5. Time-Varying Systems

Instead of setting $A_t = A$, $B_t = B$, and $D_t = D$ for all t , we now choose

$$A_t = A + \frac{(-1)^t t A}{10}, \quad B_t = B + \frac{(-1)^t t B}{10}, \quad D_t = D + \frac{(-1)^t t D}{10},$$

and set $R^w = 10 \cdot \mathbf{I}$. Rest of the parameters are set to the same as in (D.1). The horizon of the problem is again set to

$$N = 5 \text{ and we initialize our algorithm using } K_{0,t}^5 = K_0^5 = \begin{bmatrix} -0.0984 & -0.7158 & -0.1460 \\ -0.1405 & 0.0039 & 0.4544 \\ -0.1559 & -0.7595 & 0.7403 \end{bmatrix}.$$

We demonstrate in Figure 9 the convergence of exact double-loop NPG updates. The stepsizes of the inner-loop and the outer-loop updates are chosen to be $(\eta, \alpha) = (0.0097, 3.0372 \times 10^{-5})$. Also, we require the approximate inner-loop solution to have the accuracy of $\epsilon_1 = 0.001$. As shown at the top of Figure 9, the double-loop NPG updates successfully converge to the unique Nash equilibrium of the game. Also, the minimum eigenvalue of $R^w - D^\top P_{\mathbf{K}, \mathbf{L}} D$ is monotonically non-decreasing along the iterations of the outer-loop NPG update, which matches the implicit regularization property we introduced in Theorem 3.7. This guarantees that all future iterates of \mathbf{K} will stay in the interior of \mathcal{K} given that the initial \mathbf{K} does. The convergences of the average gradient norms with respect to both \mathbf{L} and \mathbf{K} are also presented at the bottom of Figure 9.

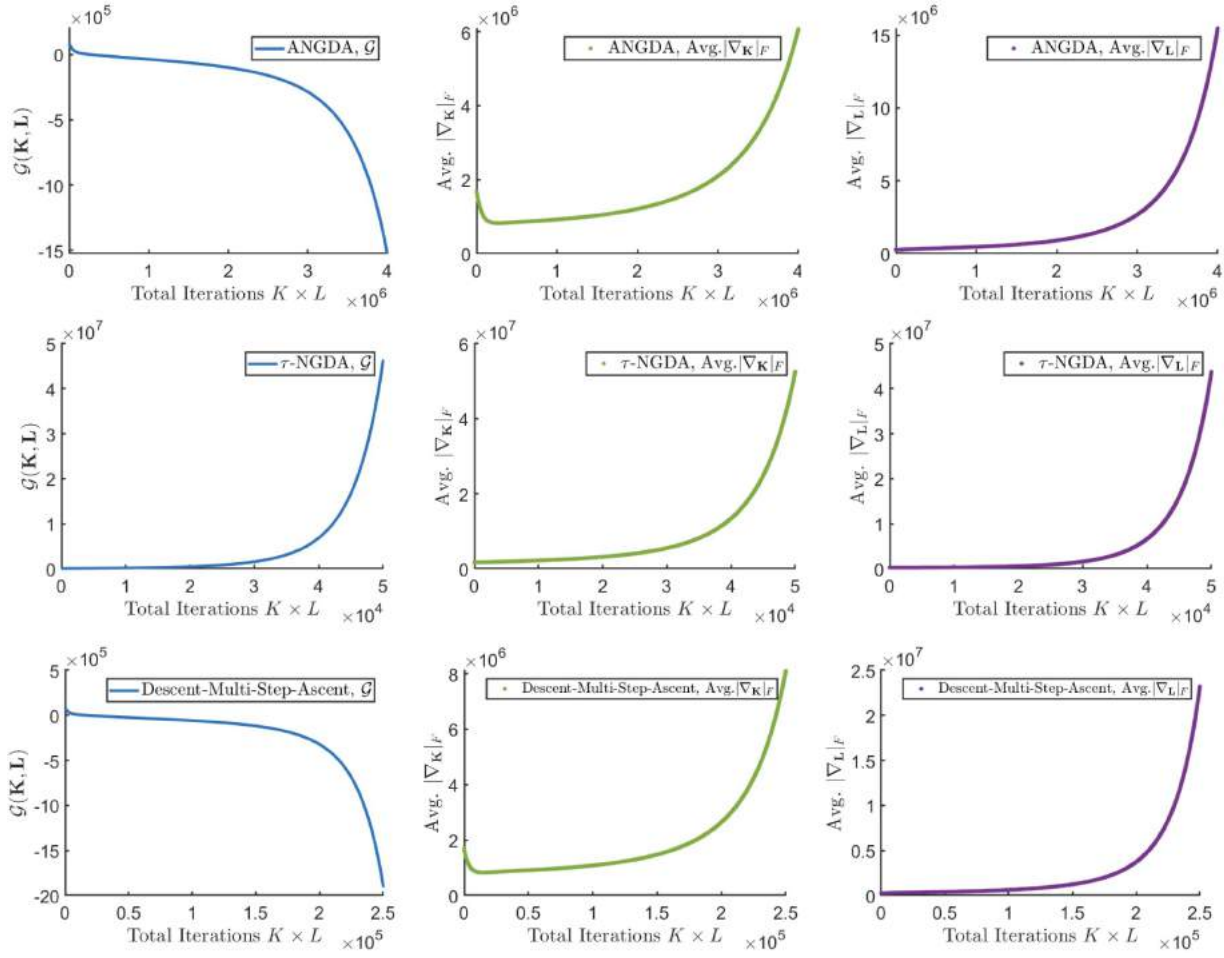


Figure 8. Divergence of $\mathcal{G}(\mathbf{K}, \mathbf{L})$, $\{\ell^{-1} \sum_{l=0}^{\ell-1} \|\nabla_{\mathbf{L}}(\mathbf{K}, \mathbf{L}_l)\|_F\}_{l \geq 1}$, and $\{k^{-1} \sum_{\kappa=0}^{k-1} \|\nabla_{\mathbf{K}} \mathcal{G}(\mathbf{K}_{\kappa}, \bar{\mathbf{L}}(\mathbf{K}_{\kappa}))\|_F\}_{k \geq 1}$. Top: ANGDA with stepsizes $\eta = \alpha = 1.7319 \times 10^{-11}$. Middle: τ -NGDA with stepsizes $\alpha = 1.7319 \times 10^{-11}$ and $\eta = 10^3 \cdot \alpha$. Bottom: Descent-multi-step-ascent with update rules (3.10) and (3.11) and stepsizes $\eta = \alpha = 1.7319 \times 10^{-9}$. For each iteration of the outer-loop update (3.11), we run 10 iterations of the inner-loop update (3.10).




Figure 9. Convergence of the exact double-loop NPG updates with initial gain matrix being \mathbf{K}_0^5 . *Top*: Convergence of $\mathcal{G}(\mathbf{K}, \mathbf{L})$ and $\lambda_{\min}(\mathbf{H}_{\mathbf{K}, \mathbf{L}})$. *Bottom*: Convergence of $\{k^{-1} \sum_{\kappa=0}^{k-1} \|\nabla_{\mathbf{K}} \mathcal{G}(\mathbf{K}_{\kappa}, \bar{\mathbf{L}}(\mathbf{K}_{\kappa}))\|_F\}_{k \geq 1}$ and $\{l^{-1} \sum_{\iota=0}^{l-1} \|\nabla_{\mathbf{L}}(\mathbf{K}, \mathbf{L}_{\iota})\|_F\}_{l \geq 1}$. The stepsizes of the inner-loop and the outer-loop NPG updates are chosen to be 0.0097 and 3.0372×10^{-5} , respectively. For each iteration of the outer-loop NPG update, we solve the inner-loop subproblem to the accuracy of $\epsilon_1 = 0.001$.